



Lecture 6: Clustering and Segmentation – Part 2

Professor Fei-Fei Li

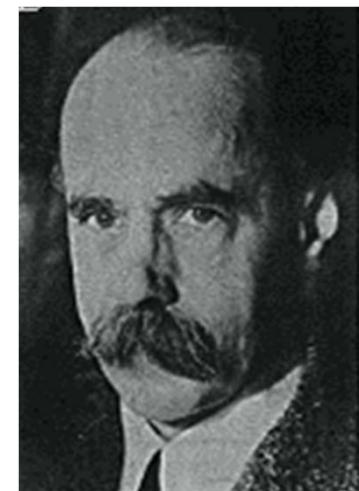
Stanford Vision Lab

Recap: Gestalt Theory

- Gestalt: whole or group
 - Whole is greater than sum of its parts
 - Relationships among parts can yield new properties/features
- Psychologists identified series of factors that predispose set of elements to be grouped (by human visual system)

*"I stand at the window and see a house, trees, sky.
Theoretically I might say there were 327 brightnesses
and nuances of colour. Do I have "327"? No. I have sky, house,
and trees."*

Max Wertheimer
(1880-1943)



Untersuchungen zur Lehre von der Gestalt,
Psychologische Forschung, Vol. 4, pp. 301-350, 1923
<http://psy.ed.asu.edu/~classics/Wertheimer/Forms/forms.htm>

Recap: Gestalt Factors



Not grouped



Proximity



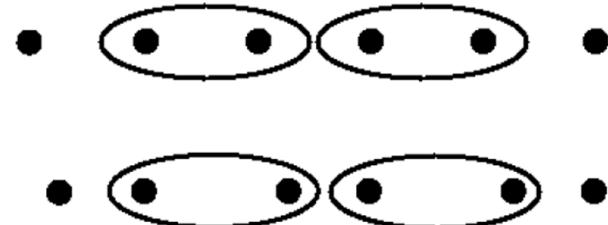
Similarity



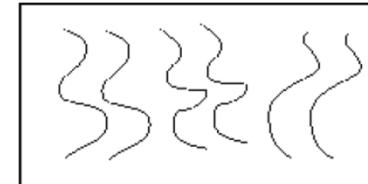
Similarity



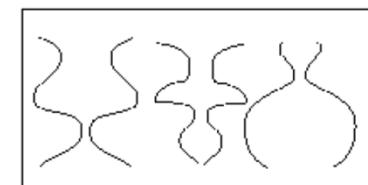
Common Fate



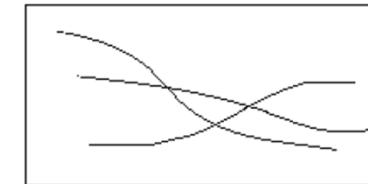
Common Region



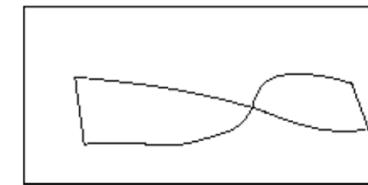
Parallelism



Symmetry



Continuity

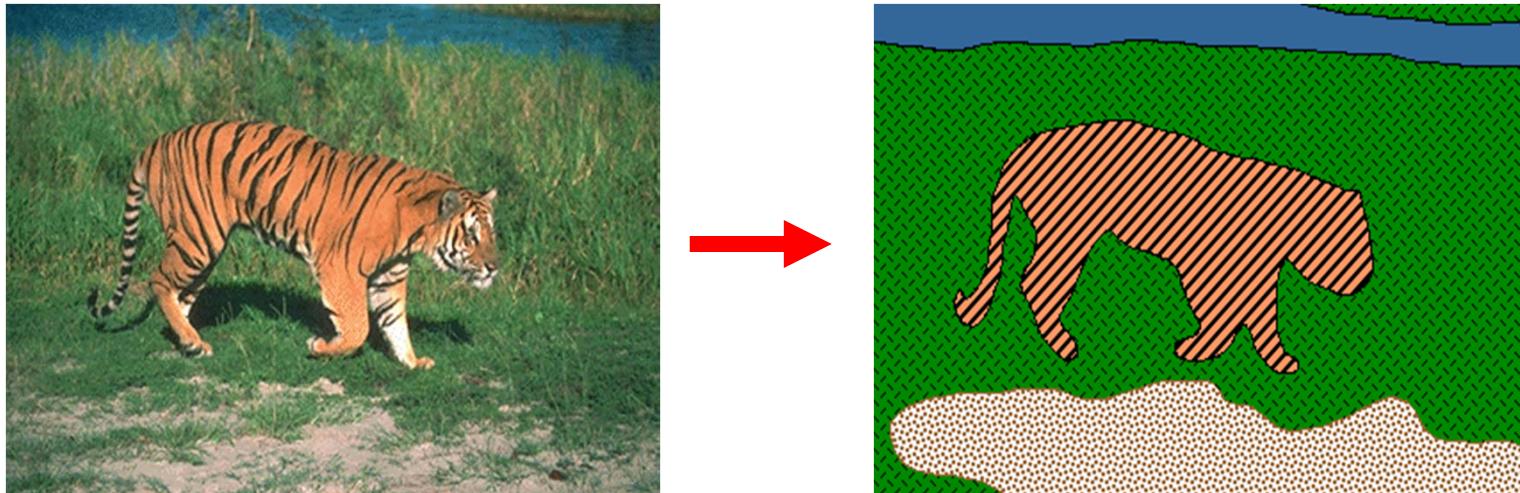


Closure

- These factors make intuitive sense, but are very difficult to translate into algorithms.

Recap: Image Segmentation

- Goal: identify groups of pixels that go together



Recap: K-Means Clustering

- Basic idea: randomly initialize the k cluster centers, and iterate between the two steps we just saw.

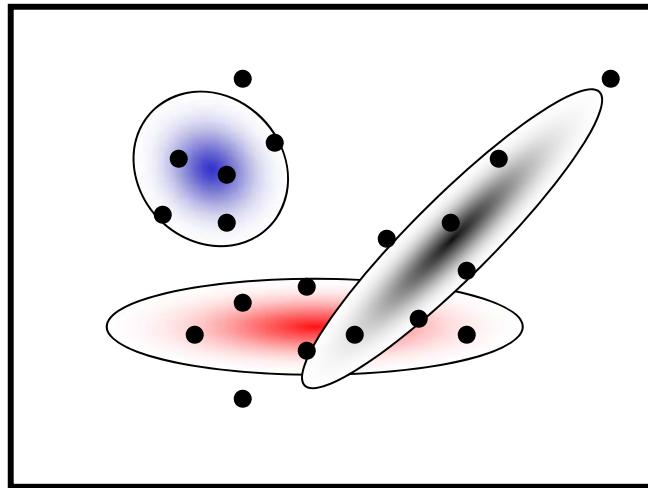
1. Randomly initialize the cluster centers, c_1, \dots, c_K
2. Given cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
4. If c_i have changed, repeat Step 2



- Properties
 - Will always converge to *some* solution
 - Can be a “local minimum”
 - Does not always find the global minimum of objective function:

$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Recap: Expectation Maximization (EM)



- Goal
 - Find blob parameters θ that maximize the likelihood function:
$$P(data|\theta) = \prod_x P(x|\theta)$$
- Approach:
 1. E-step: given current guess of blobs, compute ownership of each point
 2. M-step: given ownership probabilities, update blobs to maximize likelihood function
 3. Repeat until convergence

What we will learn today?

- Model free clustering
 - Mean-shift
- Graph theoretic segmentation
 - Normalized Cuts
 - Using texture features
- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation (supp. materials)
 - s-t mincut algorithm (supp. materials)
 - Extension to non-binary case (supp. materials)
 - Applications

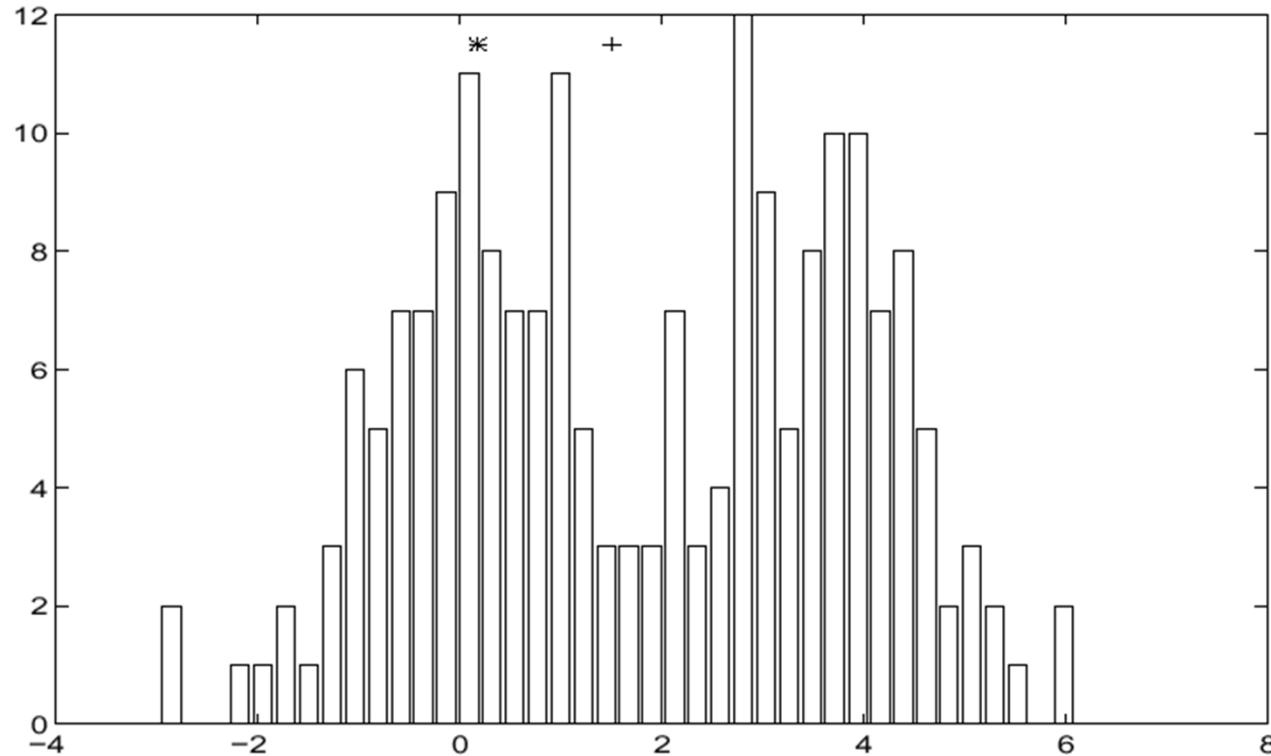


(Midterm materials)

What we will learn today?

- “Model free” clustering
 - Mean-shift
- Graph theoretic segmentation
 - Normalized Cuts
 - Using texture features
- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation (supp. materials)
 - s-t mincut algorithm (supp. materials)
 - Extension to non-binary case (supp. materials)
 - Applications

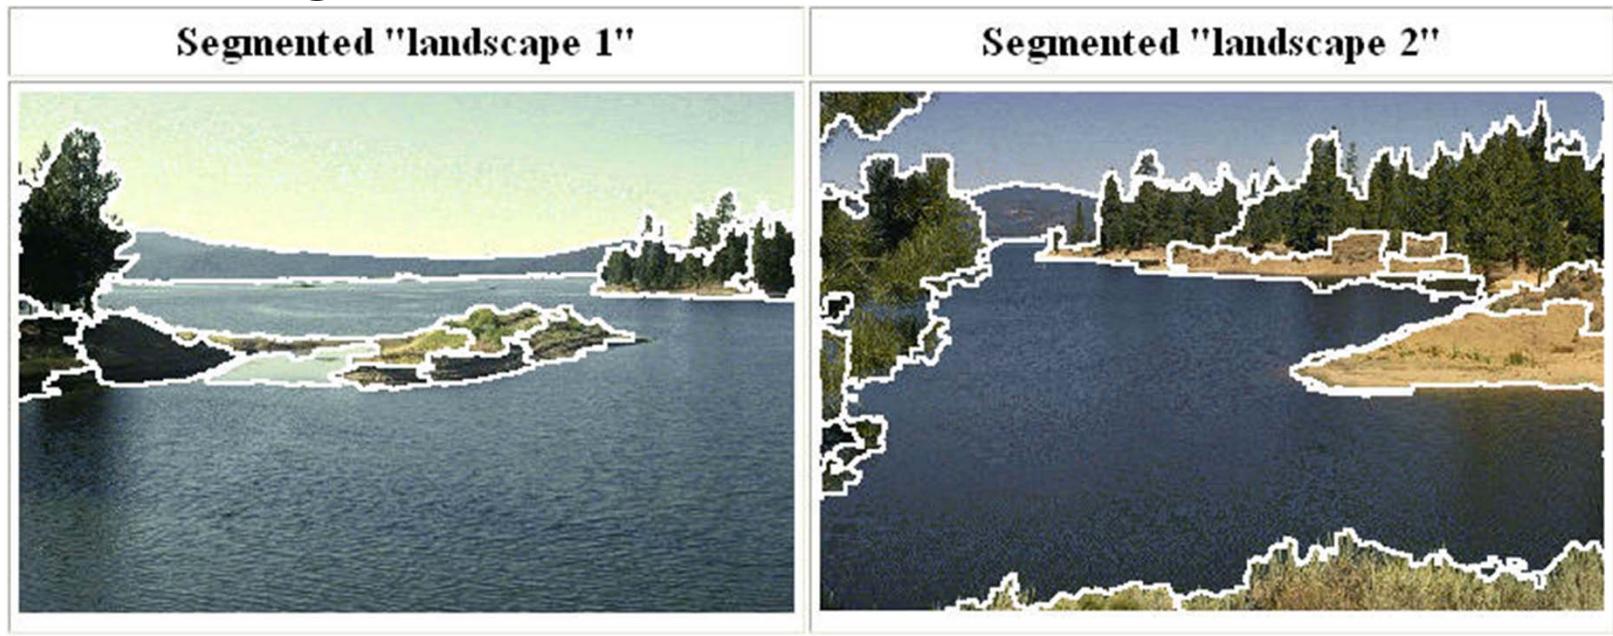
Finding Modes in a Histogram



- How many modes are there?
 - *Mode* = local maximum of the density of a given distribution
 - Easy to see, hard to compute

Mean-Shift Segmentation

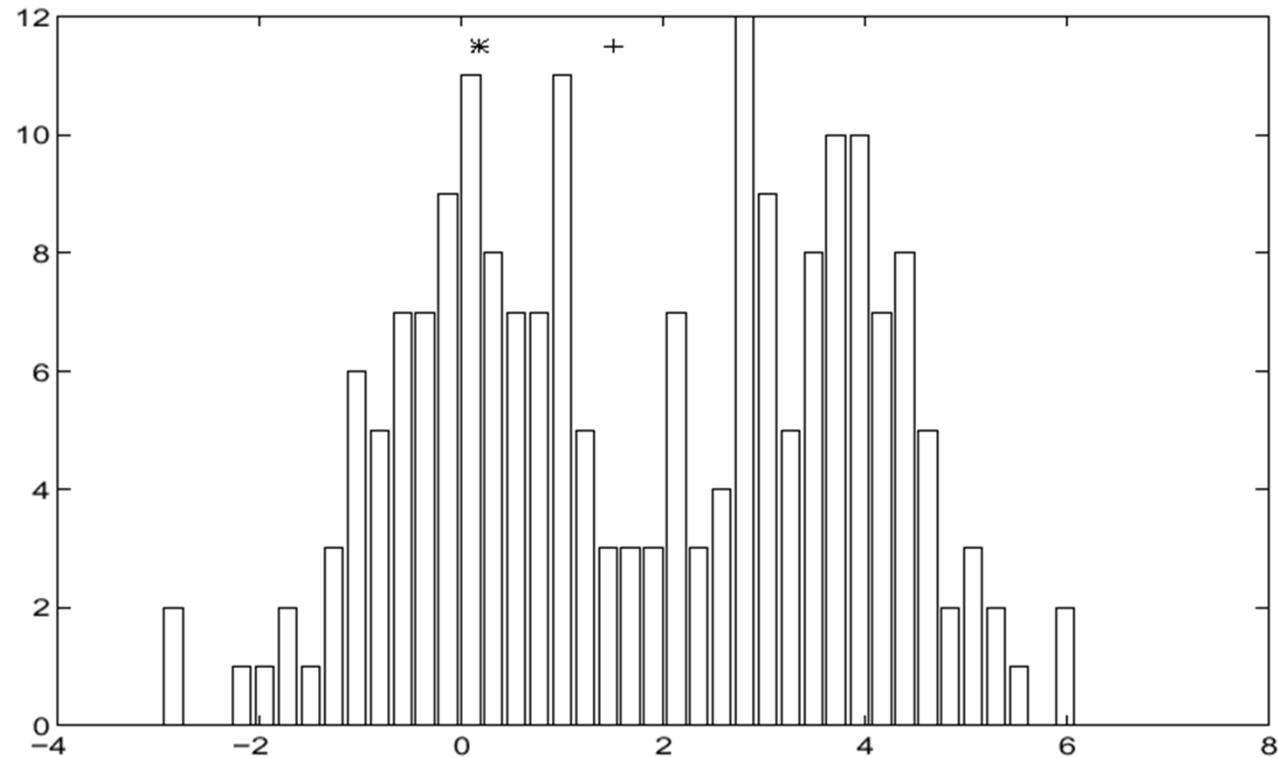
- An advanced and versatile technique for clustering-based segmentation



<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>

D. Comaniciu and P. Meer, [Mean Shift: A Robust Approach toward Feature Space Analysis](#), PAMI 2002.

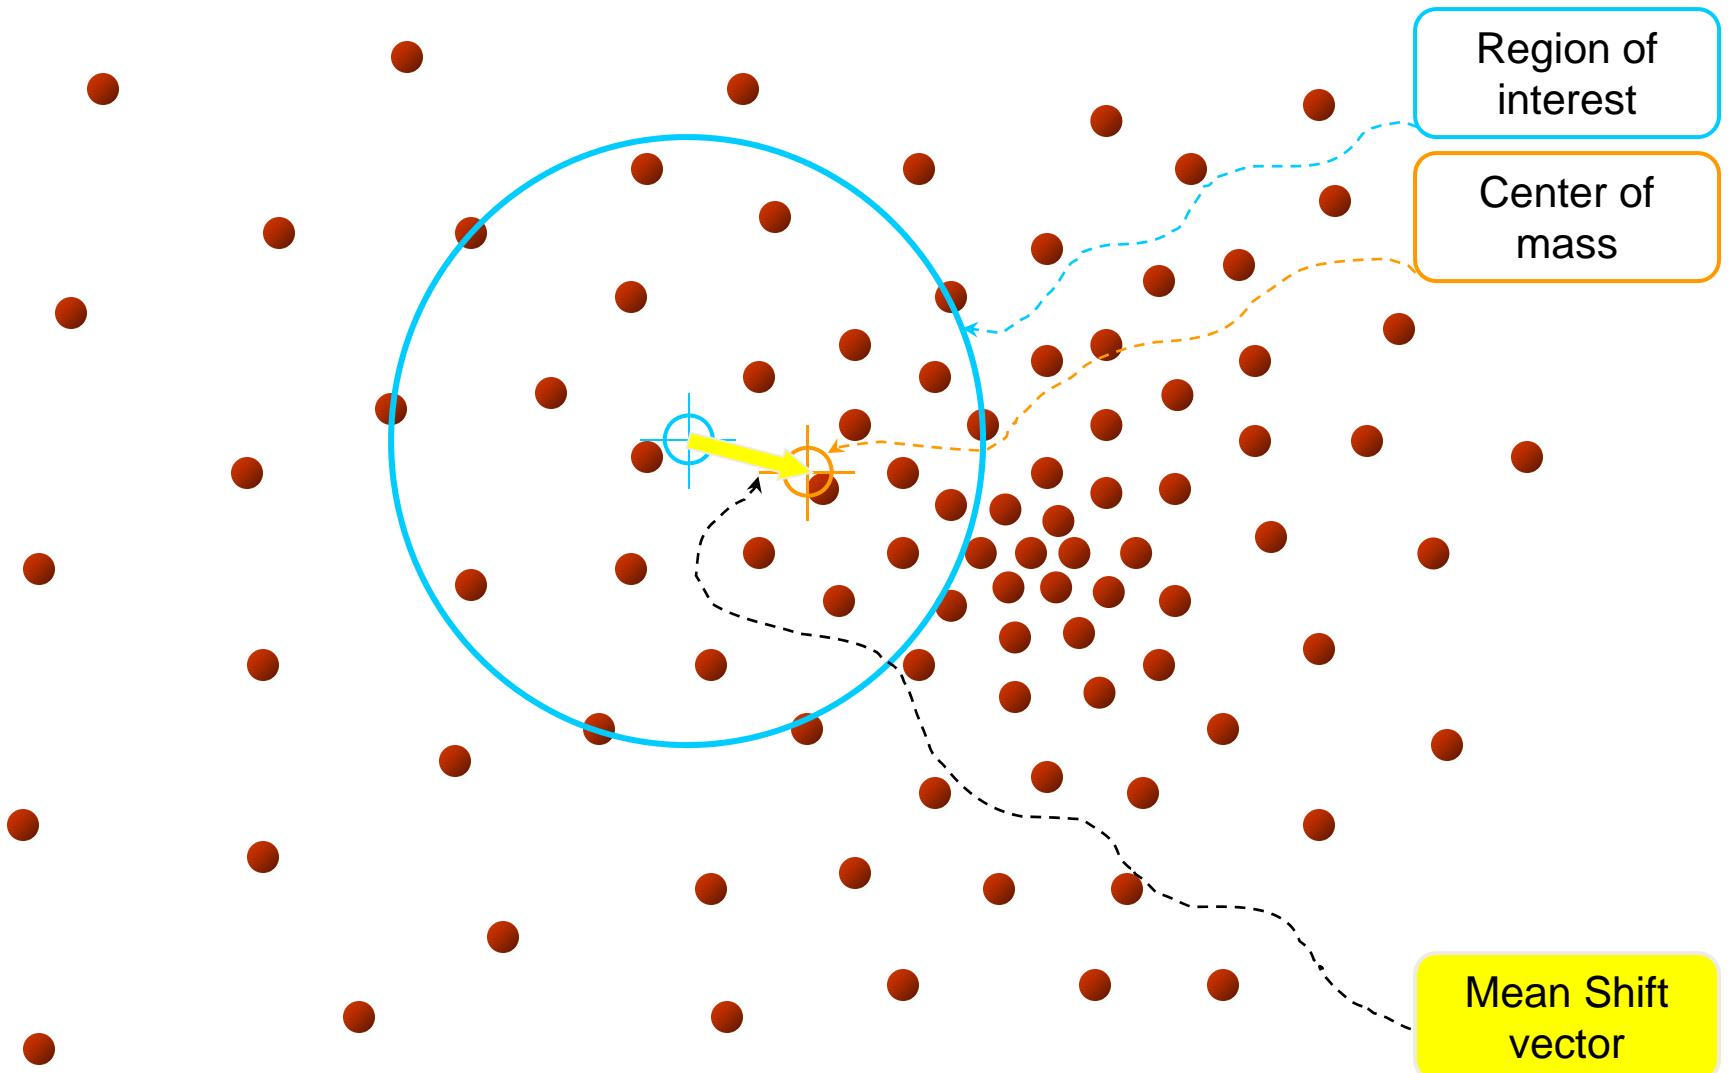
Mean-Shift Algorithm



- Iterative Mode Search

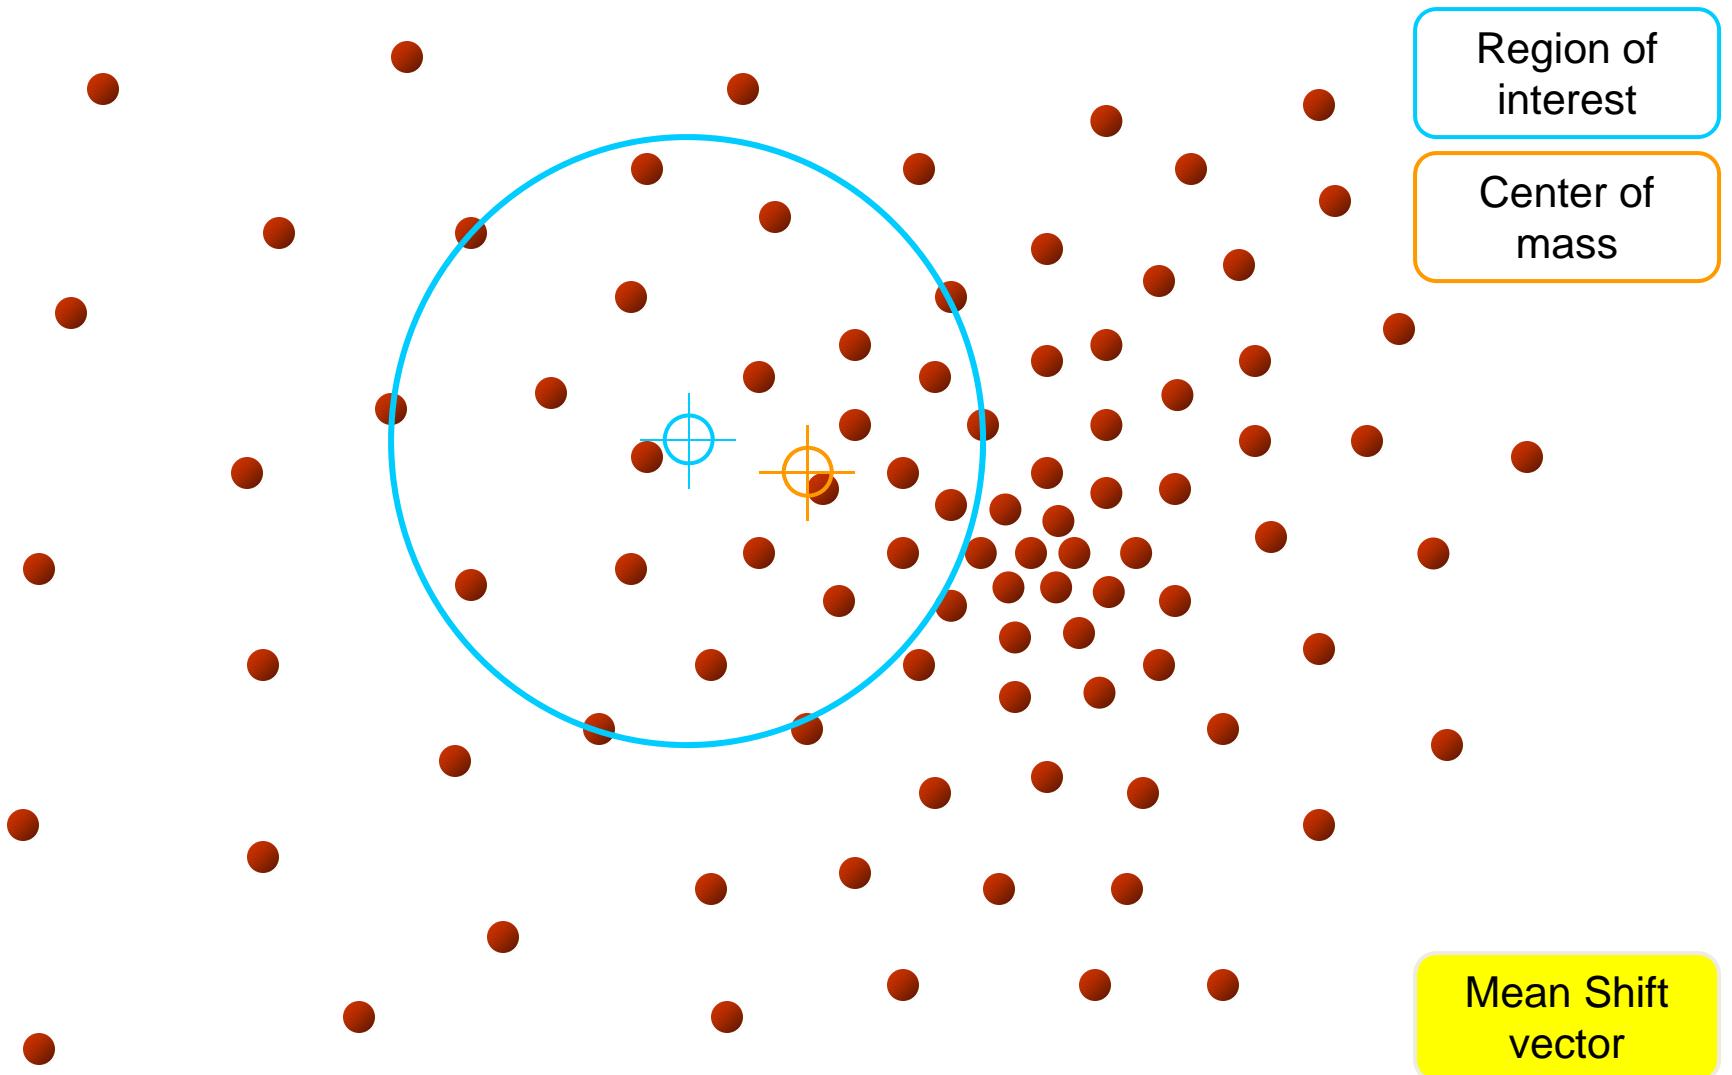
1. Initialize random seed, and window W
2. Calculate center of gravity (the “mean”) of W : $\sum_{x \in W} xH(x)$
3. Shift the search window to the mean
4. Repeat Step 2 until convergence

Mean-Shift



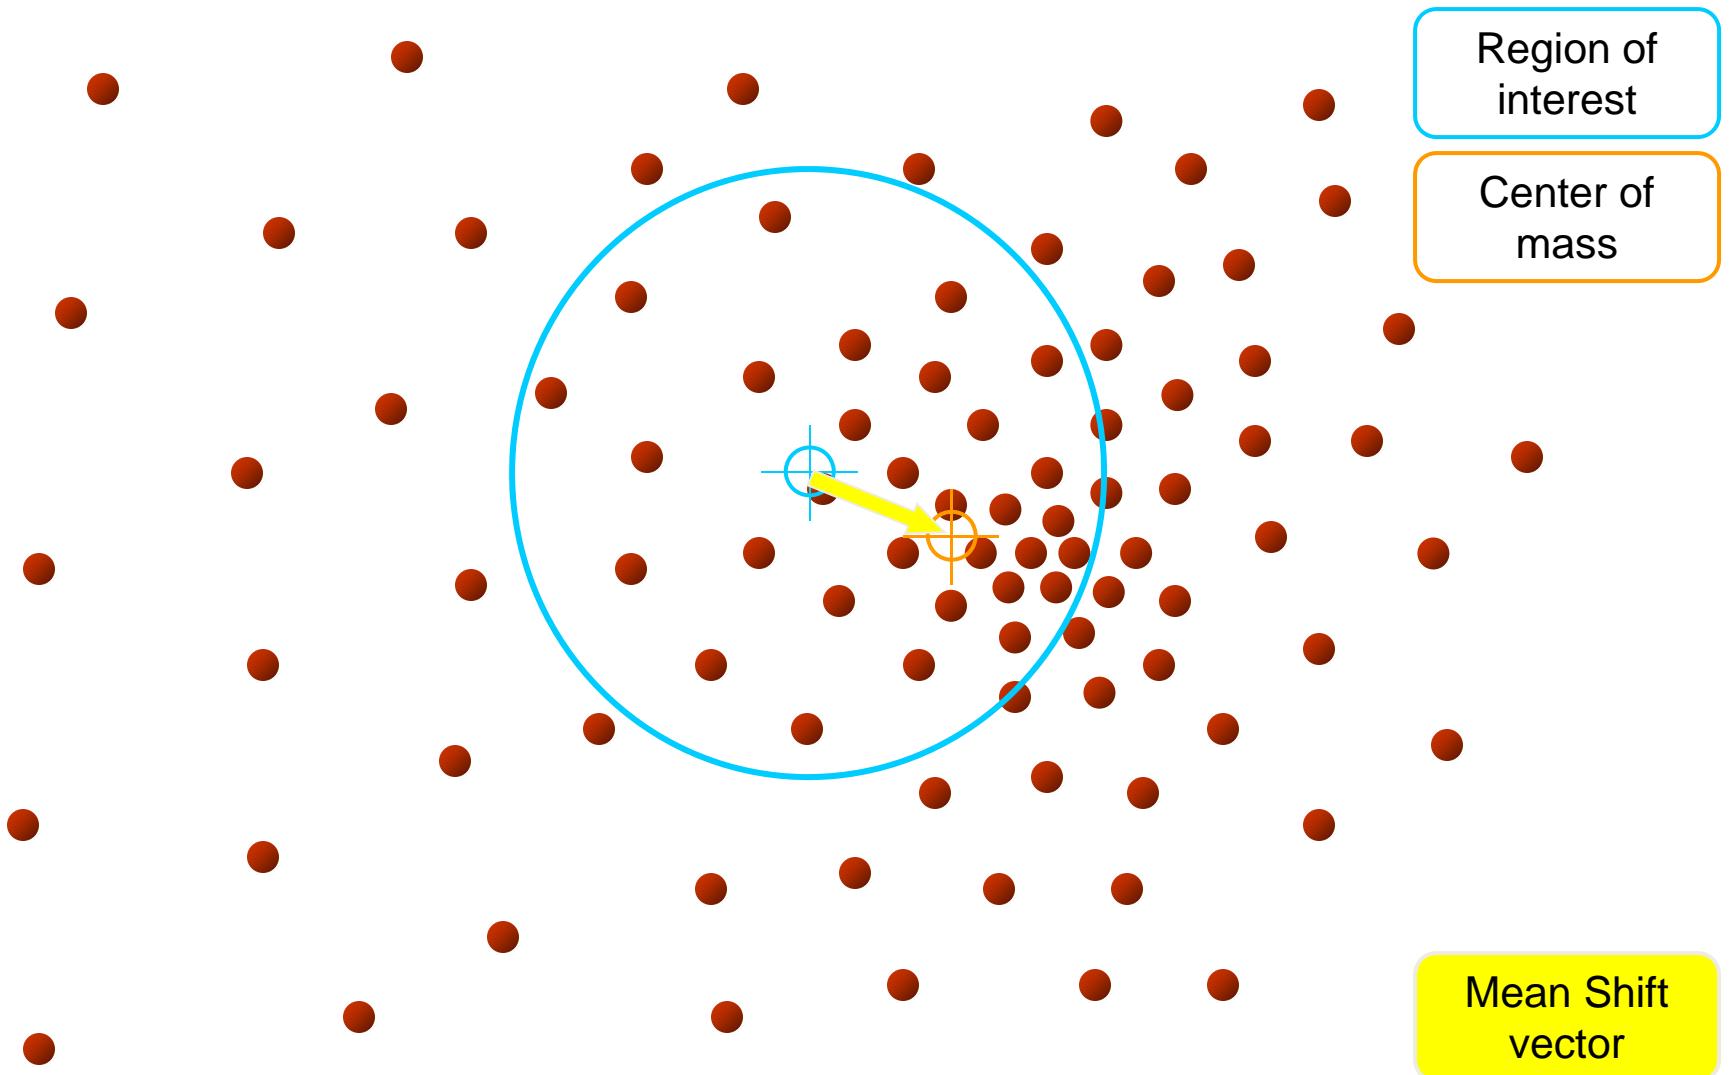
Slide by Y. Ukrainitz & B. Sarel

Mean-Shift



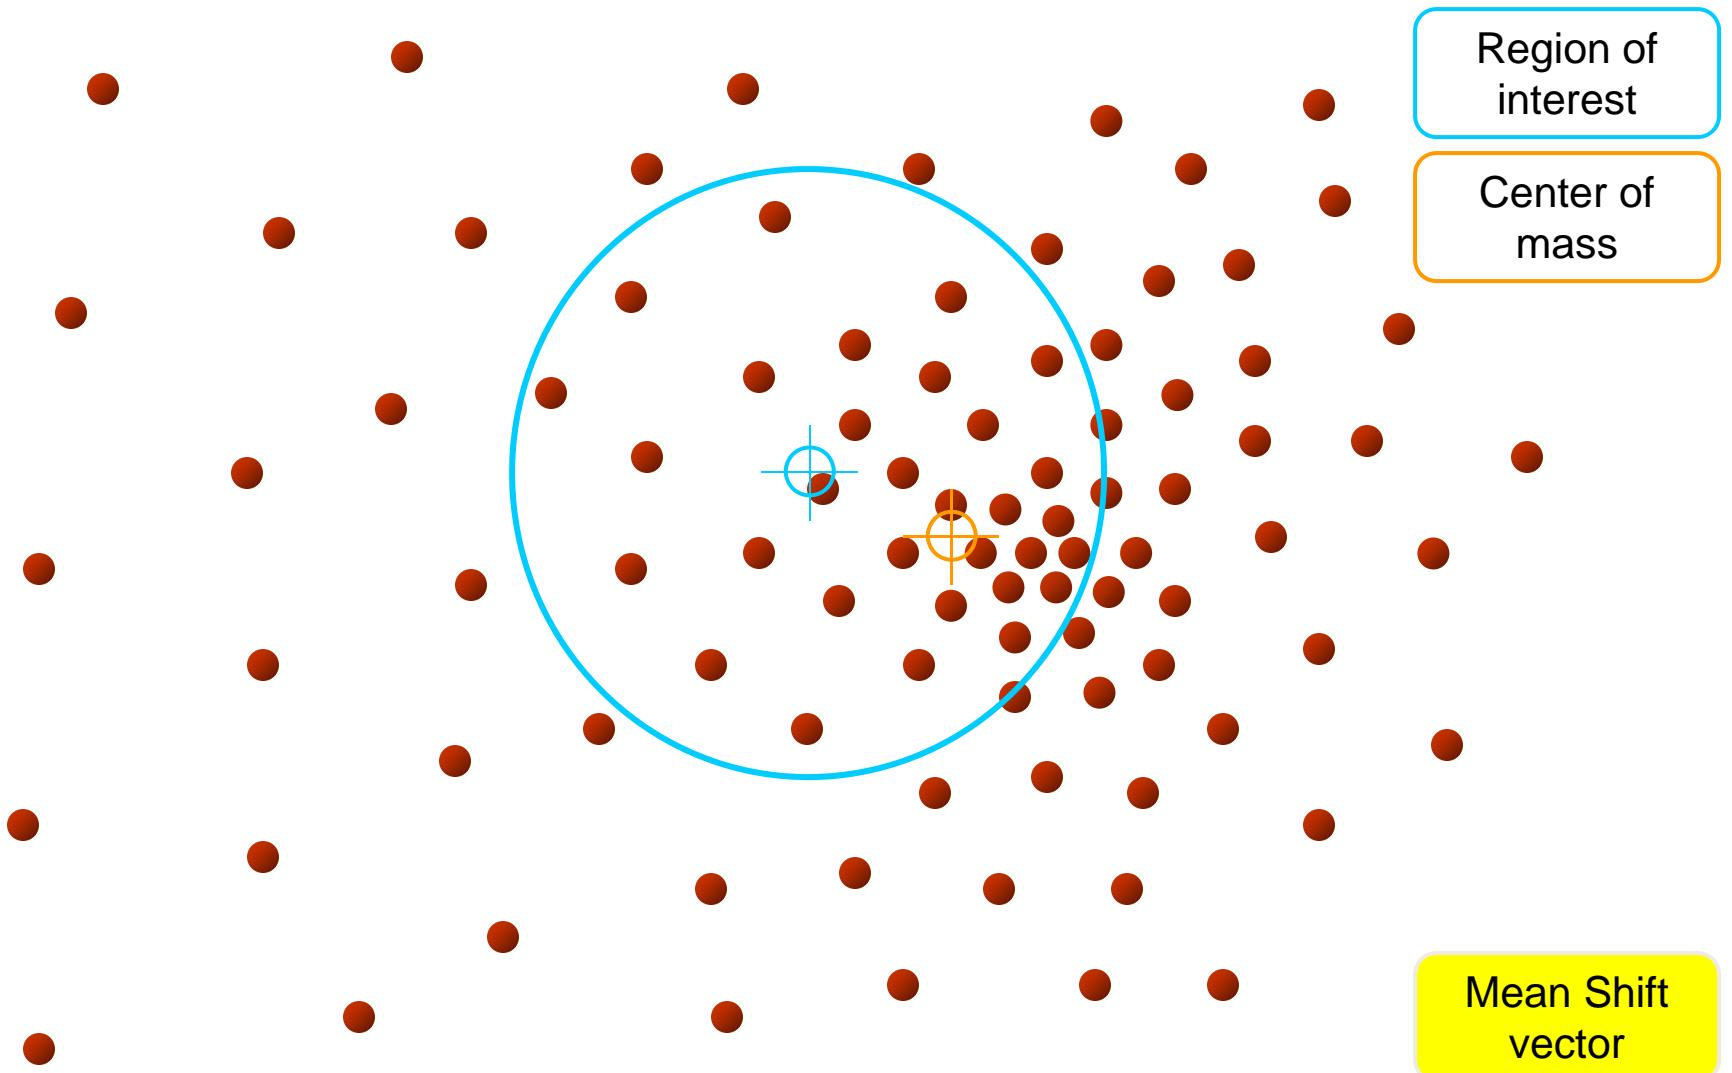
Slide by Y. Ukrainitz & B. Sarel

Mean-Shift



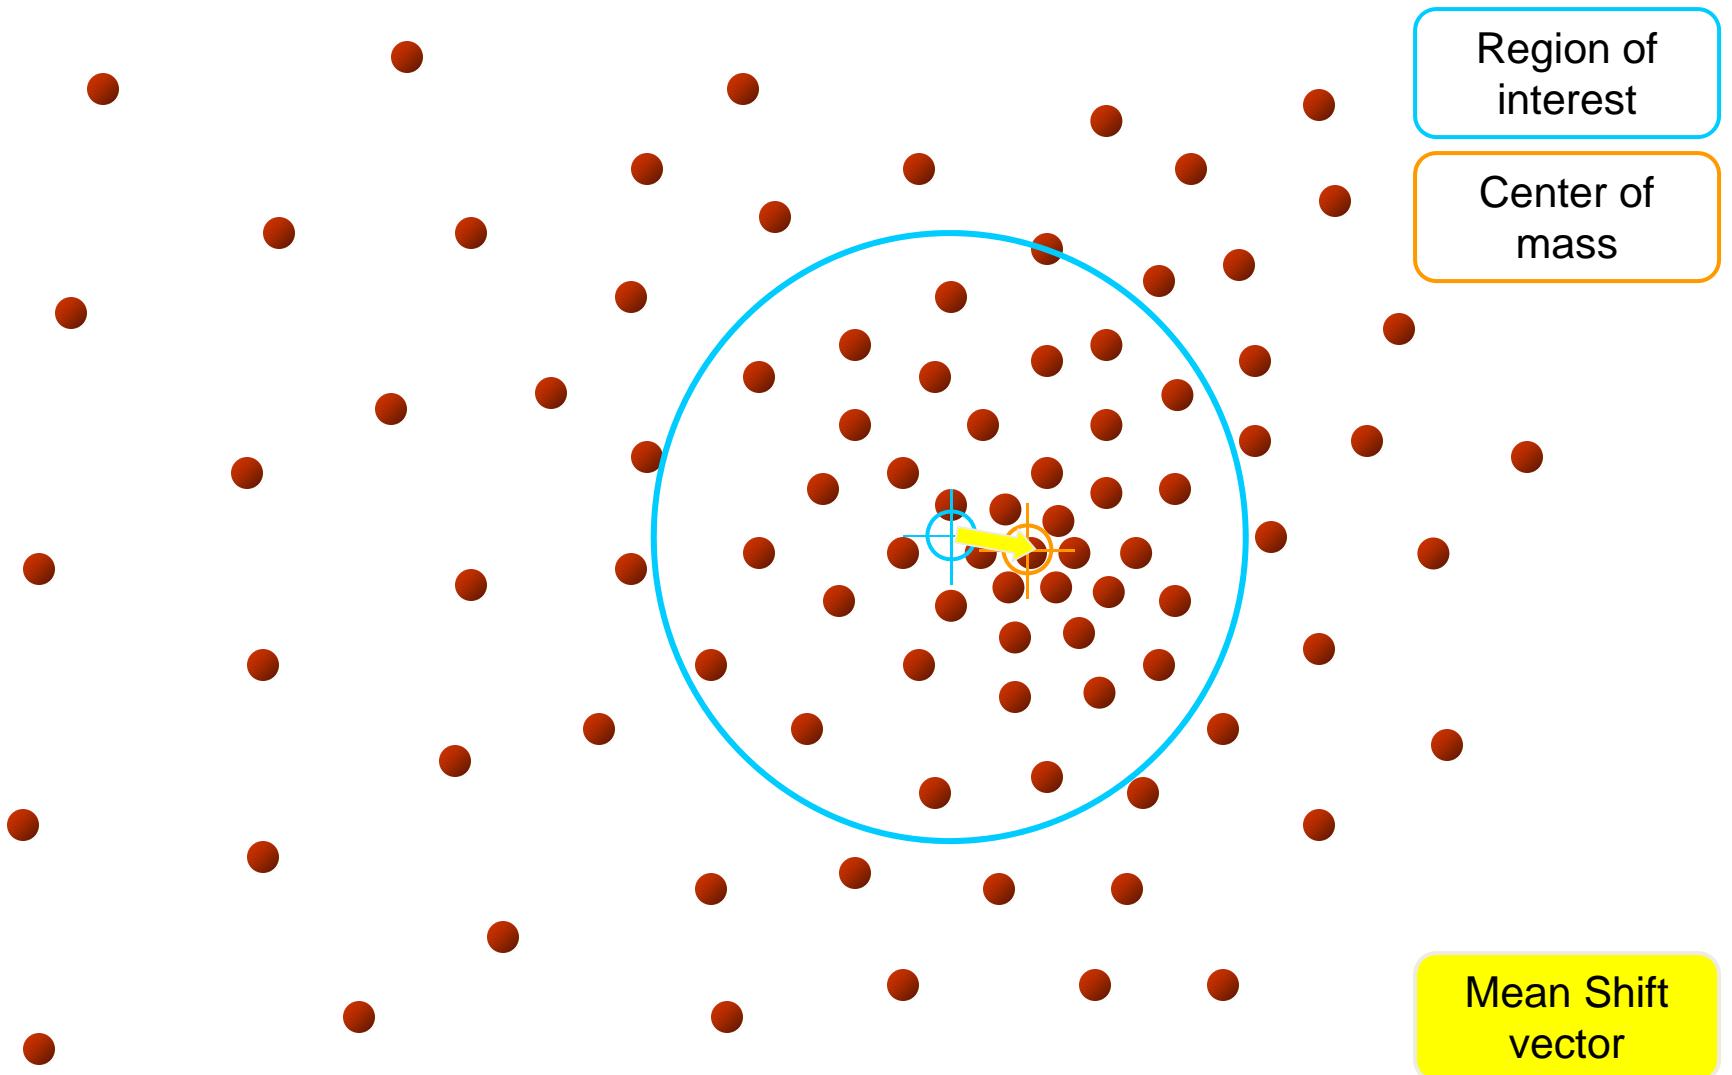
Slide by Y. Ukrainitz & B. Sarel

Mean-Shift



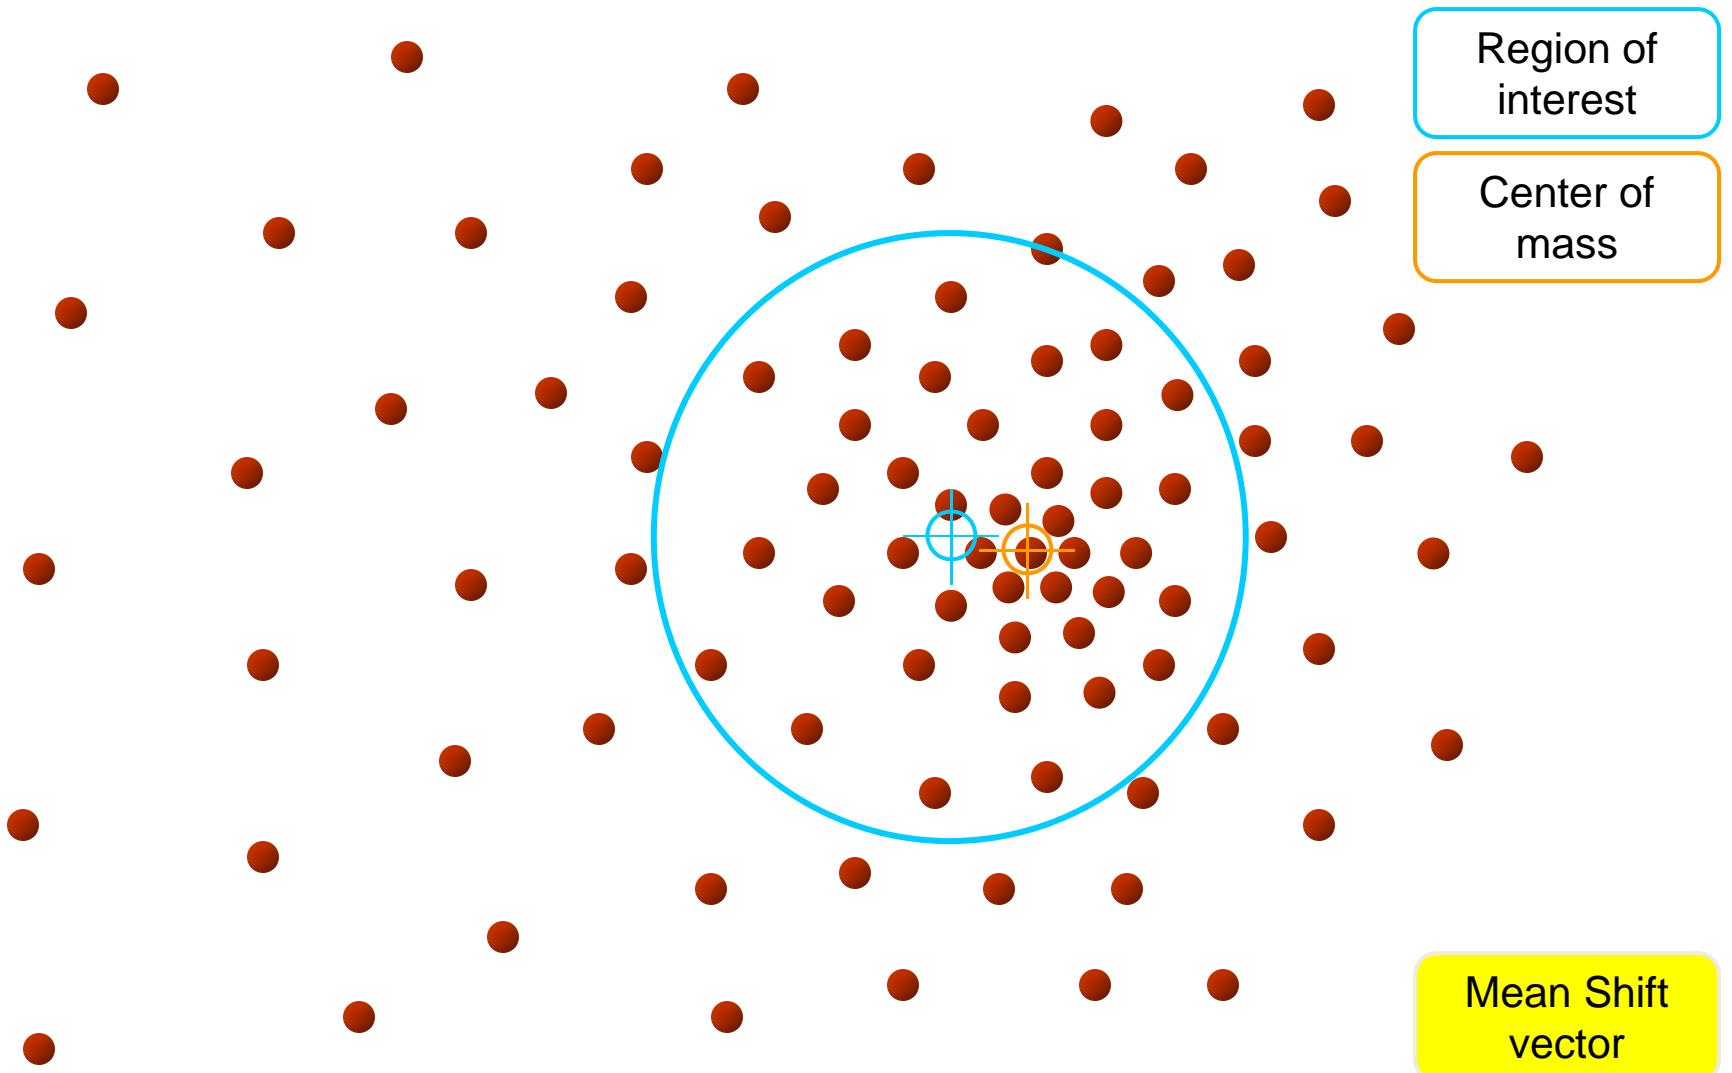
Slide by Y. Ukrainitz & B. Sarel

Mean-Shift



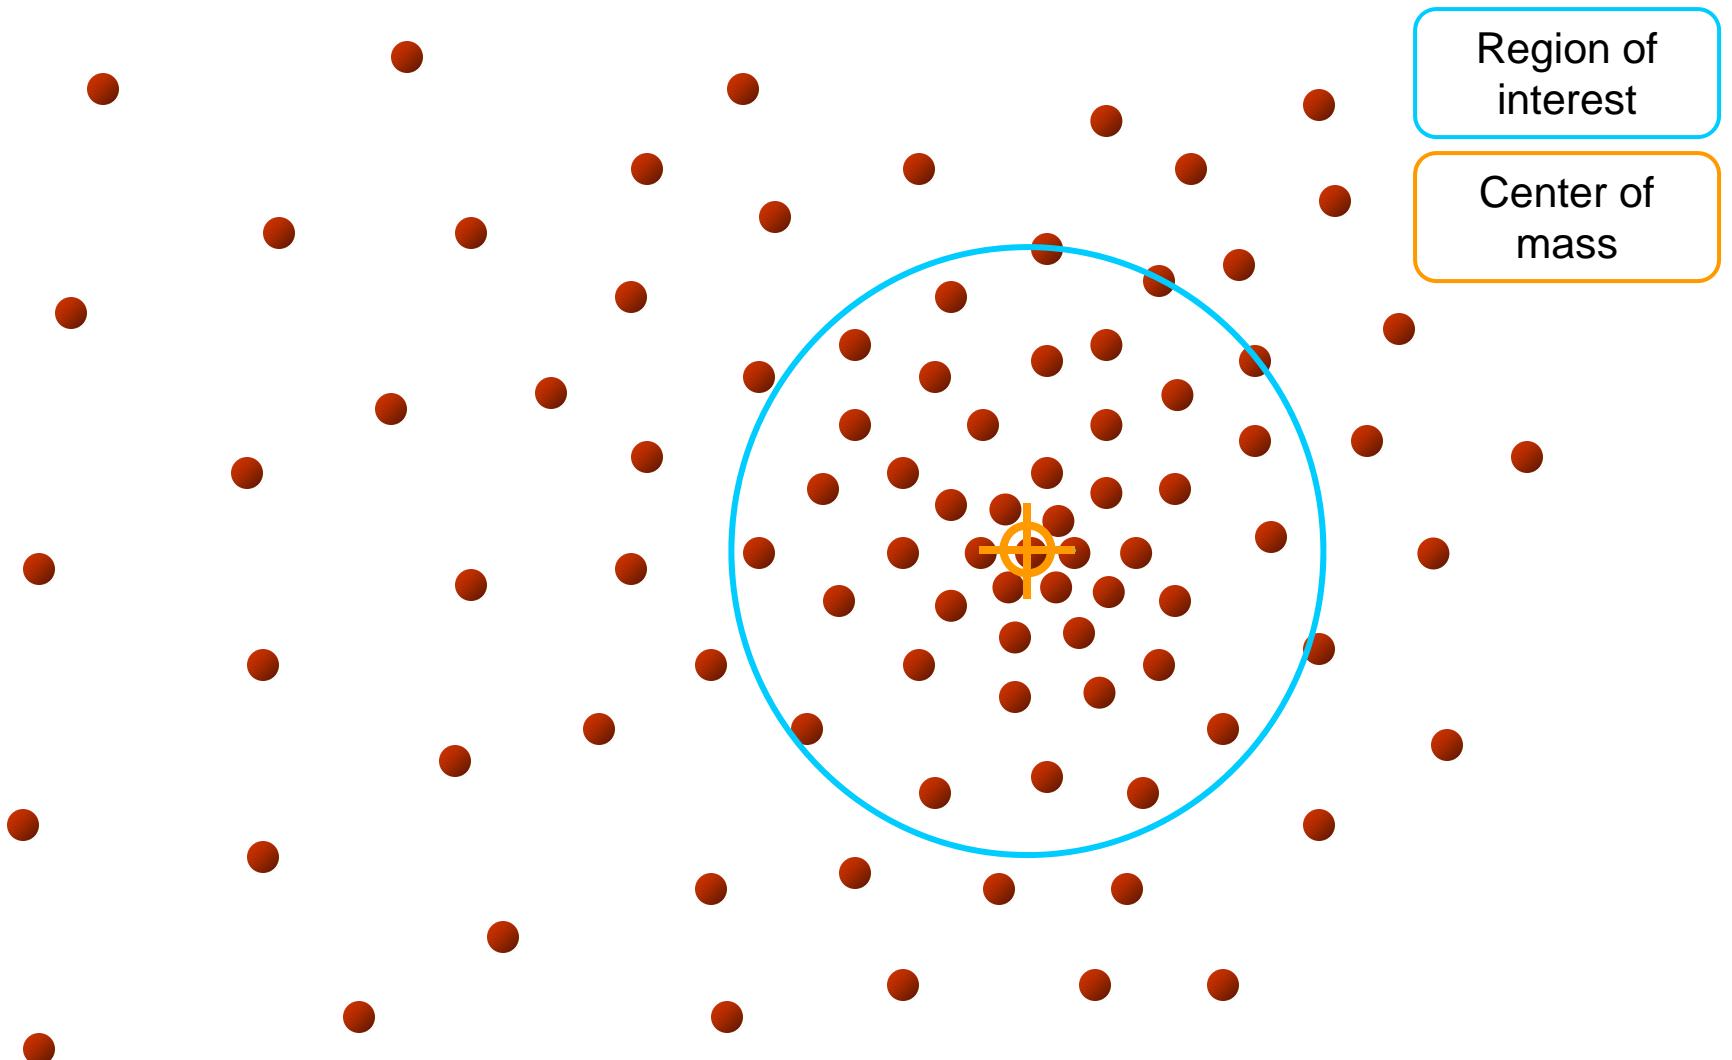
Slide by Y. Ukrainitz & B. Sarel

Mean-Shift



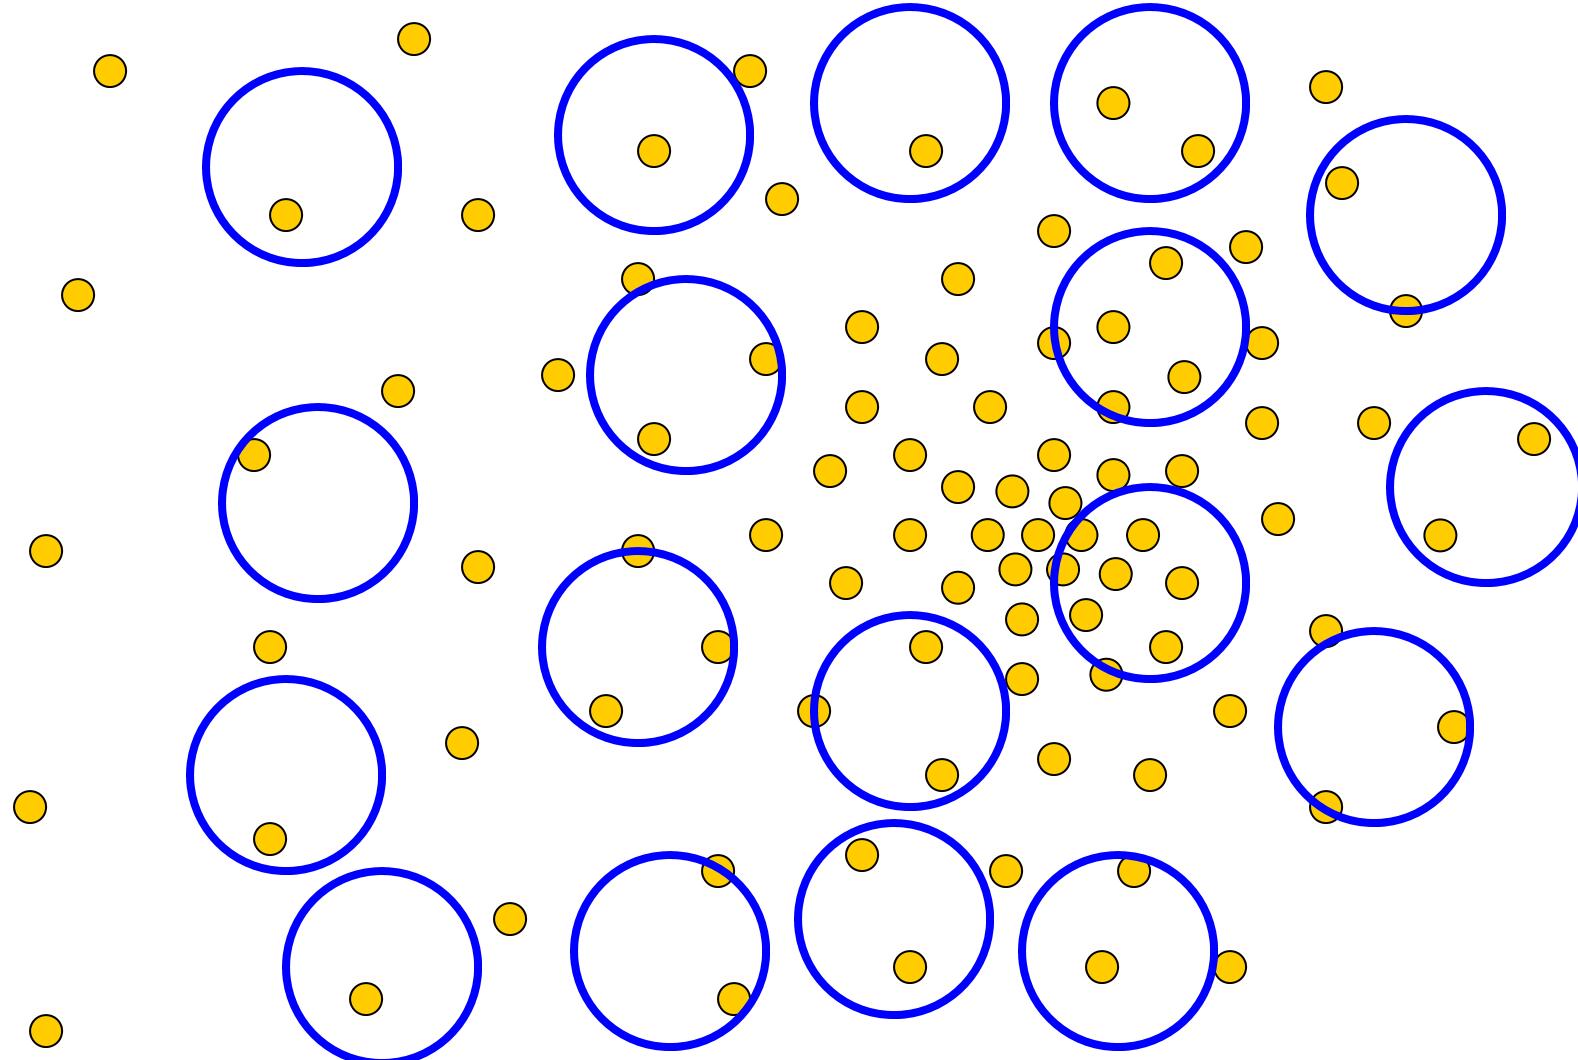
Slide by Y. Ukrainitz & B. Sarel

Mean-Shift



Slide by Y. Ukrainitz & B. Sarel

Real Modality Analysis

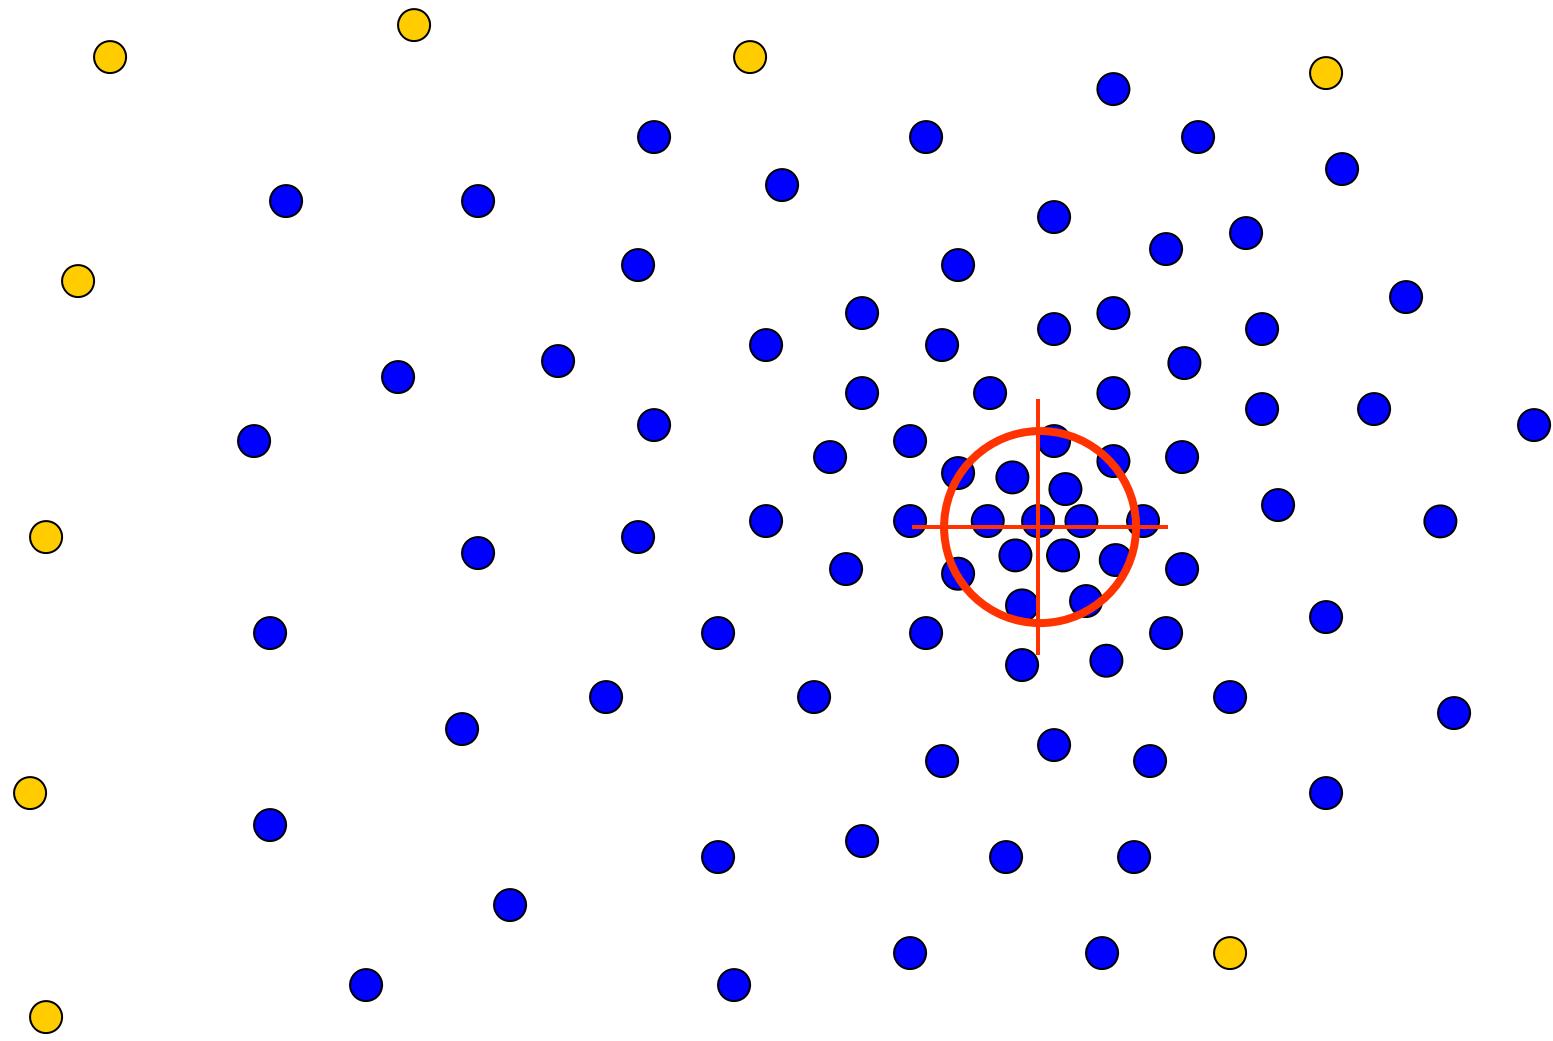


Tessellate the space with windows

Run the procedure in parallel

Slide by Y. Ukrainitz & B. Sarel

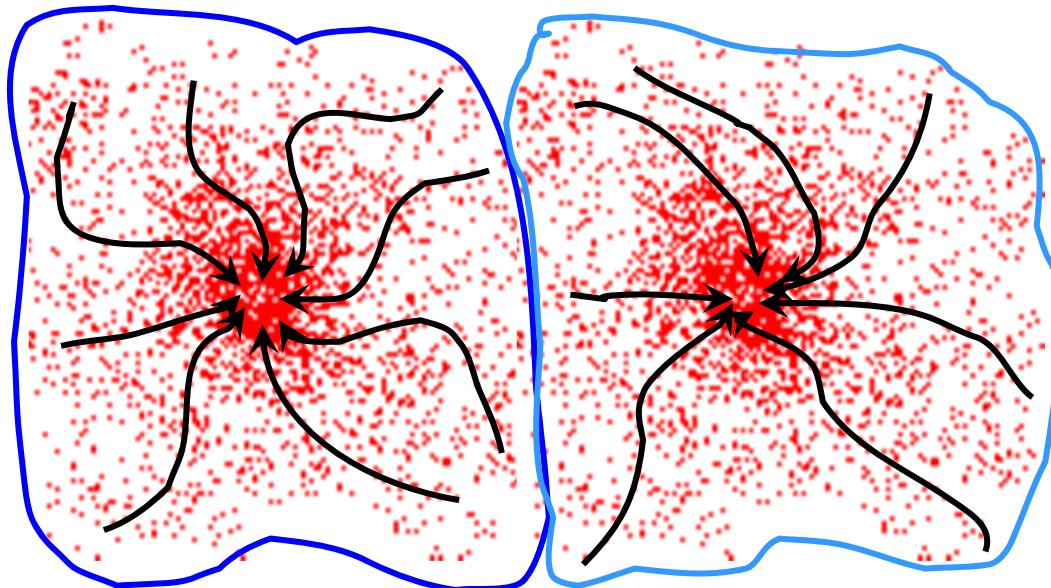
Real Modality Analysis



The blue data points were traversed by the windows towards the mode.

Mean-Shift Clustering

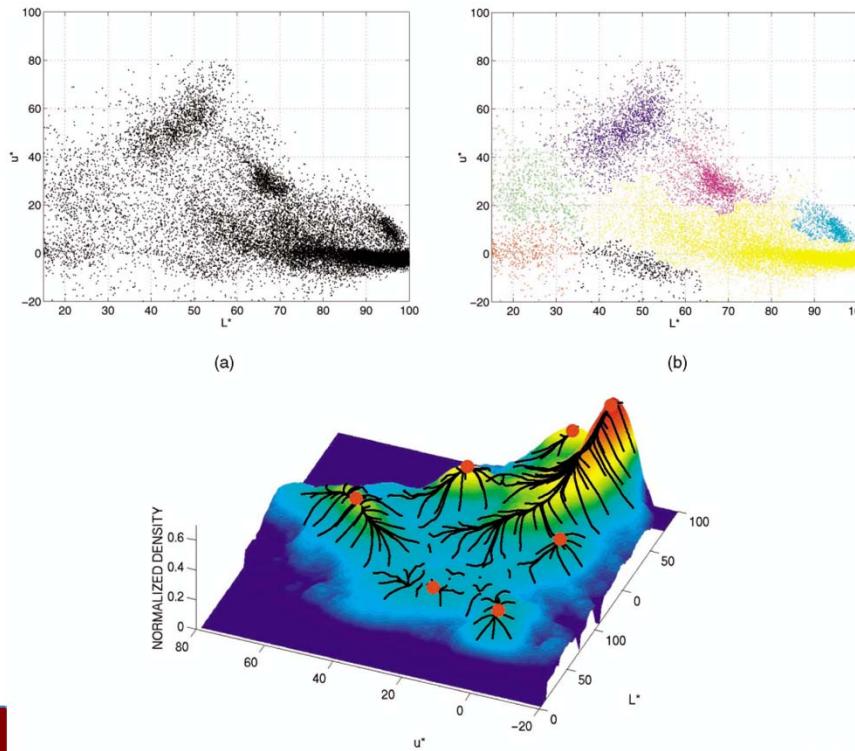
- Cluster: all data points in the attraction basin of a mode
- Attraction basin: the region for which all trajectories lead to the same mode



Slide by Y. Ukrainitz & B. Sarel

Mean-Shift Clustering/Segmentation

- Find features (color, gradients, texture, etc)
- Initialize windows at individual pixel locations
- Perform mean shift for each window until convergence
- Merge windows that end up near the same “peak” or mode

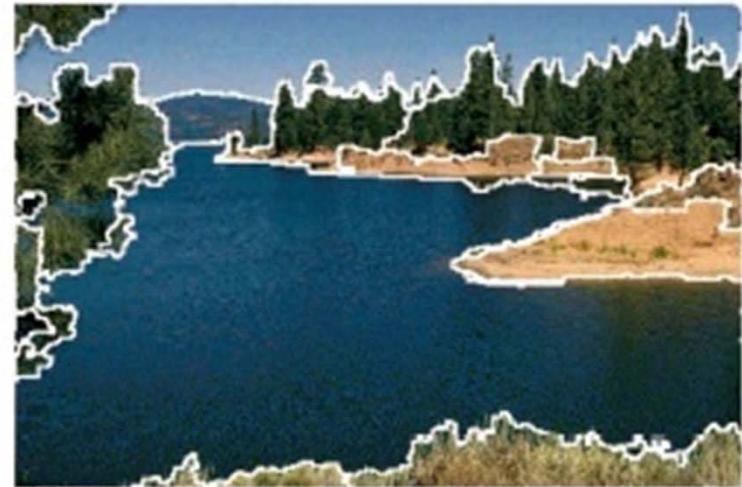


Mean-Shift Segmentation Results



<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>

More Results

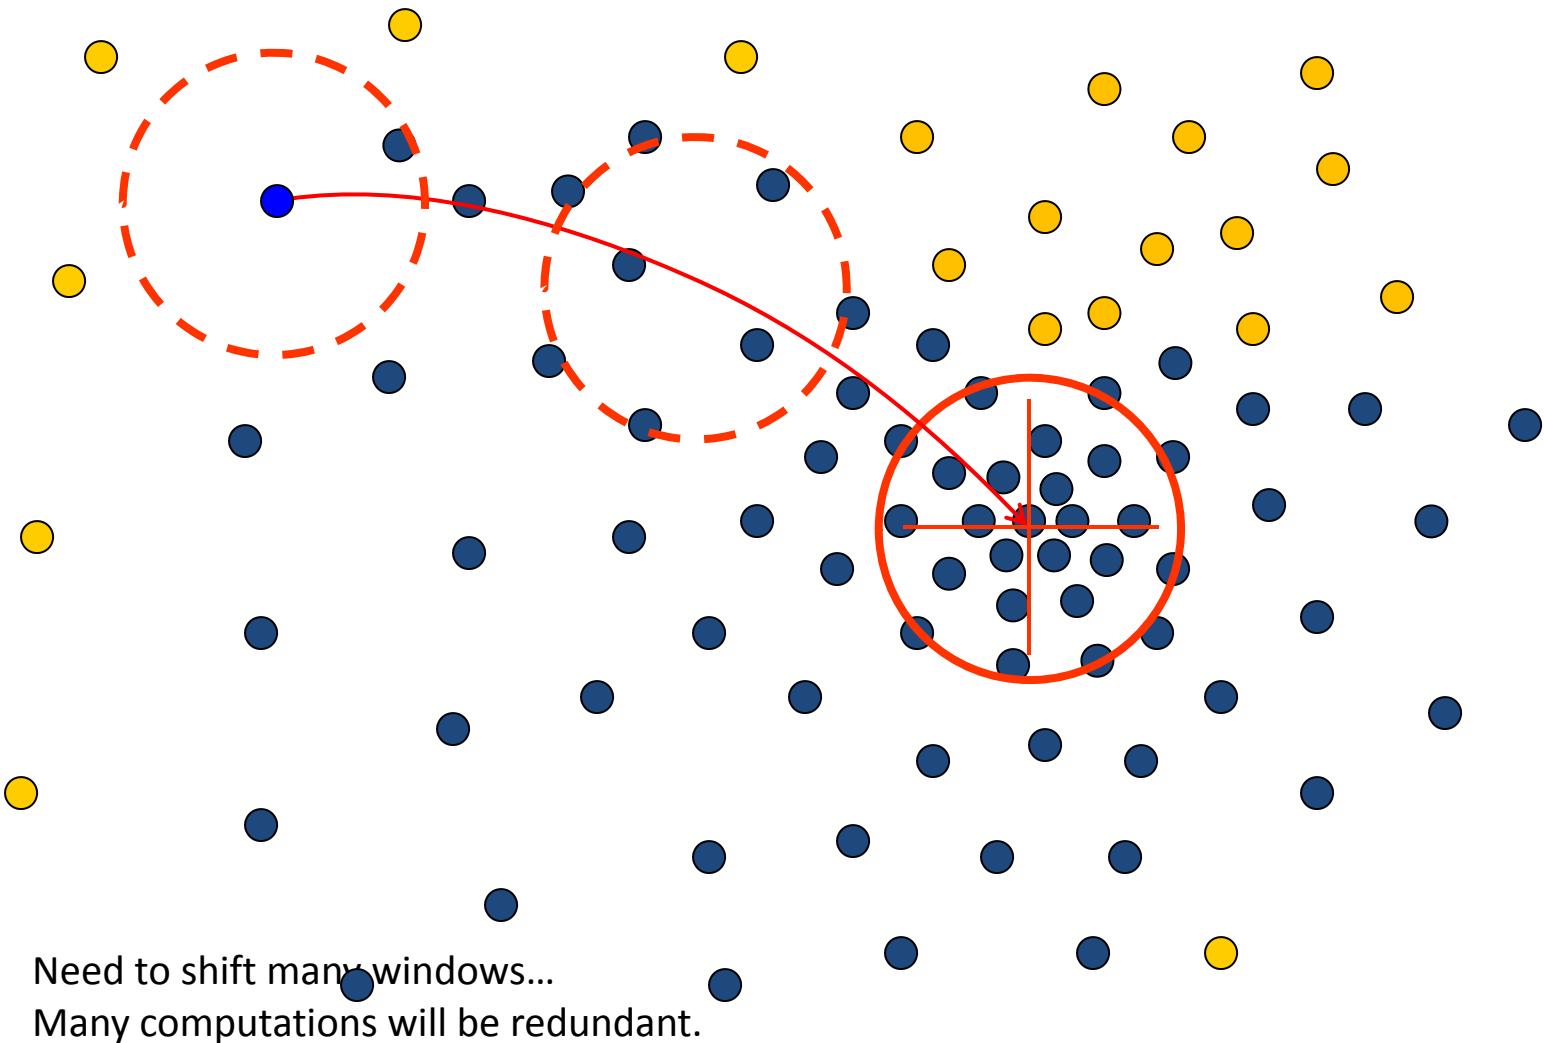


Slide credit: Svetlana Lazebnik

More Results

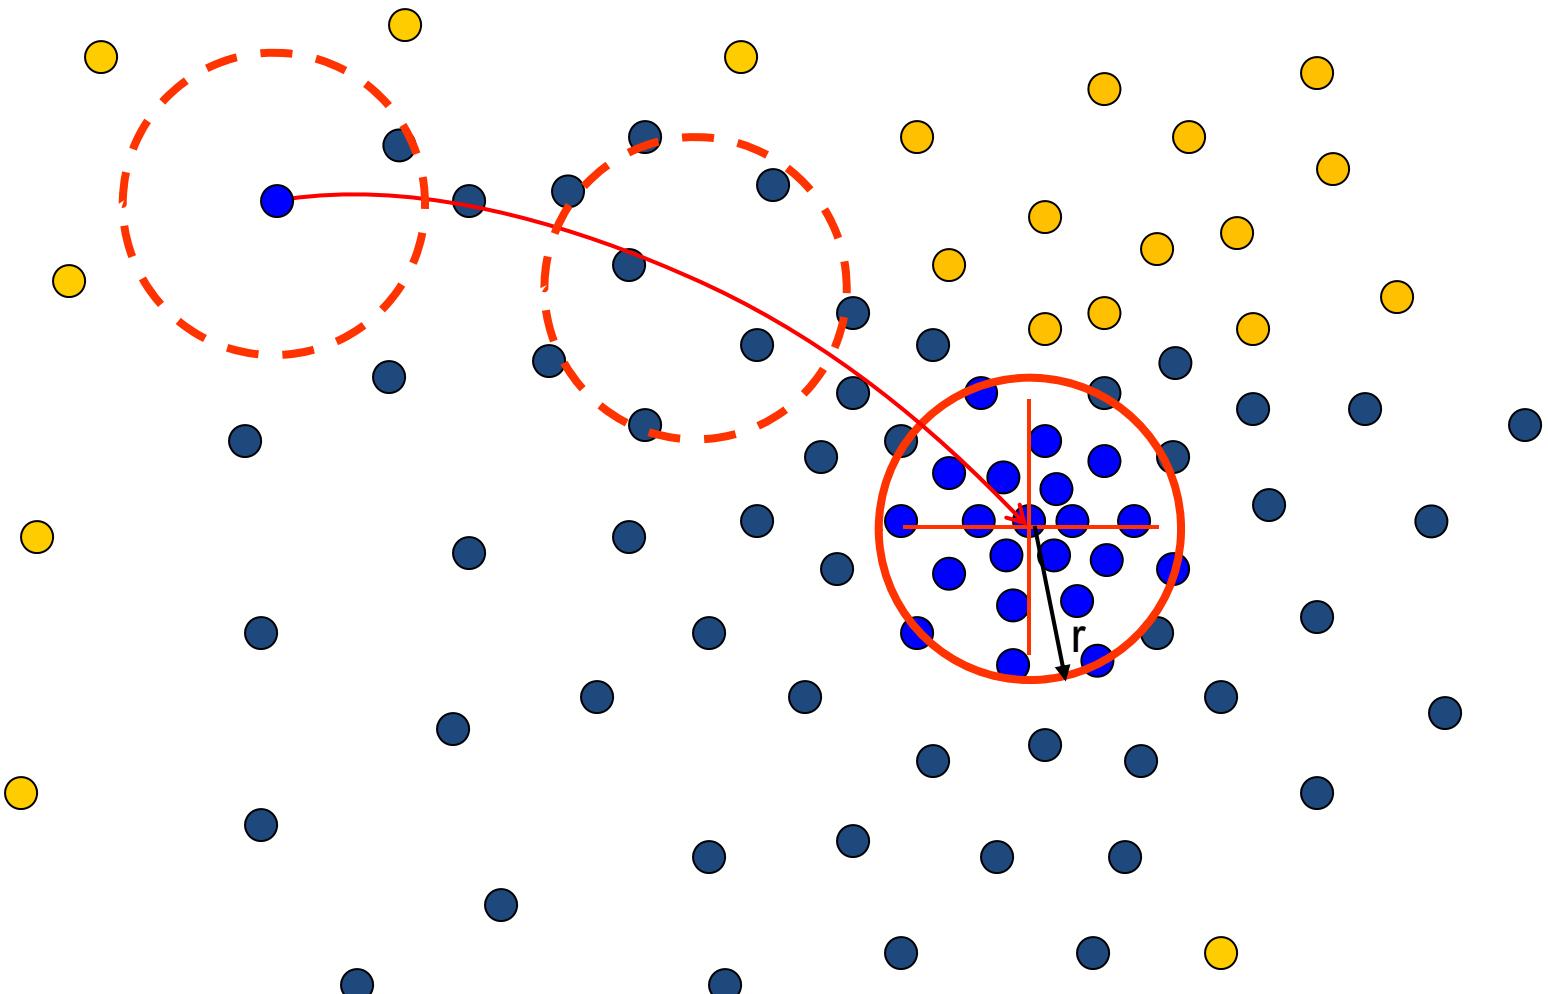


Problem: Computational Complexity



Slide credit: Bastian Leibe

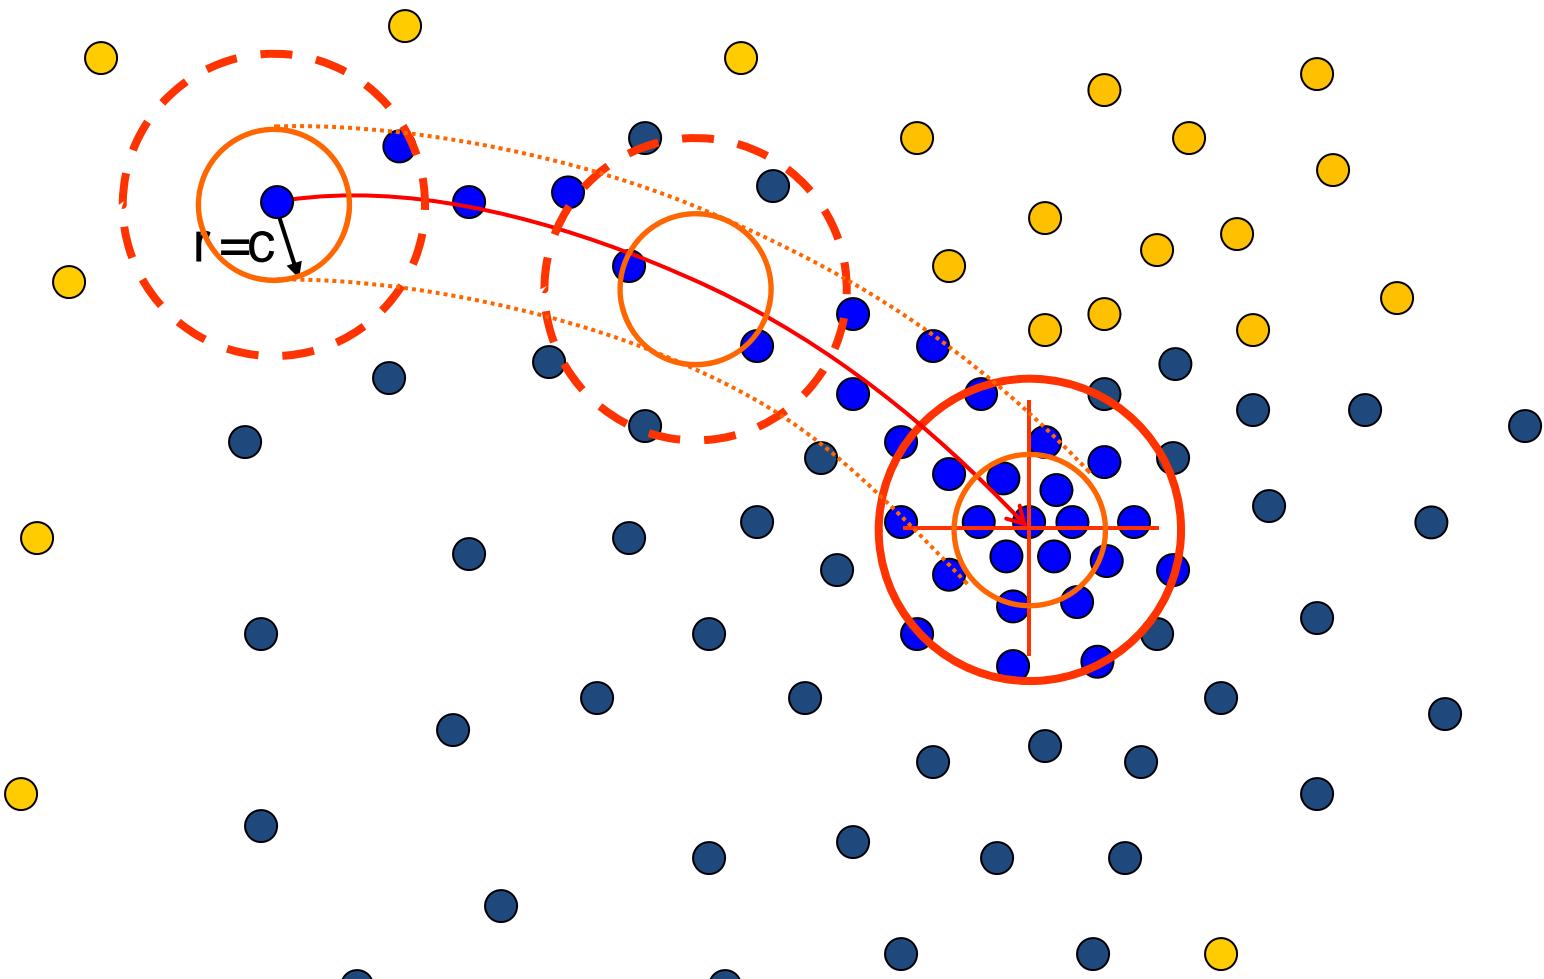
Speedups: Basin of Attraction



1. Assign all points within radius r of end point to the mode.

Slide credit: Bastian Leibe

Speedups



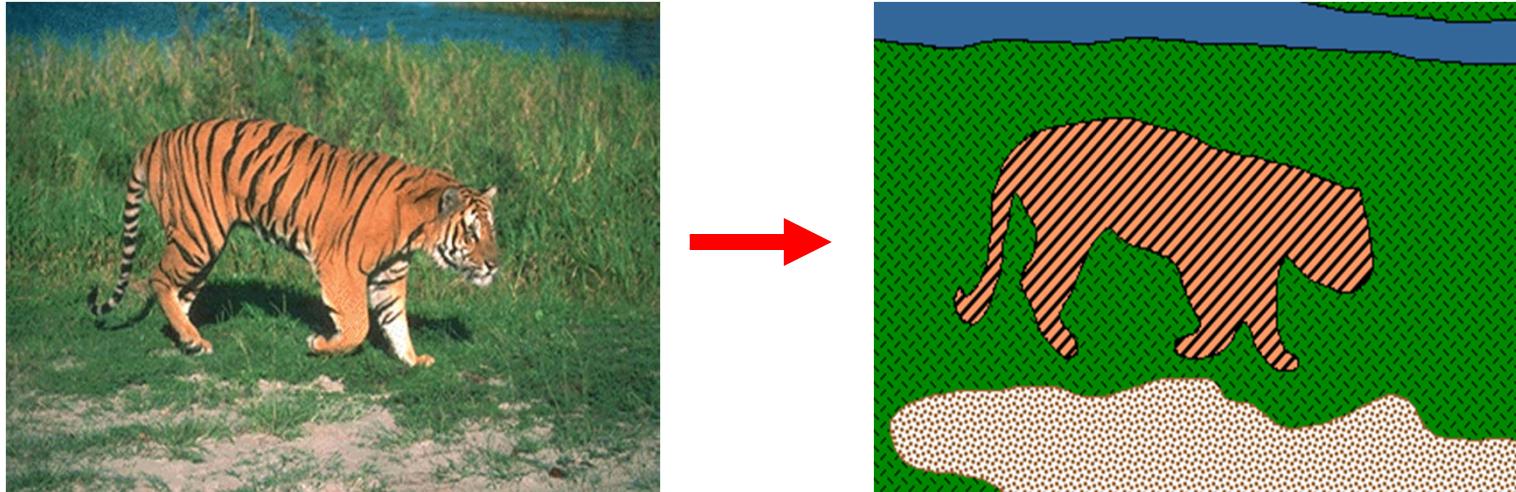
2. Assign all points within radius r/c of the search path to the mode -> reduce the number of data points to search.

Summary Mean-Shift

- Pros
 - General, application-independent tool
 - Model-free, does not assume any prior shape (spherical, elliptical, etc.) on data clusters
 - Just a single parameter (window size h)
 - h has a physical meaning (unlike k-means)
 - Finds variable number of modes
 - Robust to outliers
- Cons
 - Output depends on window size
 - Window size (bandwidth) selection is not trivial
 - Computationally (relatively) expensive ($\sim 2s/\text{image}$)
 - Does not scale well with dimension of feature space

Back to the Image Segmentation Problem...

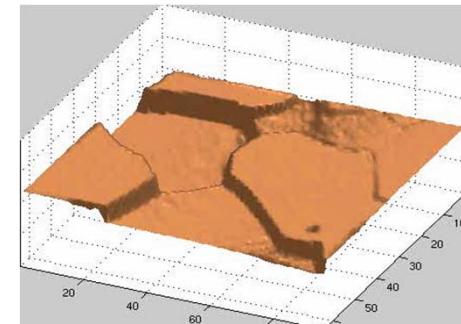
- Goal: identify groups of pixels that go together



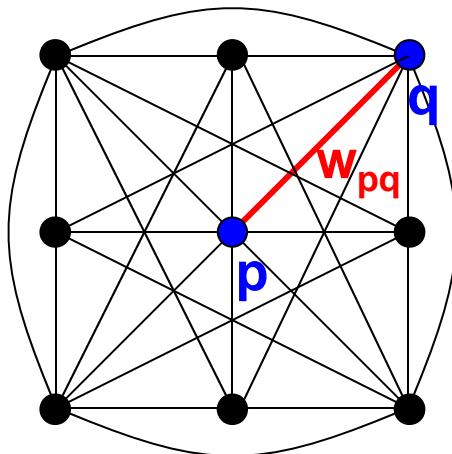
- Up to now, we have focused on ways to group pixels into image segments based on their appearance...
 - Segmentation as clustering.
- We also want to enforce region constraints.
 - Spatial consistency
 - Smooth borders

What we will learn today?

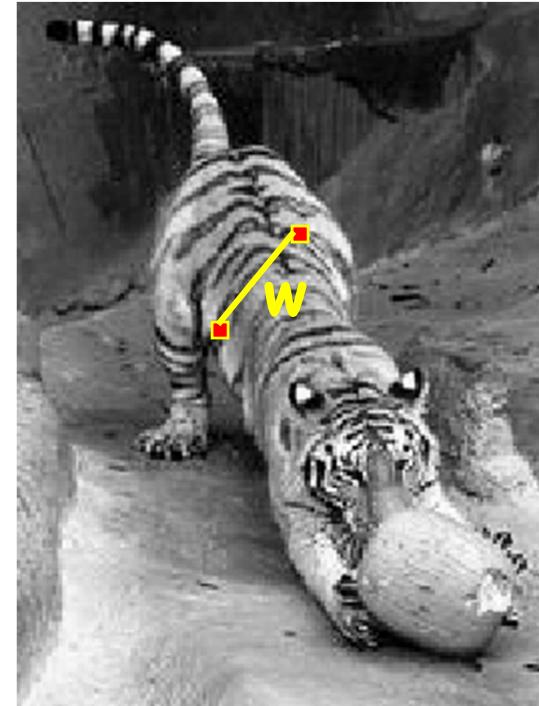
- Model free clustering
 - Mean-shift
- Graph theoretic segmentation
 - Normalized Cuts
 - Using texture features
- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation (supp. materials)
 - s-t mincut algorithm (supp. materials)
 - Extension to non-binary case (supp. materials)
 - Applications



Images as Graphs

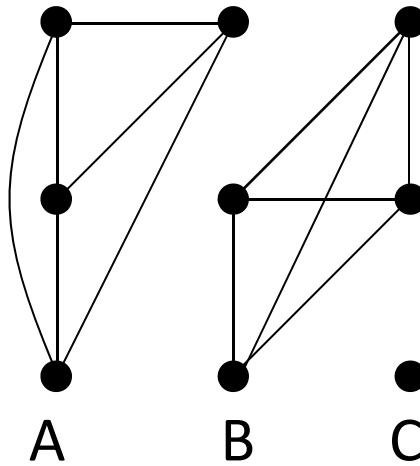


- *Fully-connected graph*
 - Node (vertex) for every pixel
 - Link between *every* pair of pixels, (p,q)
 - Affinity weight w_{pq} for each link (edge)
 - w_{pq} measures similarity
 - Similarity is *inversely proportional* to difference (in color and position...)



Slide credit: Steve Seitz

Segmentation by Graph Cuts



- Break Graph into Segments
 - Delete links that cross between segments
 - Easiest to break links that have low similarity (low weight)
 - Similar pixels should be in the same segments
 - Dissimilar pixels should be in different segments

Slide credit: Steve Seitz

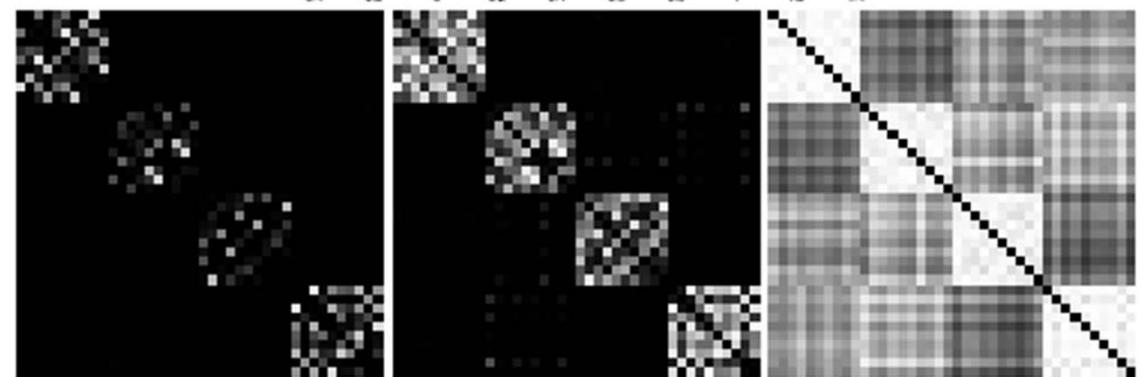
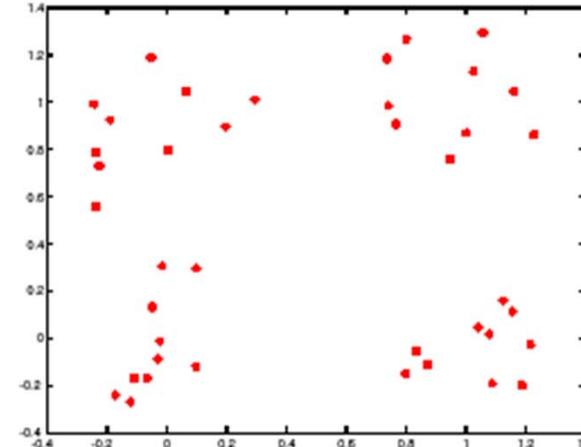
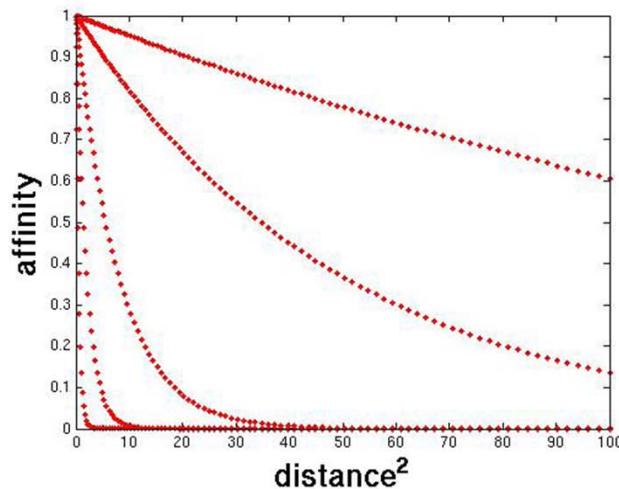
Measuring Affinity

- Distance
$$aff(x, y) = \exp\left\{-\frac{1}{2\sigma_d^2} \|x - y\|^2\right\}$$
- Intensity
$$aff(x, y) = \exp\left\{-\frac{1}{2\sigma_d^2} \|I(x) - I(y)\|^2\right\}$$
- Color
$$aff(x, y) = \exp\left\{-\frac{1}{2\sigma_d^2} \underbrace{\text{dist}\left(c(x), c(y)\right)^2}_{\text{(some suitable color space distance)}}\right\}$$
- Texture
$$aff(x, y) = \exp\left\{-\frac{1}{2\sigma_d^2} \underbrace{\|f(x) - f(y)\|^2}_{\text{(vectors of filter outputs)}}\right\}$$

Source: Forsyth & Ponce

Scale Affects Affinity

- Small σ : group only nearby points
- Large σ : group far-away points



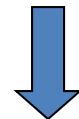
Slide credit: Svetlana Lazebnik

Graph Cut: using Eigenvalues

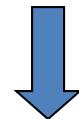
- Extract a single good cluster
 - Where elements have high affinity values with each other

$$\mathbf{w}_n^T \mathcal{A} \mathbf{w}_n$$

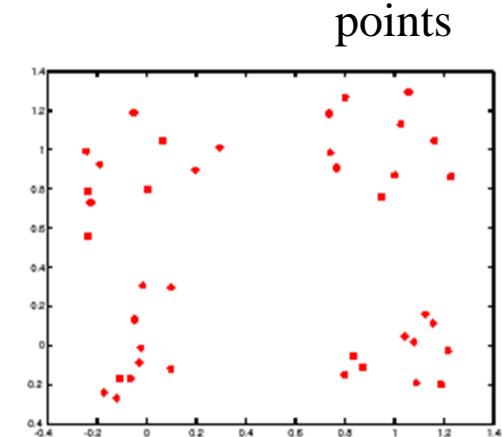
{association of element i with cluster n } \times
{affinity between i and j } \times
{association of element j with cluster n }



$$\mathbf{w}_n^T \mathcal{A} \mathbf{w}_n + \lambda (\mathbf{w}_n^T \mathbf{w}_n - 1)$$



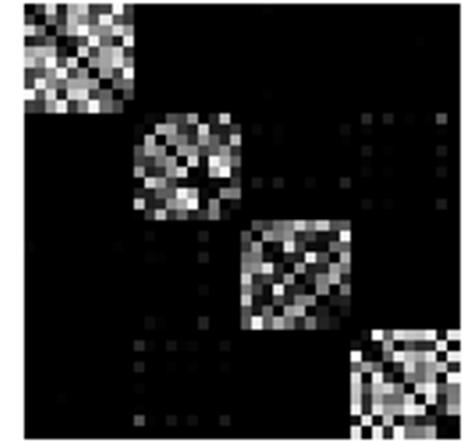
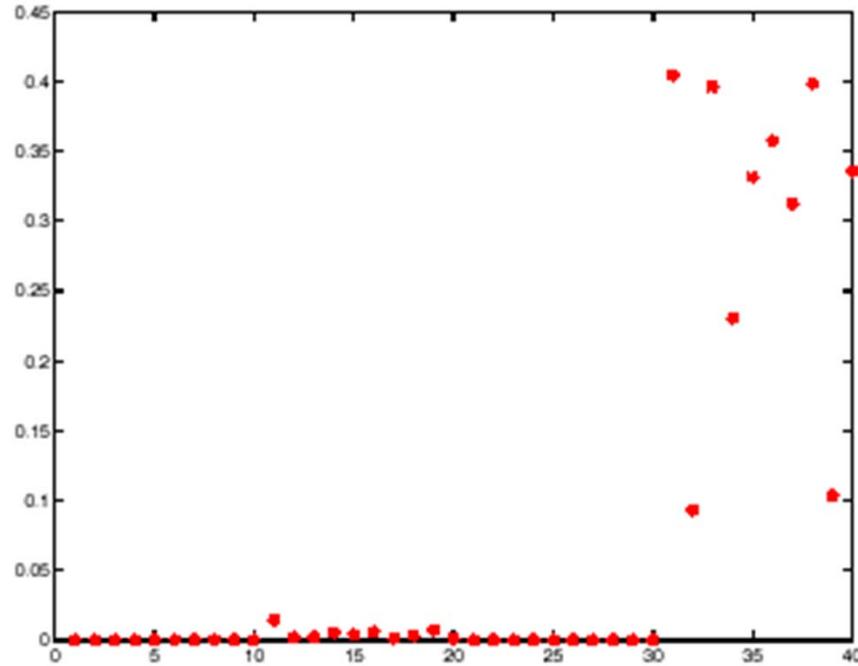
$$\mathcal{A} \mathbf{w}_n = \lambda \mathbf{w}_n$$



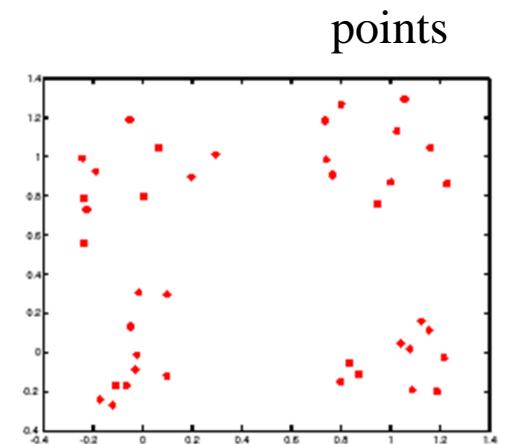
Graph Cut: using Eigenvalues

- Extract a single good cluster

Eigenvector associated w/ the largest eigenvalue



matrix

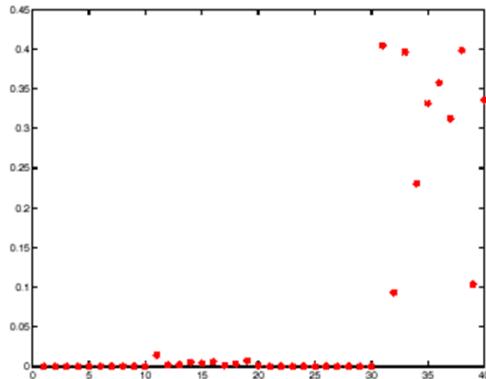


points

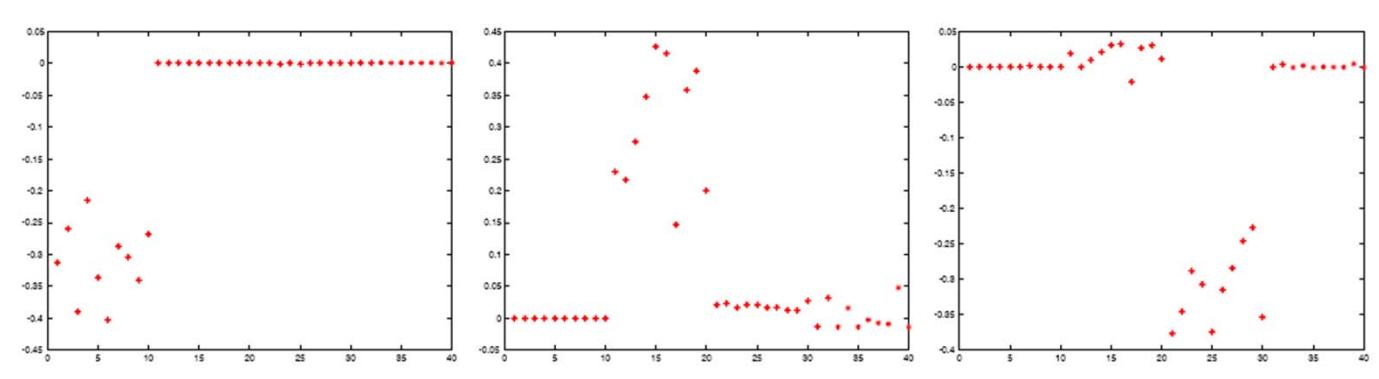
Graph Cut: using Eigenvalues

- Extract a single good cluster
- Extract weights for a set of clusters

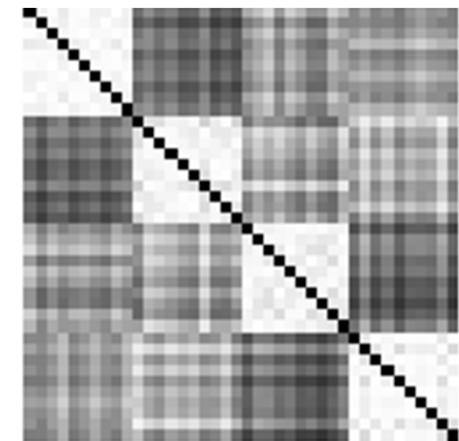
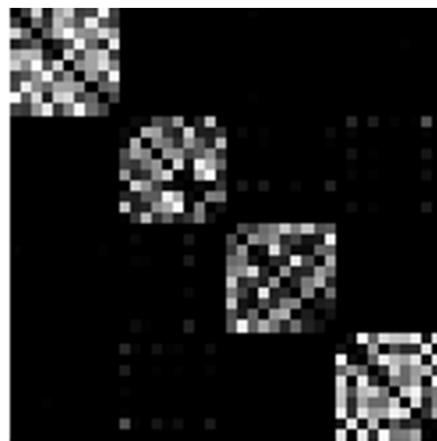
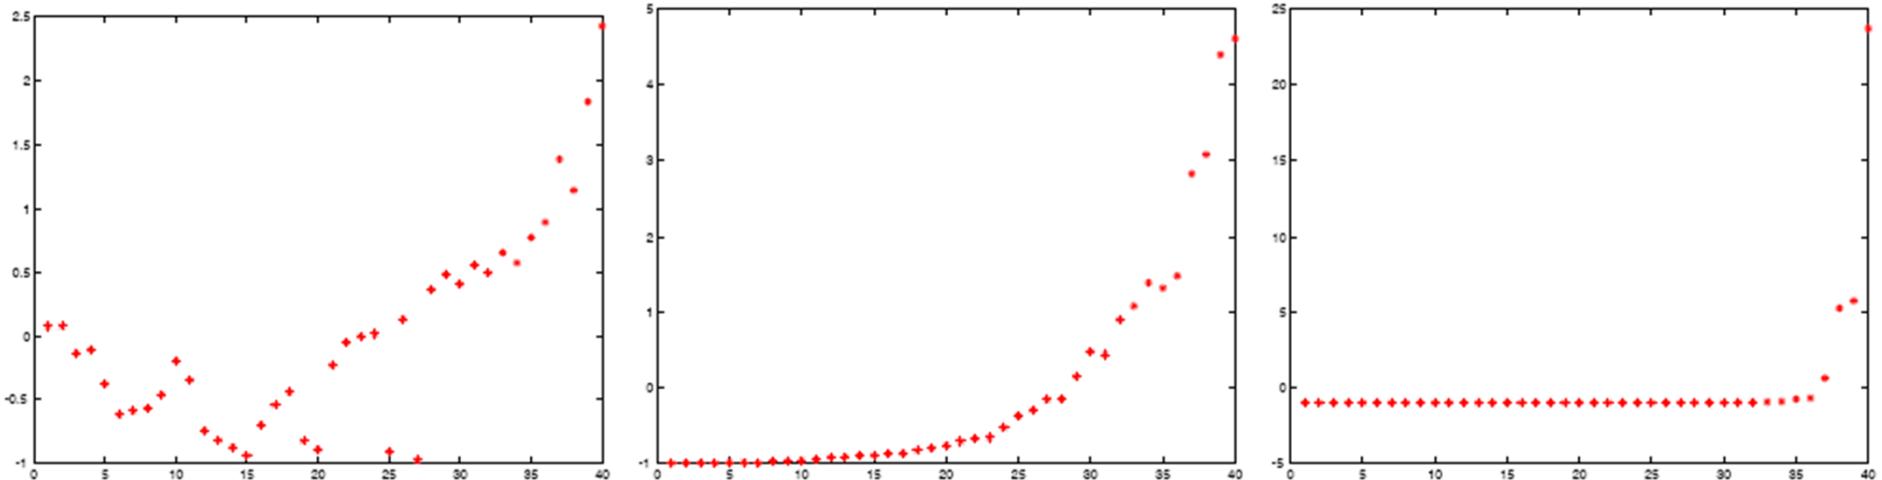
Eigenvector associated
w/ the largest eigenvalue



Eigenvectors associated with other eigenvalues



Graph Cut: using Eigenvalues (effect of the scaling factor)



Algorithm 14.6: Clustering by Graph Eigenvectors

Construct an affinity matrix

Compute the eigenvalues and eigenvectors of the affinity matrix

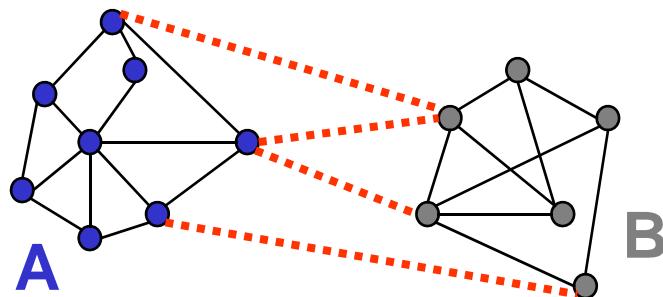
Until there are sufficient clusters

 Take the eigenvector corresponding to the
 largest unprocessed eigenvalue; zero all components corresponding
 to elements that have already been clustered, and threshold the
 remaining components to determine which element
 belongs to this cluster, choosing a threshold by
 clustering the components, or
 using a threshold fixed in advance.

 If all elements have been accounted for, there are
 sufficient clusters

end

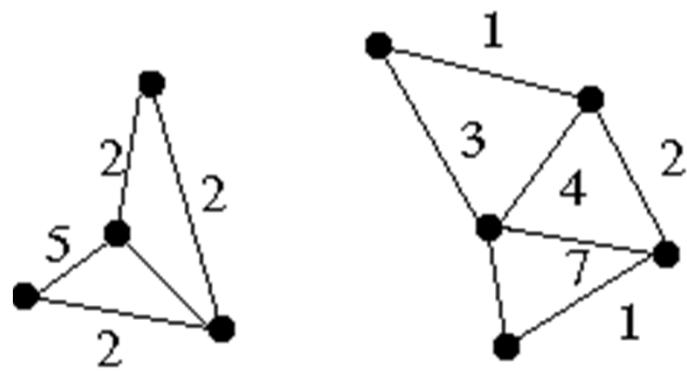
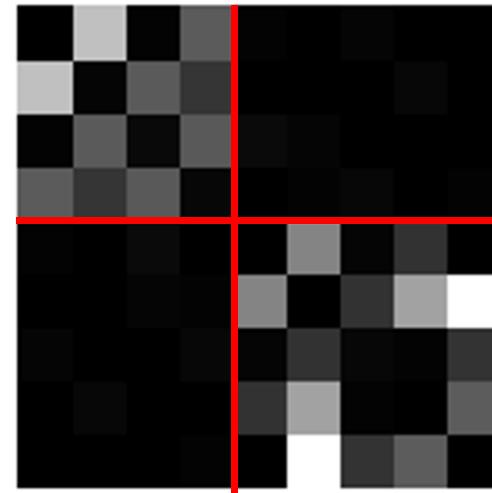
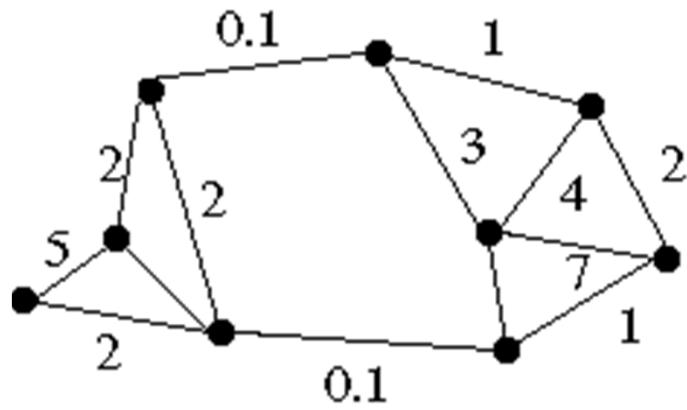
Graph Cut



- Set of edges whose removal makes a graph disconnected
- Cost of a cut
 - Sum of weights of cut edges: $cut(A, B) = \sum_{p \in A, q \in B} w_{p,q}$
- A graph cut gives us a segmentation
 - What is a “good” graph cut and how do we find one?

Slide credit: Steve Seitz

Graph Cut

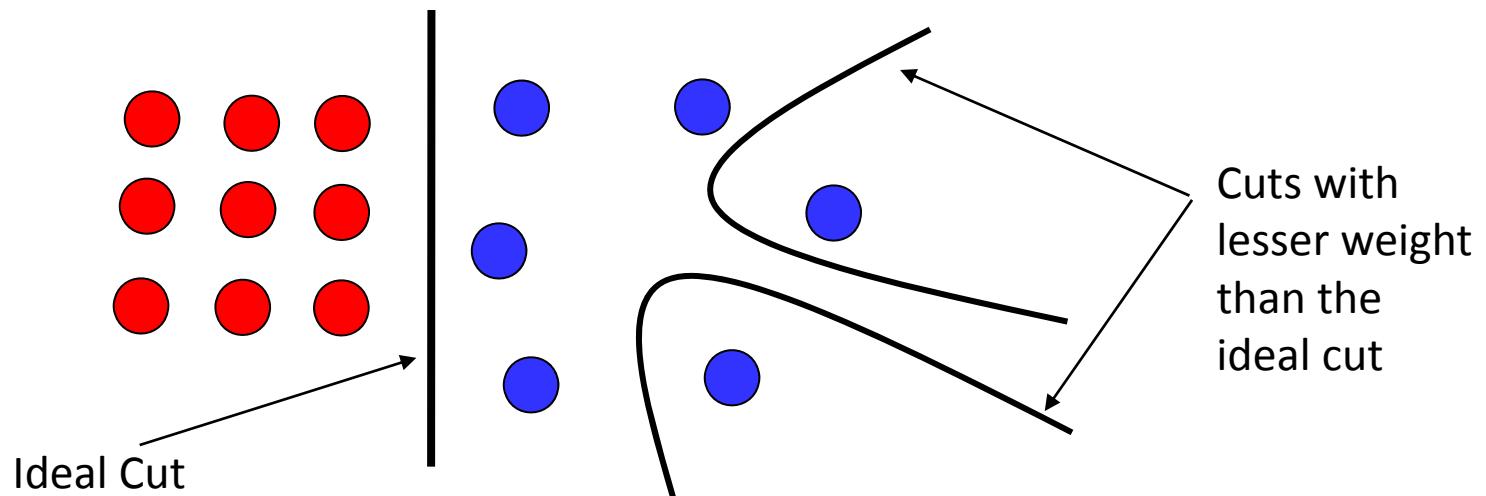


Here, the cut is nicely defined by the block-diagonal structure of the affinity matrix.

Image Source: Forsyth & Ponce

Minimum Cut

- We can do segmentation by finding the *minimum cut* in a graph
 - a **minimum cut** of a graph is a cut whose cutset has the smallest number of elements (unweighted case) or smallest sum of weights possible.
 - Efficient algorithms exist for doing this
- Drawback:
 - Weight of cut proportional to number of edges in the cut
 - Minimum cut tends to cut off very small, isolated components



Normalized Cut (NCut)

- A minimum cut penalizes large segments
- This can be fixed by normalizing for size of segments
- The **normalized cut** cost is:

$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

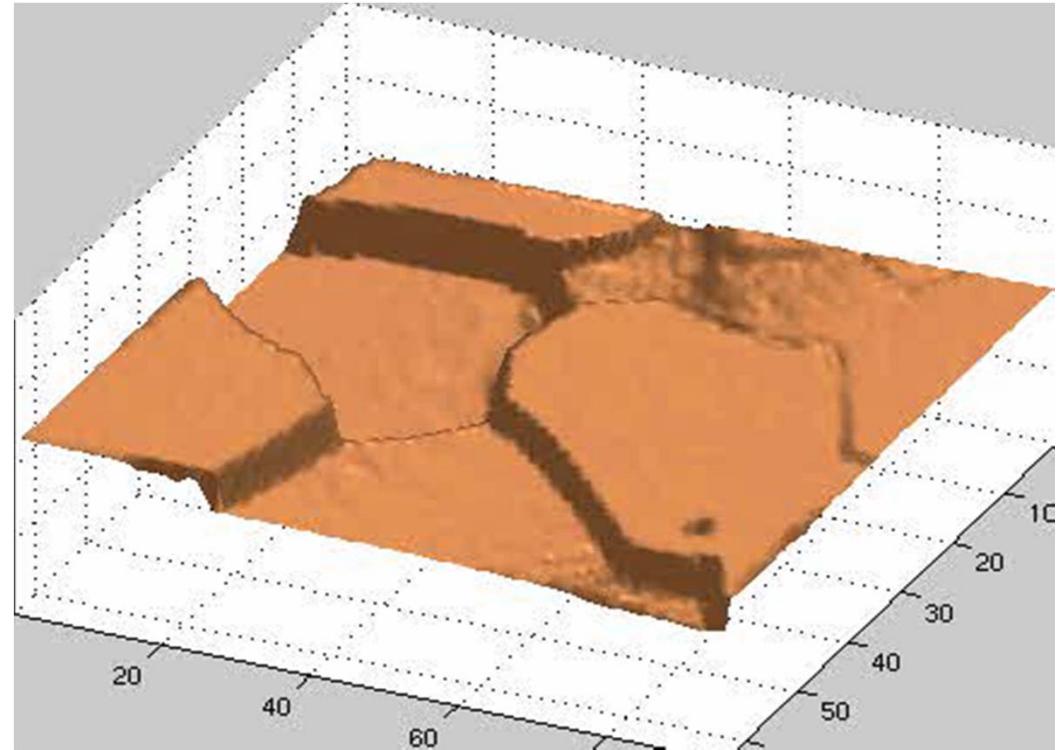
$assoc(A, V)$ = sum of weights of all edges in V that touch A

$$= cut(A, B) \left[\frac{1}{\sum_{p \in A} w_{p,q}} + \frac{1}{\sum_{q \in B} w_{p,q}} \right]$$

- The exact solution is NP-hard but an approximation can be computed by solving a *generalized eigenvalue* problem.

J. Shi and J. Malik. [Normalized cuts and image segmentation](#). PAMI 2000

Interpretation as a Dynamical System



- Treat the links as springs and shake the system
 - Elasticity proportional to cost
 - Vibration “modes” correspond to segments
 - Can compute these by solving a generalized eigenvector problem

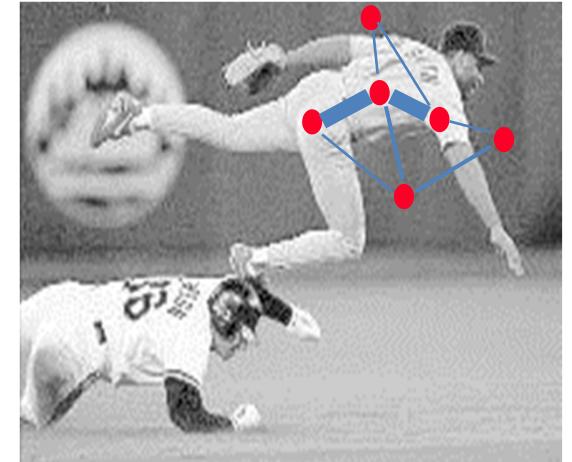
NCuts as a Generalized Eigenvector Problem

- Definitions

W : the affinity matrix, $W(i, j) = w_{i,j}$;

D : the diag. matrix, $D(i, i) = \sum_j W(i, j)$;

x : a vector in $\{1, -1\}^N$, $x(i) = 1 \Leftrightarrow i \in A$.



- Rewriting Normalized Cut in matrix form:

$$\begin{aligned} NCut(A, B) &= \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)} \\ &= \frac{(1+x)^T (D-W)(1+x)}{k l^T D l} + \frac{(1-x)^T (D-W)(1-x)}{(1-k) l^T D l}; \quad k = \frac{\sum_{x_i > 0} D(i, i)}{\sum_i D(i, i)} \\ &= \dots \end{aligned}$$

Slide credit: Jitendra Malik

Some More Math...

We see again this is an unbiased measure, which reflects how tightly on average nodes within the group are connected to each other.

Another important property of this definition of association and disassociation of a partition is that they are naturally related:

$$\begin{aligned} Ncut(A, B) &= \frac{cut(A, B)}{\alpha\alpha\alpha(A, V)} + \frac{cut(A, B)}{\alpha\alpha\alpha(B, V)} \\ &= \frac{\alpha\alpha\alpha(A, V) - \alpha\alpha\alpha(A, A)}{\alpha\alpha\alpha(A, V)} \\ &\quad + \frac{\alpha\alpha\alpha(B, V) - \alpha\alpha\alpha(B, B)}{\alpha\alpha\alpha(B, V)} \\ &= 2 - \left(\frac{\alpha\alpha\alpha(A, A)}{\alpha\alpha\alpha(A, V)} + \frac{\alpha\alpha\alpha(B, B)}{\alpha\alpha\alpha(B, V)} \right) \\ &= 2 - Nassoc(A, B) \end{aligned}$$

Hence the two partition criteria that we seek in our grouping algorithm, minimizing the disassociation between the groups and maximizing the association within the group, are in fact identical, and can be satisfied simultaneously. In our algorithm, we will use this normalized one as the partition criterion.

Having defined the graph partition criterion that we want to optimize, we will show how such an optimal partition can be computed efficiently.

2.1 Computing the optimal partition

Given a partition of nodes of a graph, V , into two sets A and B , let \mathbf{z} be an $N = |V|$ dimensional indicator vector, $z_i = 1$ if node i is in A , and -1 otherwise. Let $d(i) = \sum_j w(i, j)$, be the total connection from node i to all other nodes. With the definitions \mathbf{a} and \mathbf{d} we can rewrite $Ncut(A, B)$ as:

$$\begin{aligned} Ncut(A, B) &= \frac{cut(A, B)}{\alpha\alpha\alpha(A, V)} + \frac{cut(B, A)}{\alpha\alpha\alpha(B, V)} \\ &= \frac{\sum_{(a_i>0, b_j<0)} -w_{ij}\mathbf{a}_i \cdot \mathbf{a}_j}{\sum_{a_i>0} d_i} \\ &\quad + \frac{\sum_{(a_i<0, b_j>0)} -w_{ij}\mathbf{a}_i \cdot \mathbf{a}_j}{\sum_{a_i<0} d_i} \end{aligned}$$

Let \mathbf{D} be an $N \times N$ diagonal matrix with d on its diagonal, \mathbf{W} be an $N \times N$ symmetrical matrix with $w(i, j) = w_{ij}$, $k = \frac{\sum_{i>0} d_i}{\sum_i d_i}$, and \mathbf{x} be an $N \times 1$ vector of all ones. Using the fact $\frac{1+\mathbf{a}}{2}$ and $\frac{1-\mathbf{a}}{2}$ are indicator vectors for $a_i > 0$ and $a_i < 0$ respectively, we can rewrite $\frac{1}{4}[Ncut(\mathbf{a})]$ as:

$$\begin{aligned} &= \frac{(\mathbf{1}+\mathbf{a})^T(\mathbf{D}-\mathbf{W})(\mathbf{1}+\mathbf{a})}{k\mathbf{x}^T\mathbf{D}\mathbf{x}} + \frac{(\mathbf{1}-\mathbf{a})^T(\mathbf{D}-\mathbf{W})(\mathbf{1}-\mathbf{a})}{(1-k)\mathbf{x}^T\mathbf{D}\mathbf{x}} \\ &= \frac{(\mathbf{a}^T(\mathbf{D}-\mathbf{W})\mathbf{a} + \mathbf{x}^T(\mathbf{D}-\mathbf{W})\mathbf{x})}{k(1-k)\mathbf{x}^T\mathbf{D}\mathbf{x}} + \frac{\mathbf{x}^T(\mathbf{D}-\mathbf{W})\mathbf{a} - \mathbf{a}^T(\mathbf{D}-\mathbf{W})\mathbf{x}}{k(1-k)\mathbf{x}^T\mathbf{D}\mathbf{x}} \end{aligned}$$

Let $\alpha(\mathbf{a}) = \mathbf{a}^T(\mathbf{D}-\mathbf{W})\mathbf{a}$, $\beta(\mathbf{a}) = \mathbf{x}^T(\mathbf{D}-\mathbf{W})\mathbf{a}$, $\gamma = \mathbf{x}^T(\mathbf{D}-\mathbf{W})\mathbf{x}$, and $M = \mathbf{x}^T\mathbf{D}\mathbf{x}$, we can then further expand the above equation as:

$$\begin{aligned} &= \frac{(\alpha(\mathbf{a}) + \gamma) + 2(1-2k)\beta(\mathbf{a})}{k(1-k)M} \\ &= \frac{(\alpha(\mathbf{a}) + \gamma) + 2(1-2k)\beta(\mathbf{a})}{k(1-k)M} - \frac{2(\alpha(\mathbf{a}) + \gamma)}{M} \\ &\quad + \frac{2\alpha(\mathbf{a})}{M} + \frac{2\gamma}{M} \end{aligned}$$

Dropping the last constant term, which in this case equals 0, we get

$$\begin{aligned} &= \frac{(1-2k+2k^2)(\alpha(\mathbf{a}) + \gamma) + 2(1-2k)\beta(\mathbf{a})}{k(1-k)M} + \frac{2\alpha(\mathbf{a})}{M} \\ &= \frac{(1-2k+2k^2)}{k(1-k)^2} (\alpha(\mathbf{a}) + \gamma) + \frac{2(1-2k)}{(1-k)^2} \beta(\mathbf{a}) + \frac{2\alpha(\mathbf{a})}{M} \end{aligned}$$

Letting $b = \frac{k}{1-k}$, and since $\gamma = 0$, it becomes,

$$\begin{aligned} &= \frac{(1+b^2)(\alpha(\mathbf{a}) + \gamma) + 2(1-b^2)\beta(\mathbf{a})}{bM} + \frac{2b\alpha(\mathbf{a})}{bM} \\ &= \frac{(1+b^2)(\alpha(\mathbf{a}) + \gamma)}{bM} + \frac{2(1-b^2)\beta(\mathbf{a})}{bM} + \frac{2b\alpha(\mathbf{a})}{bM} - 2b\gamma \\ &= \frac{(1+b^2)(\mathbf{a}^T(\mathbf{D}-\mathbf{W})\mathbf{a} + \mathbf{x}^T(\mathbf{D}-\mathbf{W})\mathbf{x})}{b\mathbf{x}^T\mathbf{D}\mathbf{x}} \\ &\quad + \frac{2(1-b^2)\mathbf{x}^T(\mathbf{D}-\mathbf{W})\mathbf{a}}{b\mathbf{x}^T\mathbf{D}\mathbf{x}} \\ &\quad + \frac{2b\mathbf{a}^T(\mathbf{D}-\mathbf{W})\mathbf{a} - 2b\mathbf{x}^T(\mathbf{D}-\mathbf{W})\mathbf{x}}{b\mathbf{x}^T\mathbf{D}\mathbf{x}} \\ &= \frac{(\mathbf{1}+\mathbf{a})^T(\mathbf{D}-\mathbf{W})(\mathbf{1}+\mathbf{a})}{b\mathbf{x}^T\mathbf{D}\mathbf{x}} \\ &\quad + \frac{b^2(\mathbf{1}-\mathbf{a})^T(\mathbf{D}-\mathbf{W})(\mathbf{1}-\mathbf{a})}{b\mathbf{x}^T\mathbf{D}\mathbf{x}} \\ &\quad - \frac{2b(\mathbf{1}-\mathbf{a})^T(\mathbf{D}-\mathbf{W})(\mathbf{1}+\mathbf{a})}{b\mathbf{x}^T\mathbf{D}\mathbf{x}} \\ &= \frac{[(\mathbf{1}+\mathbf{a}) - b(\mathbf{1}-\mathbf{a})]^T(\mathbf{D}-\mathbf{W})[(\mathbf{1}+\mathbf{a}) - b(\mathbf{1}-\mathbf{a})]}{b\mathbf{x}^T\mathbf{D}\mathbf{x}} \end{aligned}$$

Setting $\mathbf{y} = (\mathbf{1}+\mathbf{a}) - b(\mathbf{1}-\mathbf{a})$, it is easy to see that

$$\mathbf{y}^T \mathbf{D} \mathbf{x} = \sum_{i>0} d_i - b \sum_{i<0} d_i = 0 \quad (1)$$

since $b = \frac{k}{1-k} = \frac{\sum_{i>0} d_i}{\sum_{i<0} d_i}$, and

$$\begin{aligned} \mathbf{y}^T \mathbf{D} \mathbf{y} &= \sum_{i>0} d_i + b^2 \sum_{i<0} d_i \\ &= b \sum_{i<0} d_i + b^2 \sum_{i<0} d_i \\ &= b(\sum_{i<0} d_i + b \sum_{i<0} d_i) \\ &= b\mathbf{x}^T \mathbf{D} \mathbf{x}. \end{aligned}$$

NCuts as a Generalized Eigenvalue Problem

- After simplification, we get

$$NCut(A, B) = \frac{y^T(D-W)y}{y^T Dy}, \quad \text{with } y_i \in \{1, -b\}, \quad y^T D 1 = 0.$$

- This is a Rayleigh Quotient
 - Solution given by the “generalized” eigenvalue problem

$$(D - W)y = \lambda Dy$$

- Solved by converting to standard eigenvalue problem
- Optimal solution is second smallest eigenvector
- Gives continuous result—must convert into discrete values of y

This is hard,
as y is discrete!

Relaxation:
continuous y .

Slide credit: Alyosha Efros

NCuts Example

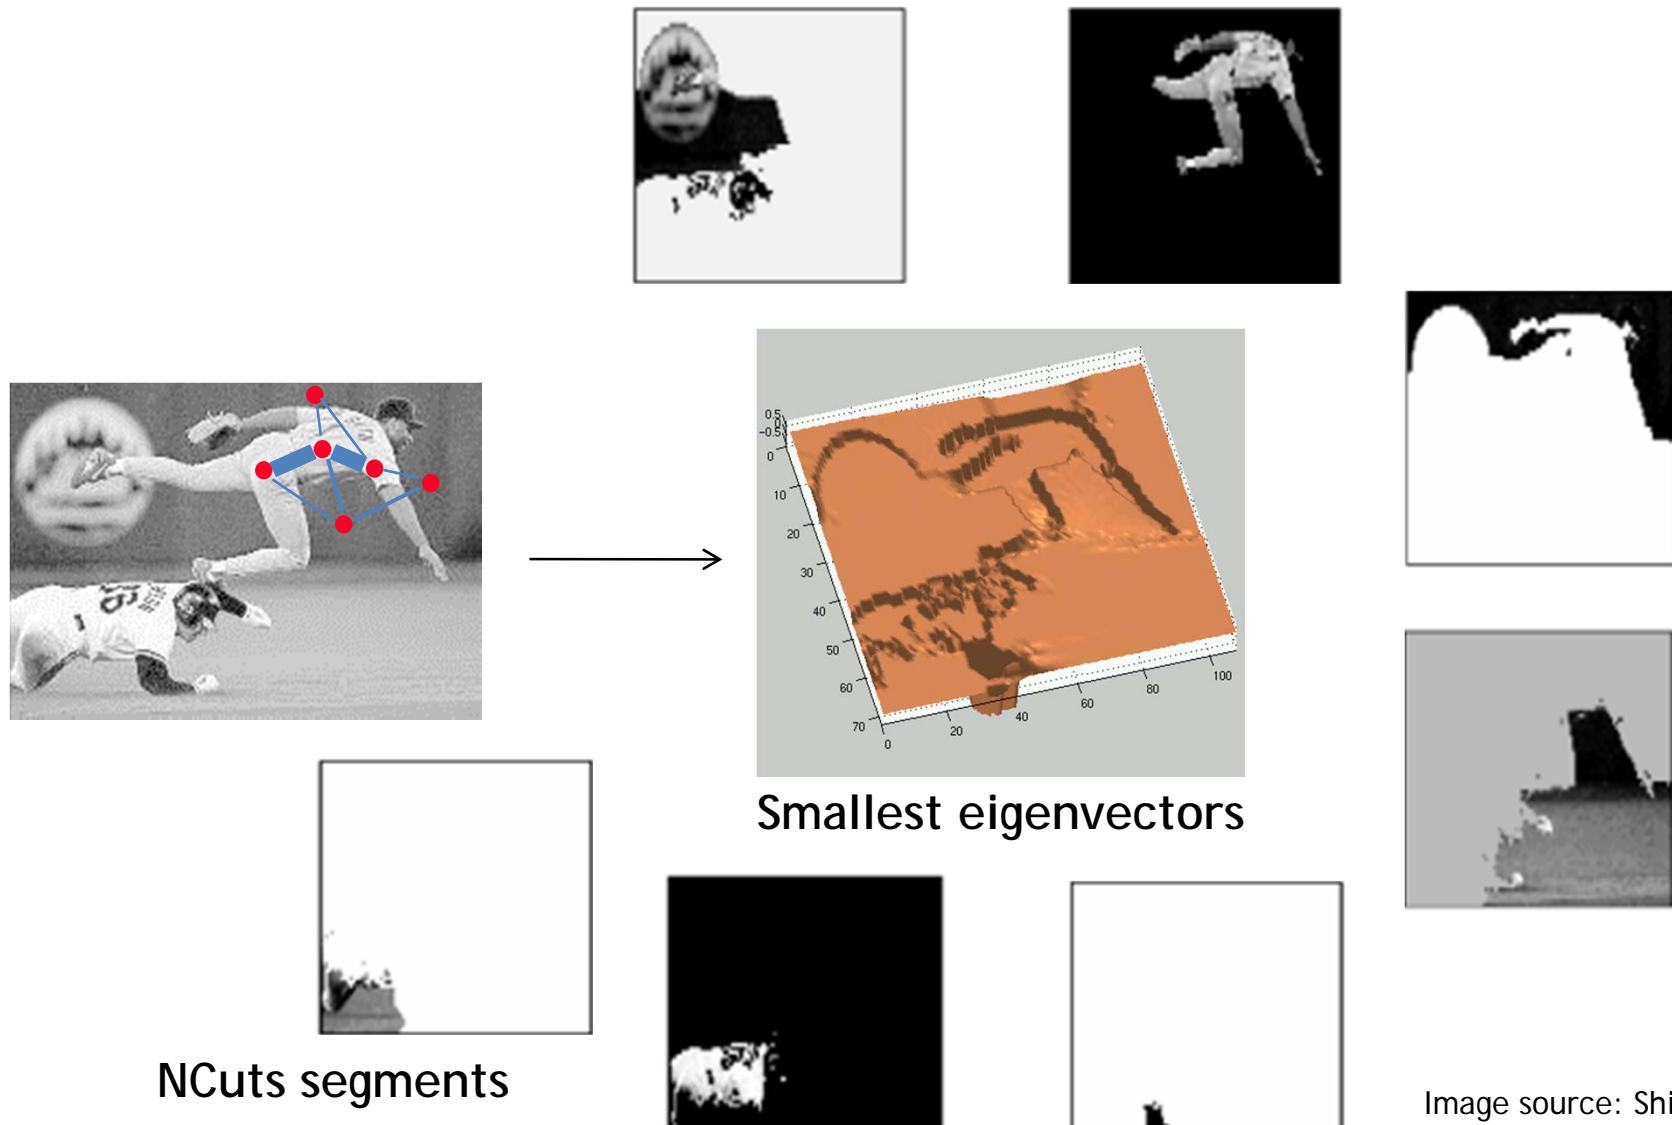
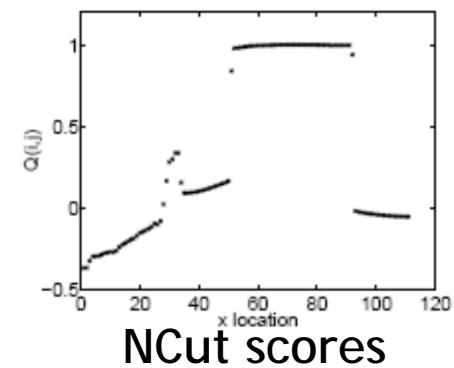
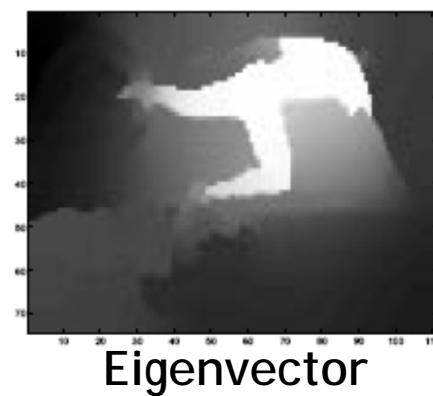
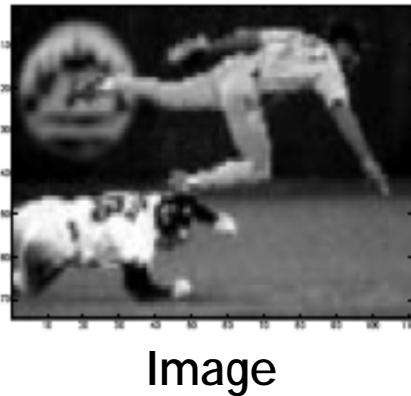


Image source: Shi & Malik

Discretization

- Problem: eigenvectors take on continuous values
 - How to choose the splitting point to binarize the image?



- Possible procedures
 - Pick a constant value (0, or 0.5).
 - Pick the median value as splitting point.
 - Look for the splitting point that has the minimum $NCut$ value:
 - Choose n possible splitting points.
 - Compute $NCut$ value.
 - Pick minimum.

NCuts: Overall Procedure

1. Construct a weighted graph $G=(V,E)$ from an image.
2. Connect each pair of pixels, and assign graph edge weights,
 $W(i,j) = \text{Prob. that } i \text{ and } j \text{ belong to the same region.}$
3. Solve $(D-W)y = \lambda Dy$ for the smallest few eigenvectors. This yields a continuous solution.
4. Threshold eigenvectors to get a discrete cut
 - This is where the approximation is made (we're not solving NP).
5. Recursively subdivide if NCut value is below a pre-specified value.

NCuts Matlab code available at

<http://www.cis.upenn.edu/~jshi/software/>

Color Image Segmentation with NCuts

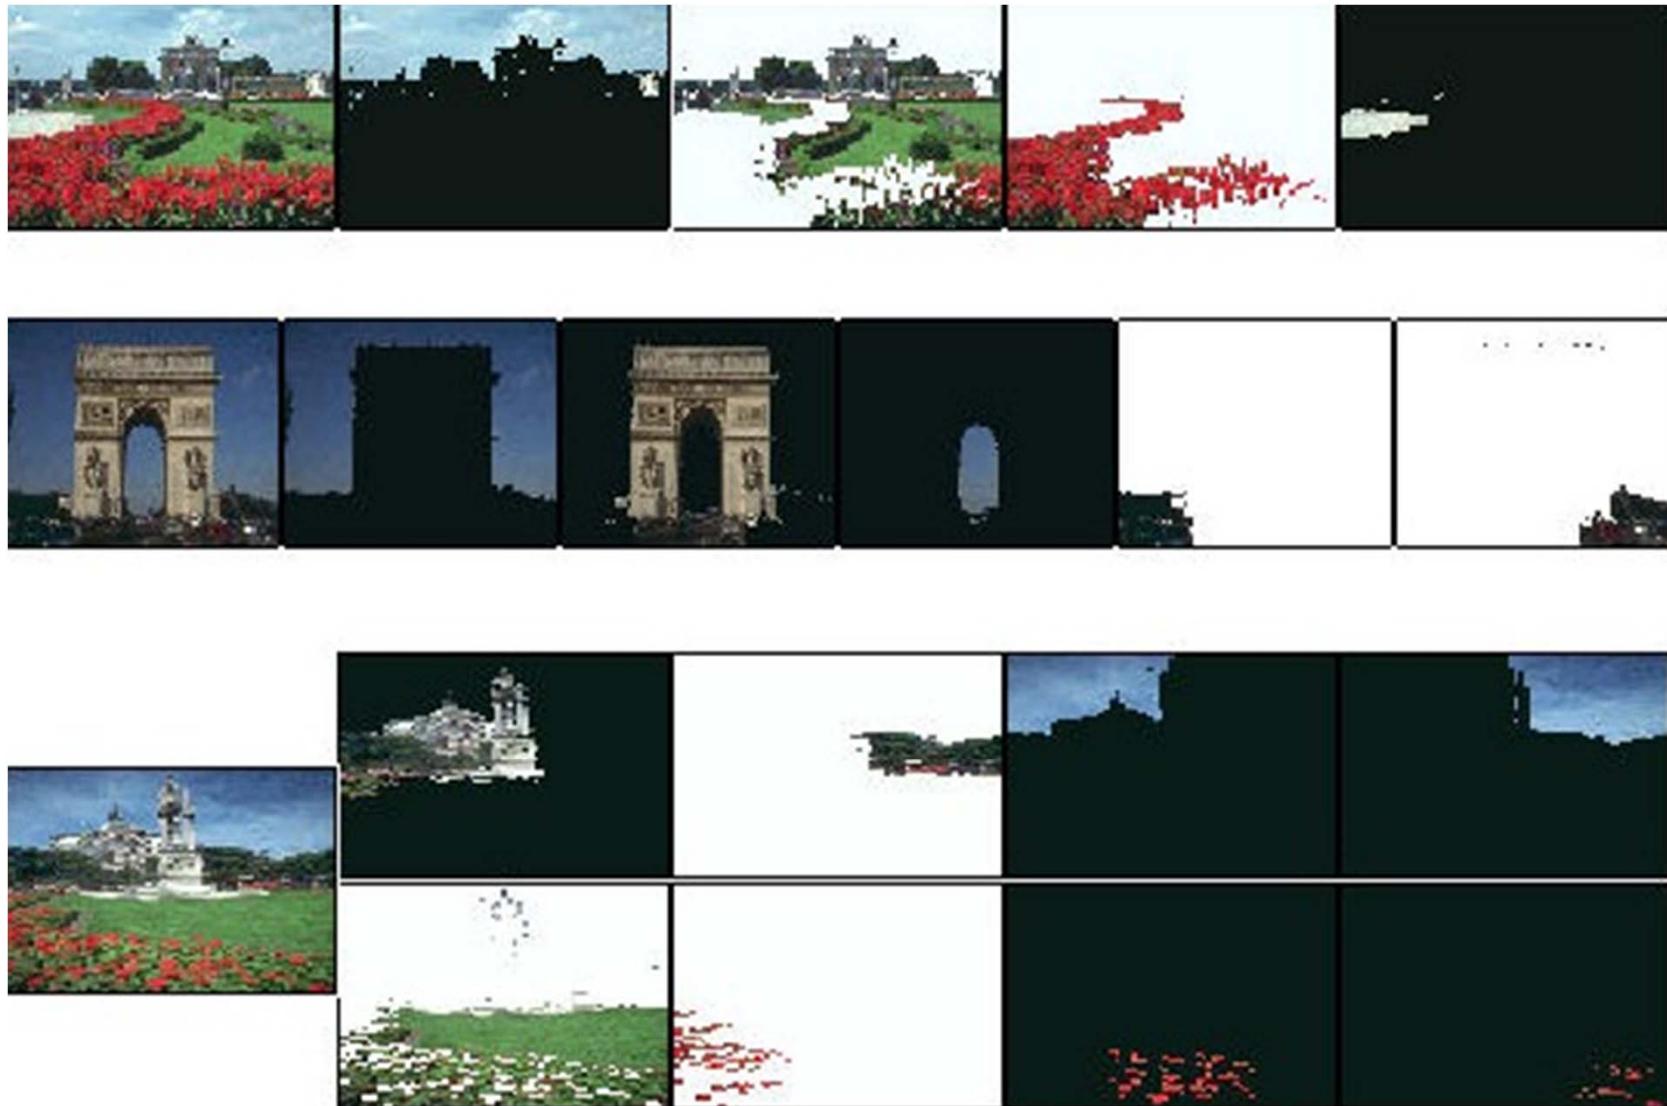
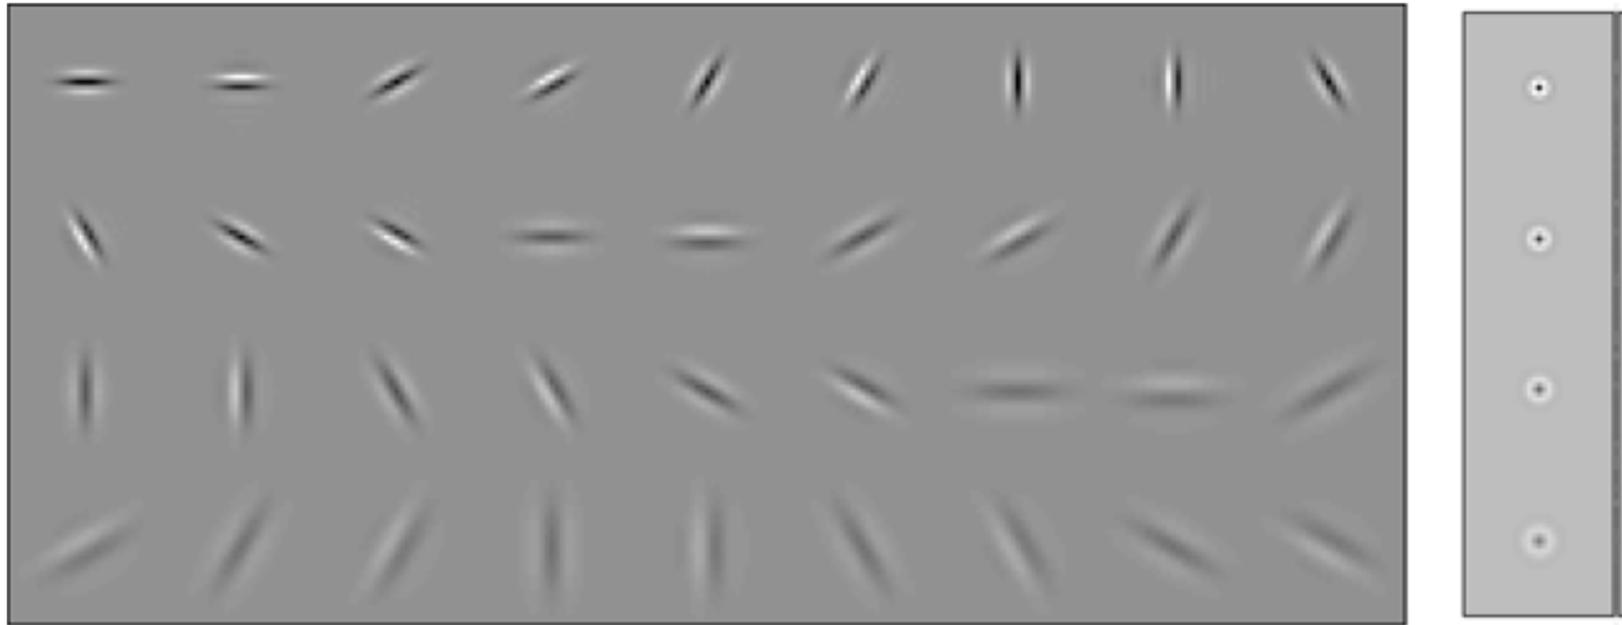


Image Source: Shi & Malik

Using Texture Features for Segmentation

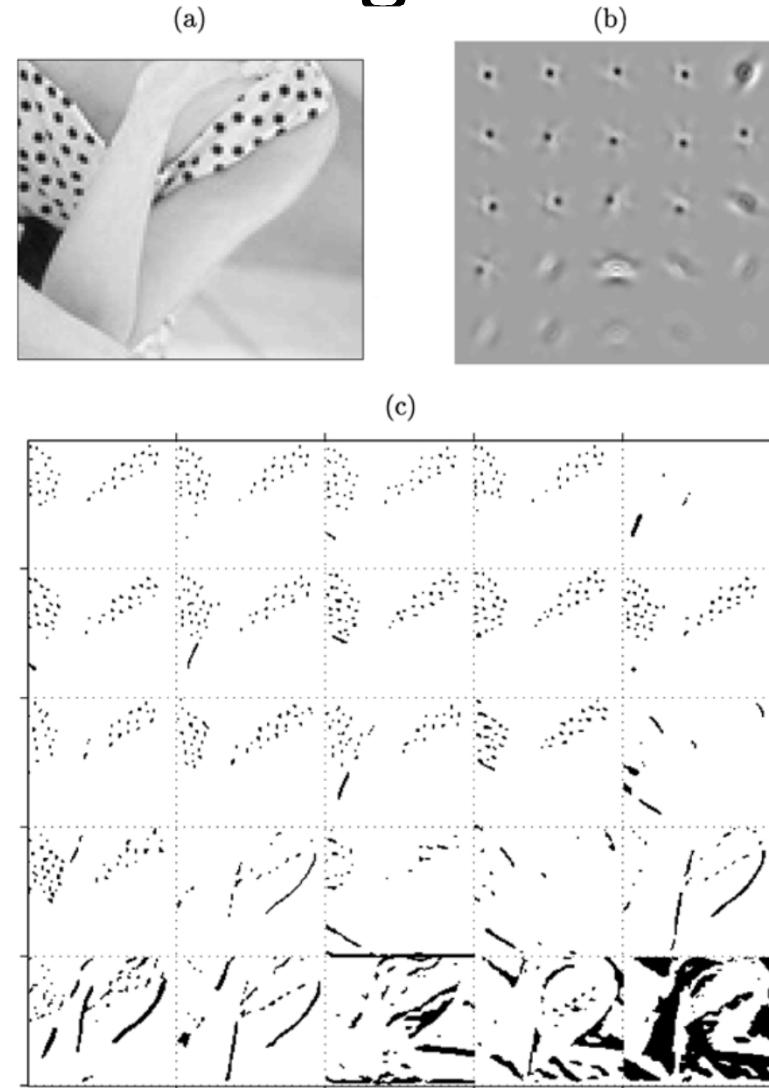
- Texture descriptor is vector of filter bank outputs



J. Malik, S. Belongie, T. Leung and J. Shi. "[Contour and Texture Analysis for Image Segmentation](#)". IJCV 43(1), 7-27, 2001.

Using Texture Features for Segmentation

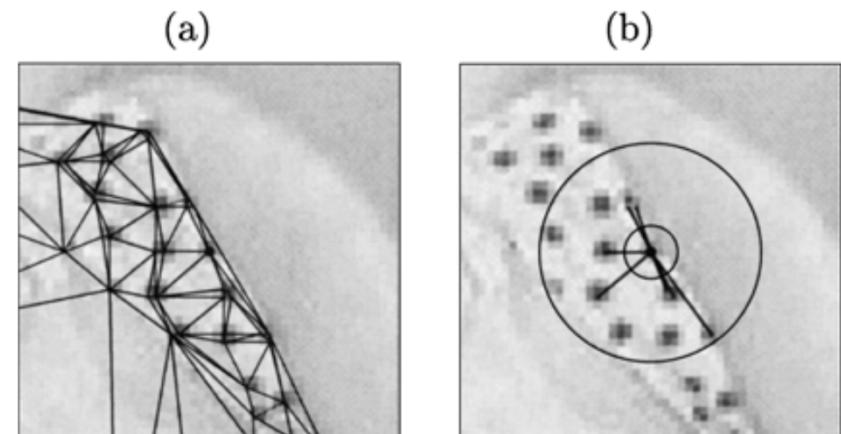
- Texture descriptor is vector of filter bank outputs.
- *Textons* are found by clustering.



Slide credit: Svetlana Lazebnik

Using Texture Features for Segmentation

- Texture descriptor is vector of filter bank outputs.
- *Textons* are found by clustering.
- Affinities are given by similarities of texton histograms over windows given by the “local scale” of the texture .



Slide credit: Svetlana Lazebnik

Results with Color & Texture



Summary: Normalized Cuts

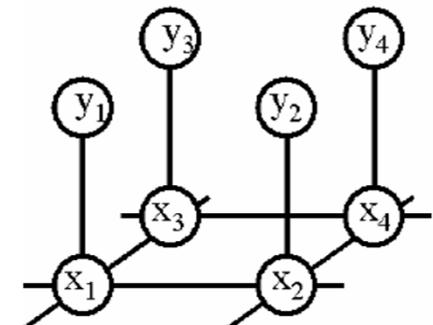
- Pros:
 - Generic framework, flexible to choice of function that computes weights (“affinities”) between nodes
 - Does not require any model of the data distribution
- Cons:
 - Time and memory complexity can be high
 - Dense, highly connected graphs \Rightarrow many affinity computations
 - Solving eigenvalue problem for each cut
 - Preference for balanced partitions
 - If a region is uniform, NCuts will find the modes of vibration of the image dimensions



Slide credit: Kristen Grauman

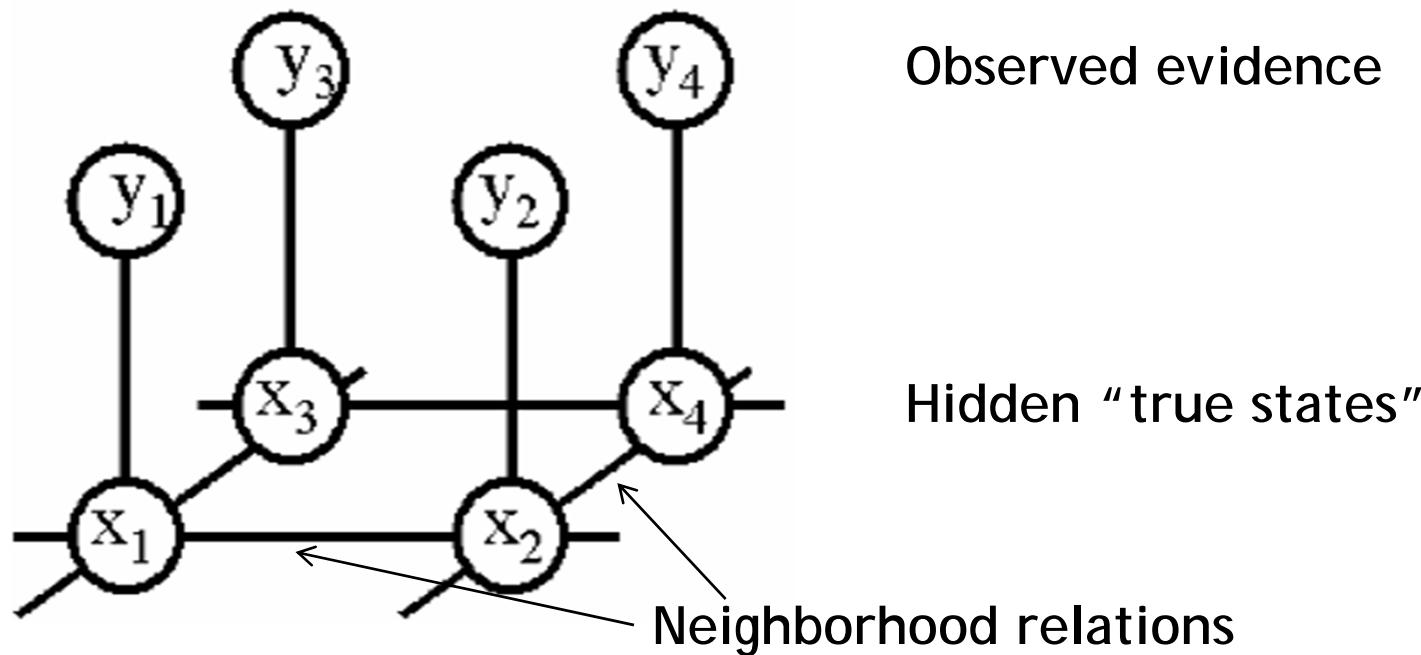
What we will learn today?

- Model free clustering
 - Mean-shift
- Graph theoretic segmentation
 - Normalized Cuts
 - Using texture features
- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation (supp. materials)
 - s-t mincut algorithm (supp. materials)
 - Extension to non-binary case (supp. materials)
 - Applications



Markov Random Fields

- Allow rich probabilistic models for images
- But built in a local, modular way
 - Learn local effects, get global effects out



MRF Nodes as Pixels



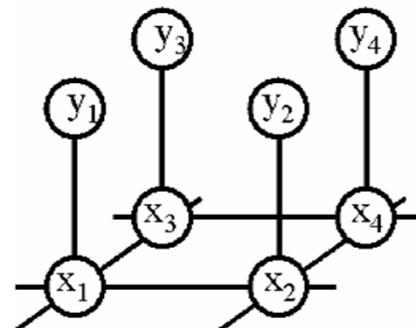
Original image



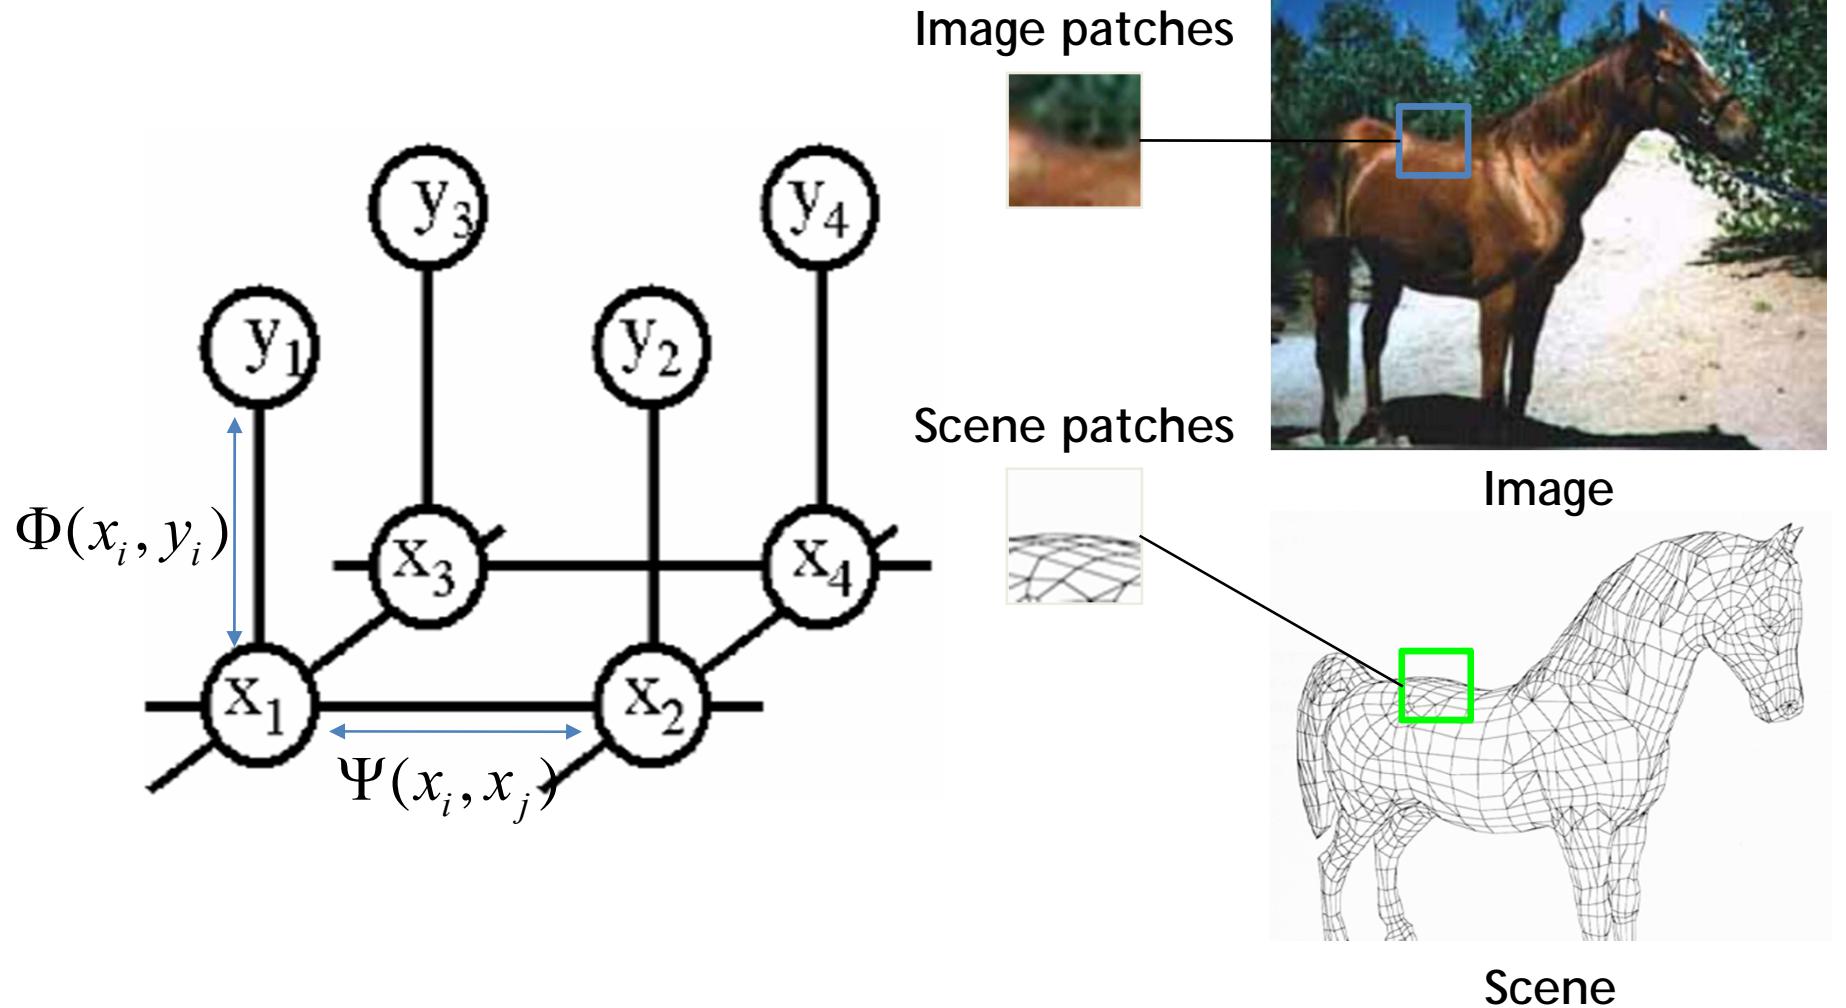
Degraded image



Reconstruction
from MRF modeling
pixel neighborhood
statistics

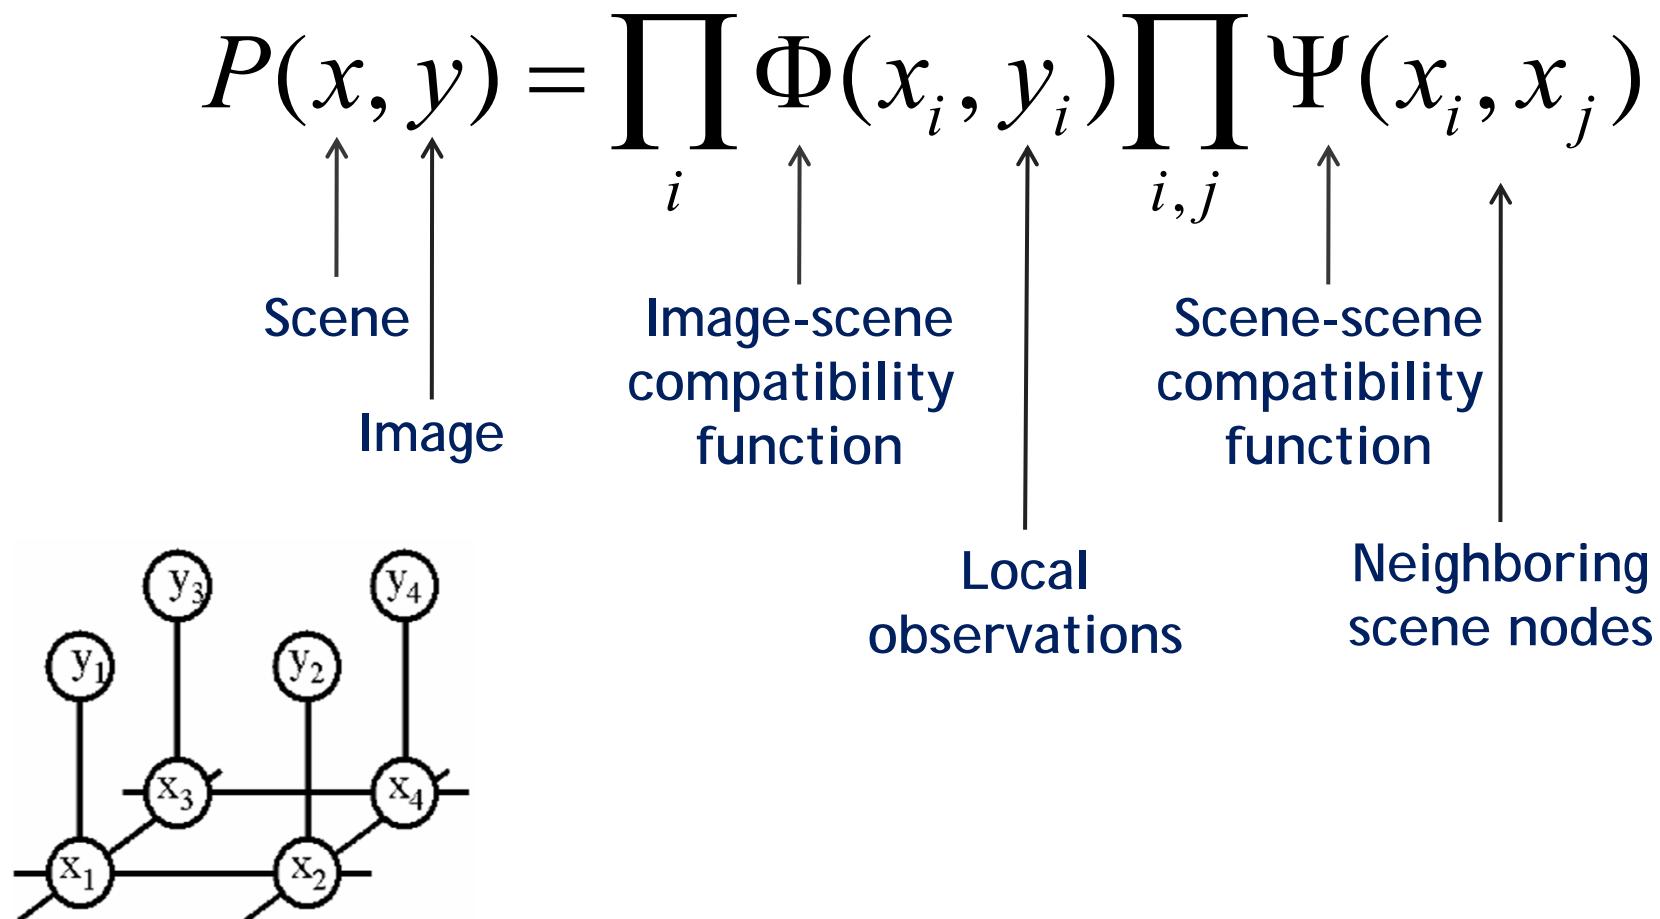


MRF Nodes as Patches



Slide credit: William Freeman

Network Joint Probability



Energy Formulation

- Joint probability

$$P(x, y) = \prod_i \Phi(x_i, y_i) \prod_{i,j} \Psi(x_i, x_j)$$

- Taking the log p(.) turns this into an Energy optimization problem

$$\log P(x, y) = \sum_i \log \Phi(x_i, y_i) + \sum_{i,j} \log \Psi(x_i, x_j)$$

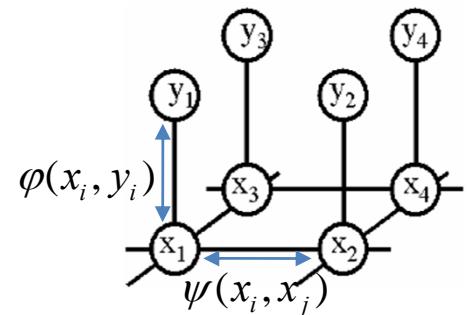
$$E(x, y) = \sum_i \varphi(x_i, y_i) + \sum_{i,j} \psi(x_i, x_j)$$

- This is similar to free-energy problems in statistical mechanics (spin glass theory). We therefore draw the analogy and call E an *energy function*.
- φ and ψ are called *potentials*.

Energy Formulation

- Energy function

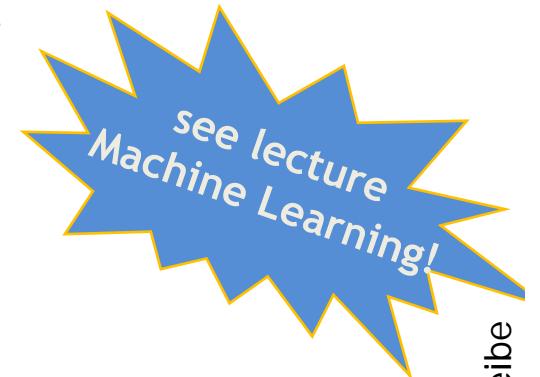
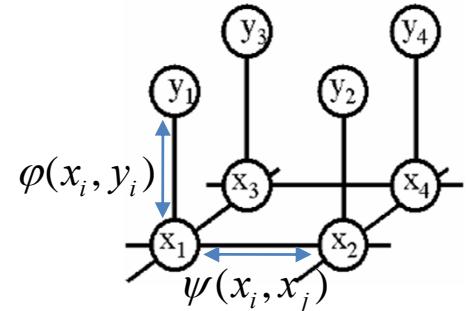
$$E(x, y) = \sum_i \underbrace{\varphi(x_i, y_i)}_{\text{Single-node potentials}} + \sum_{i,j} \underbrace{\psi(x_i, x_j)}_{\text{Pairwise potentials}}$$



- Single-node potentials φ
 - Encode local information about the given pixel/patch
 - How likely is a pixel/patch to belong to a certain class (e.g. foreground/background)?
- Pairwise potentials ψ
 - Encode neighborhood information
 - How different is a pixel/patch's label from that of its neighbor? (e.g. based on intensity/color/texture difference, edges)

Energy Minimization

- Goal:
 - Infer the optimal labeling of the MRF.
- Many inference algorithms are available, e.g.
 - Gibbs sampling, simulated annealing
 - Iterated conditional modes (ICM)
 - Variational methods
 - Belief propagation
- Recently, Graph Cuts have become a popular tool
 - Only suitable for a certain class of energy functions
 - But the solution can be obtained very fast for typical vision problems ($\sim 1\text{MPixel/sec}$).



Slide credit: Bastian Leibe

What we will learn today?

- Graph theoretic segmentation
 - Normalized Cuts
 - Using texture features
 - Extension: Multi-level segmentation
- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation
 - s-t mincut algorithm
 - Extension to non-binary case
 - Applications

GrabCut: live demo



Reported results

- *Included in MS Office 2010 (let's try it)*

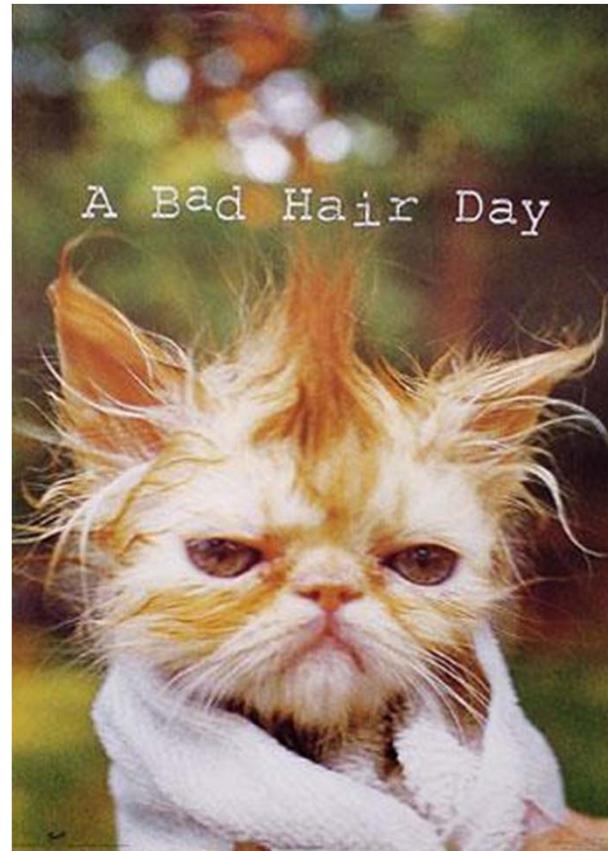
GrabCut: live demo



Reported results

- *Included in MS Office 2010 (let's try it)*

GrabCut: live demo



- *Included in MS Office 2010 (let's try it)*

GraphCut Image Synthesis Results



Source: Vivek Kwatra

Application: Texture Synthesis in the Media



- Currently, still done manually...

Slide credit: Kristen Grauman

Improving Efficiency of Segmentation

- Problem: Images contain many pixels
 - Even with efficient graph cuts, an MRF formulation has too many nodes for interactive results.
- Efficiency trick: Superpixels
 - Group together similar-looking pixels for efficiency of further processing.
 - Cheap, local oversegmentation
 - Important to ensure that superpixels do not cross boundaries
- Several different approaches possible
 - Superpixel code available here
 - <http://www.cs.sfu.ca/~mori/research/superpixels/>

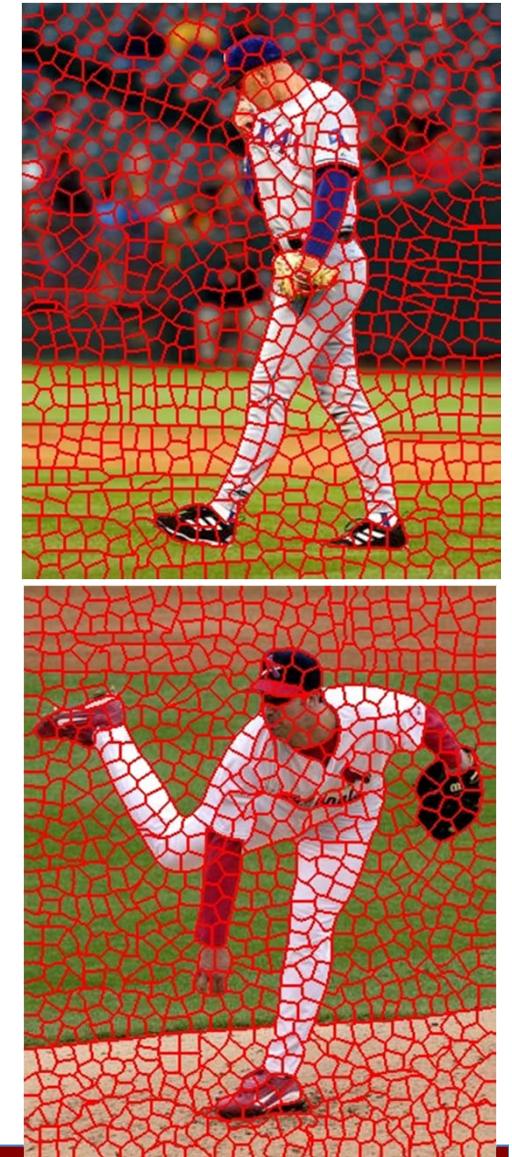


Image source: Greg Mori

Superpixels for Pre-Segmentation

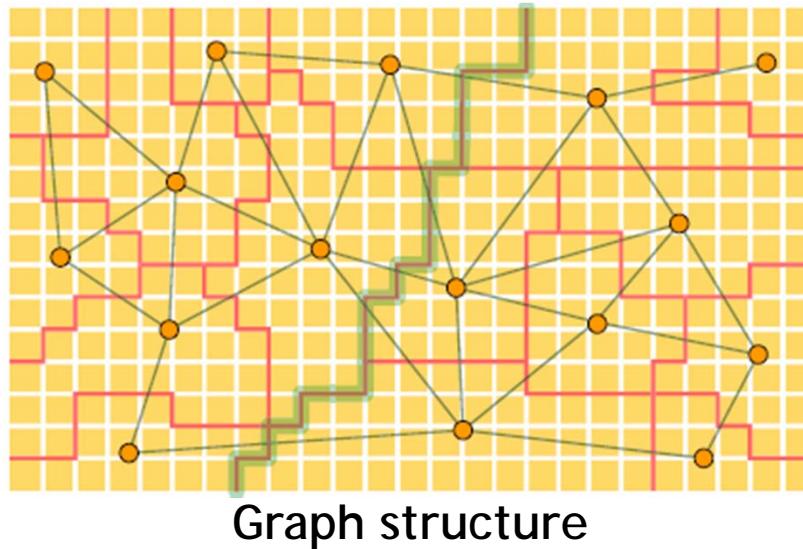
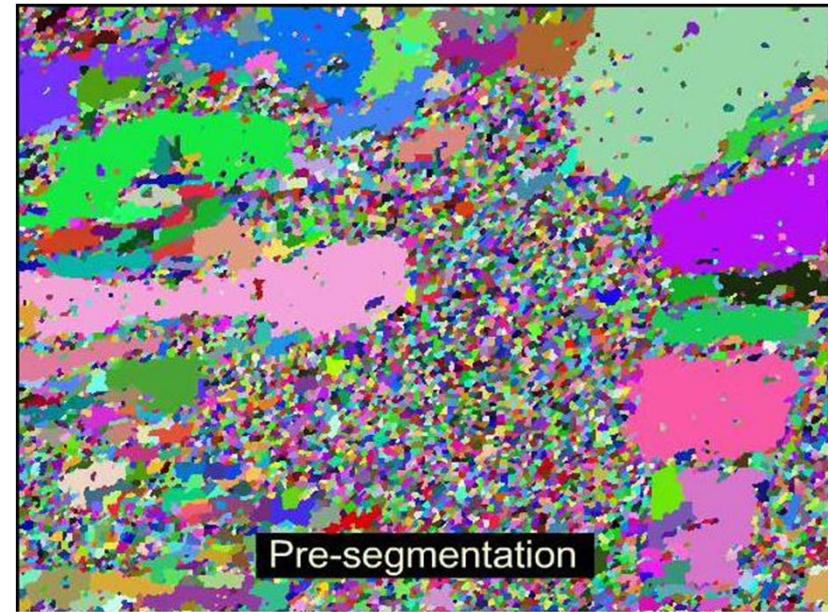


Image	Dimension	Nodes Ratio	Edges Ratio	Lag with Pre-segmentation	Lag without Pre-segmentation
Boy	(408, 600)	10.7	16.8	0.12s	0.57s
Ballet	(440, 800)	11.4	18.3	0.21s	1.39s
Twins	(1024, 768)	20.7	32.5	0.25s	1.82s
Girl	(768, 1147)	23.8	37.6	0.22s	2.49s
Grandpa	(1147, 768)	19.3	30.5	0.22s	3.56s

Speedup

Summary: Graph Cuts Segmentation

- Pros
 - Powerful technique, based on probabilistic model (MRF).
 - Applicable for a wide range of problems.
 - Very efficient algorithms available for vision problems.
 - Becoming a de-facto standard for many segmentation tasks.
- Cons/Issues
 - Graph cuts can only solve a limited class of models
 - Submodular energy functions
 - Can capture only part of the expressiveness of MRFs
 - Only approximate algorithms available for multi-label case

What we have learned today

- Model free clustering
 - Mean-shift
- Graph theoretic segmentation
 - Normalized Cuts
 - Using texture features
- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation (supp. materials)
 - s-t mincut algorithm (supp. materials)
 - Extension to non-binary case (supp. materials)
 - Applications

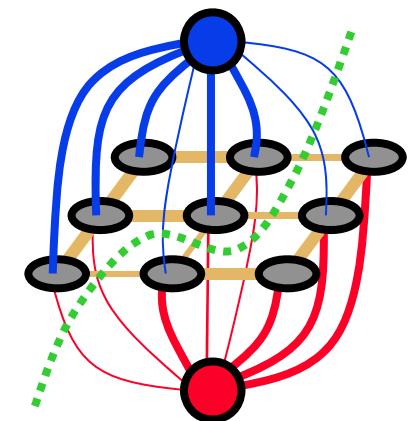


(Midterm materials)

Supplementary materials

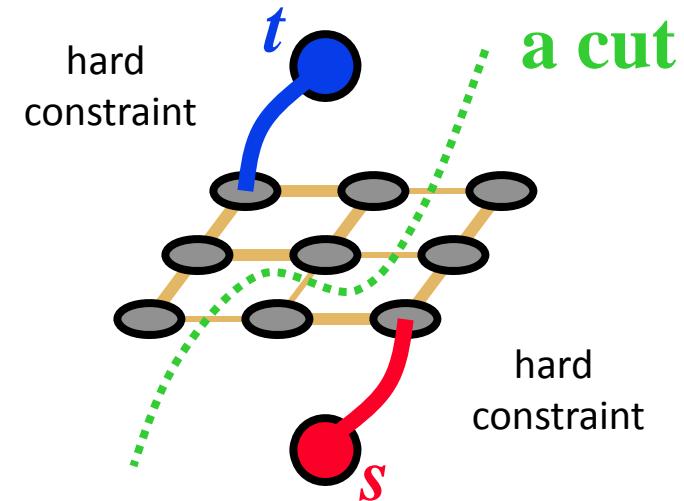
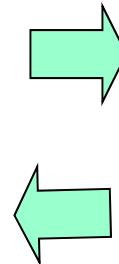
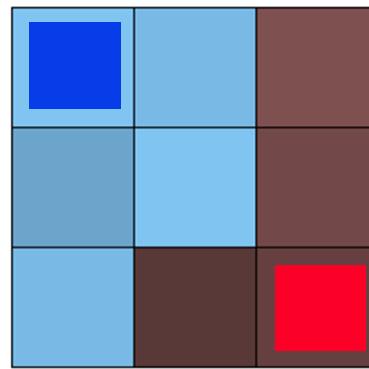
Supplementary materials

- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation (supp. materials)
 - s-t mincut algorithm (supp. materials)
 - Extension to non-binary case (supp. materials)
 - Applications

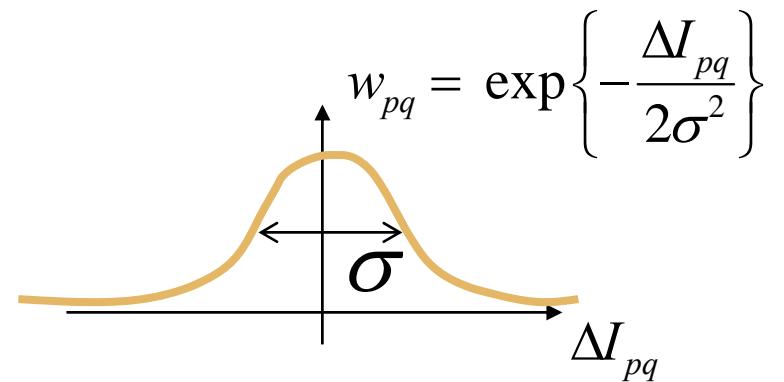


Graph Cuts for Optimal Boundary Detection

- Idea: convert MRF into source-sink graph



Minimum cost cut can be
computed in polynomial time
(max-flow/min-cut algorithms)

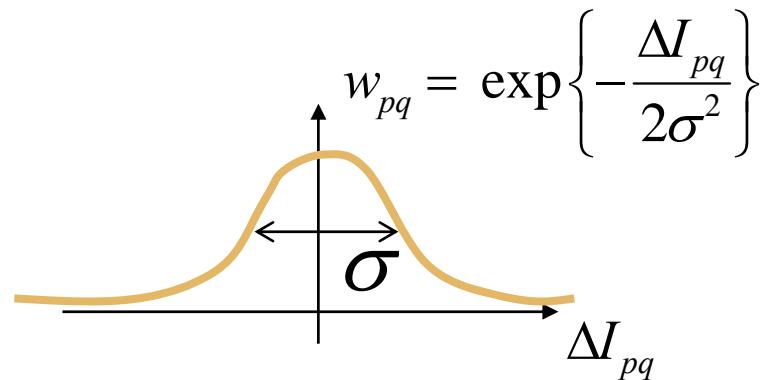
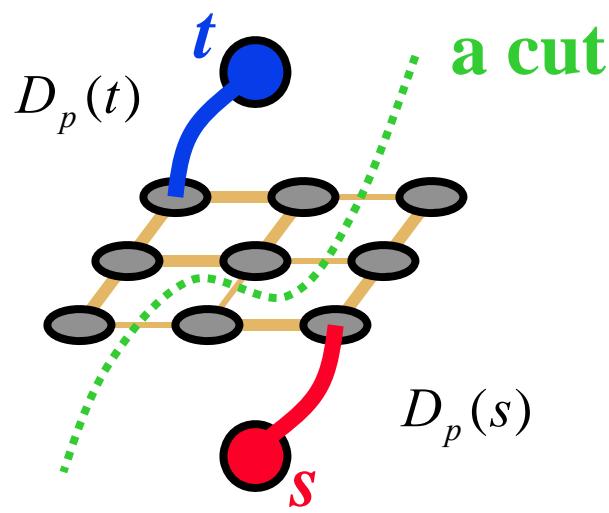


Slide credit: Yuri Boykov

Simple Example of Energy

$$E(L) = \sum_p D_p(L_p) + \sum_{pq \in N} w_{pq} \cdot \delta(L_p \neq L_q)$$

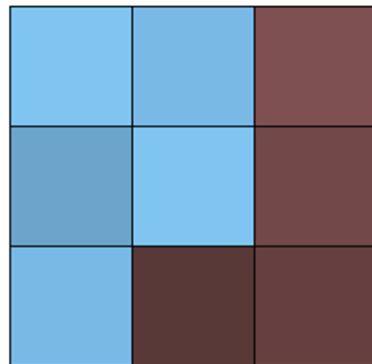
Regional term **Boundary term**
 t-links n-links



$L_p \in \{s, t\}$
(binary object segmentation)

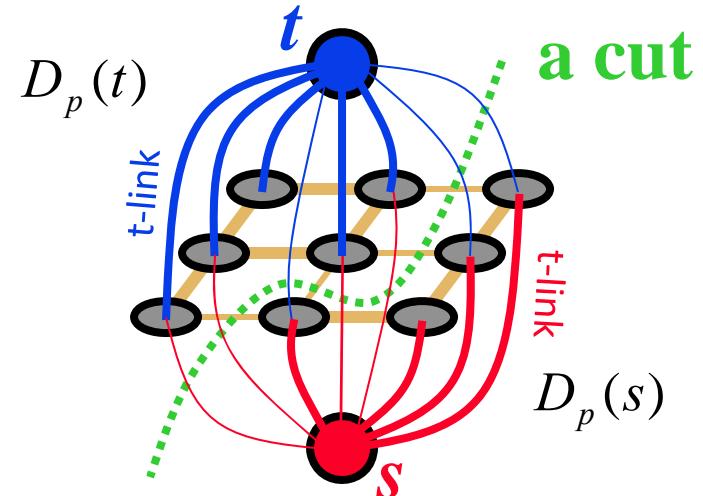
Slide credit: Yuri Boykov

Adding Regional Properties



Regional bias example

Suppose I^s and I^t are given
“expected” intensities
of **object** and **background**

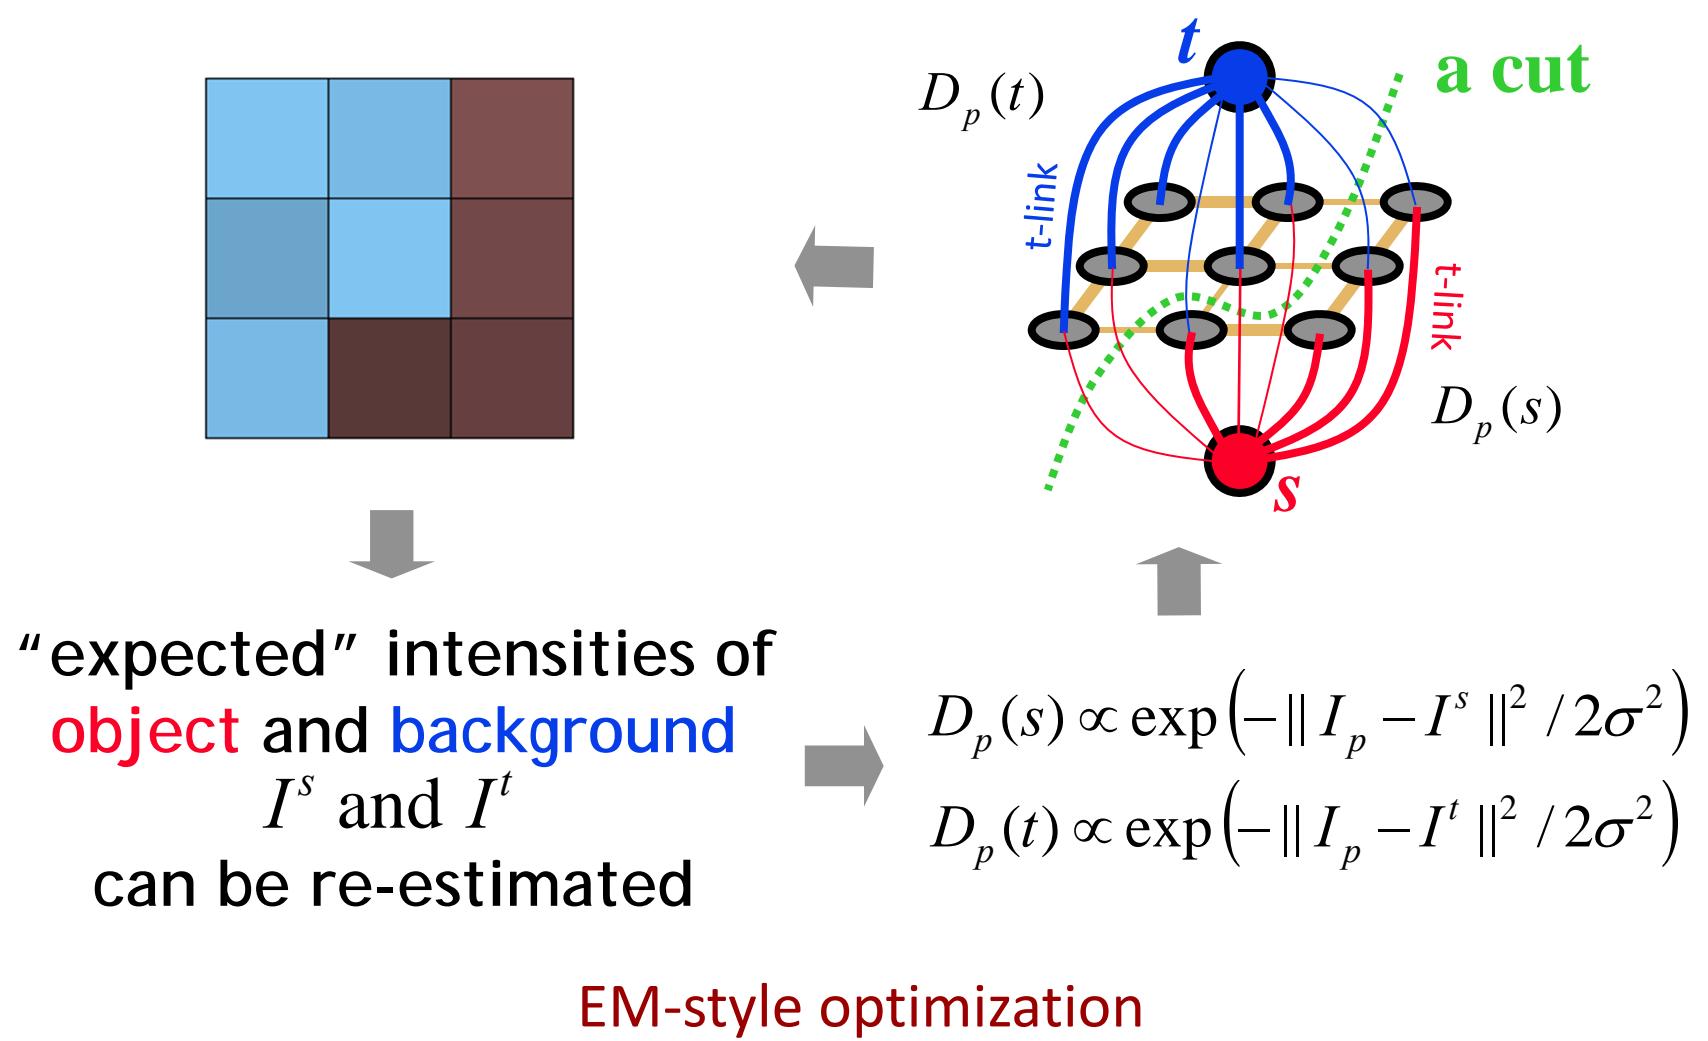


$$D_p(s) \propto \exp\left(-\|I_p - I^s\|^2 / 2\sigma^2\right)$$
$$D_p(t) \propto \exp\left(-\|I_p - I^t\|^2 / 2\sigma^2\right)$$

NOTE: hard constraints are not required, in general.

Slide credit: Yuri Boykov

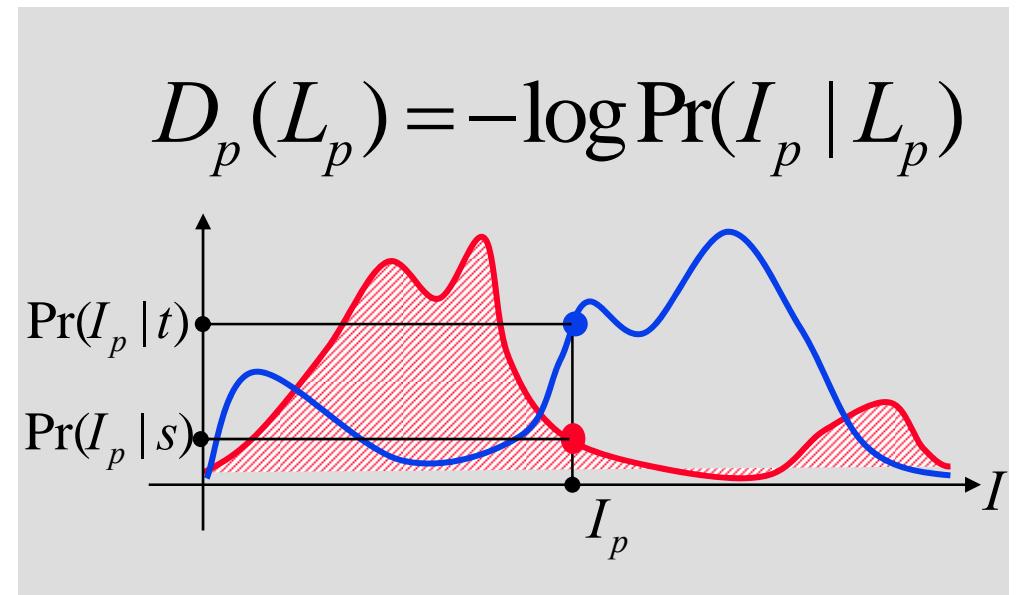
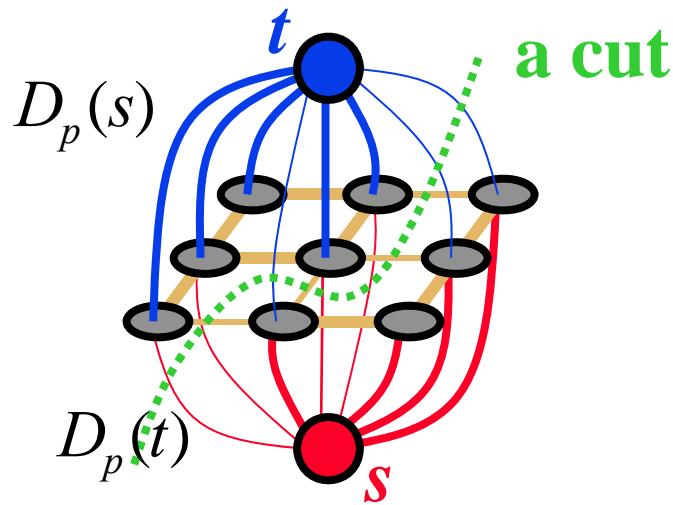
Adding Regional Properties



Slide credit: Yuri Boykov

Adding Regional Properties

- More generally, regional bias can be based on any intensity models of object and background



given object and background intensity histograms

Slide credit: Yuri Boykov

How to Set the Potentials? Some Examples

- Color potentials

- e.g. modeled with a Mixture of Gaussians

$$\pi(x_i, y_i; \theta_\pi) = \log \sum_k \theta_\pi(x_i, k) P(k | x_i) N(y_i; \bar{y}_k, \Sigma_k)$$

- Edge potentials

- e.g. a “contrast sensitive Potts model”

$$\phi(x_i, x_j, g_{ij}(y); \theta_\phi) = -\theta_\phi^T g_{ij}(y) \delta(x_i \neq x_j)$$

where $g_{ij}(y) = e^{-\beta \|y_i - y_j\|^2}$ $\beta = 2 \cdot \text{avg}(\|y_i - y_j\|^2)$

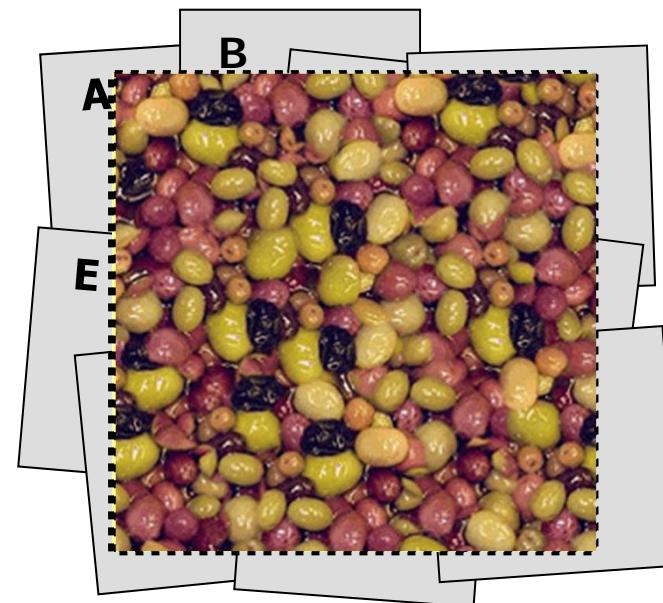
- Parameters θ_π, θ_ϕ need to be learned, too!

[Shotton & Winn, ECCV'06]

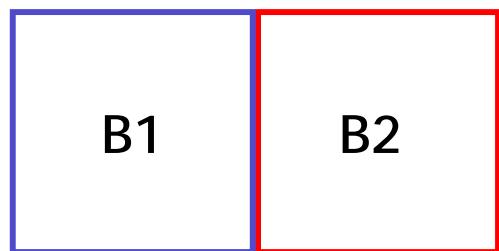
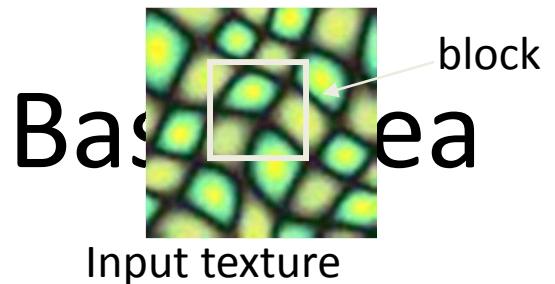
Other Applications: Texture Synthesis

Graph-cut textures

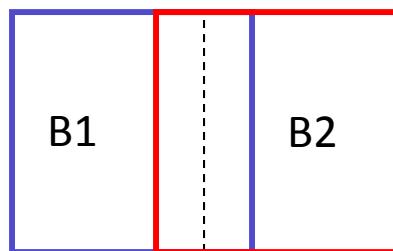
(Kwatra, Schodl, Essa, Bobick 2003)



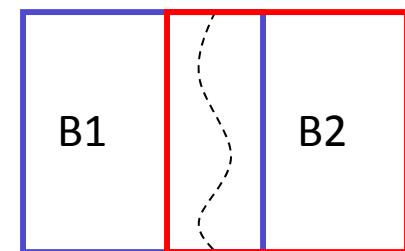
Similar to “image-quilting” (Efros & Freeman, 2001)



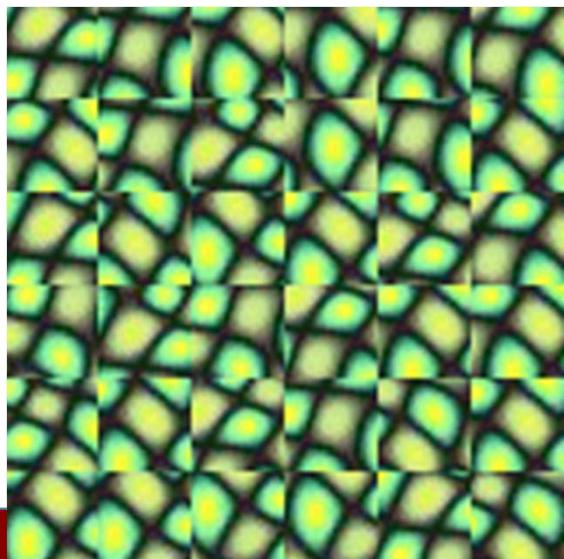
Random placement
of blocks



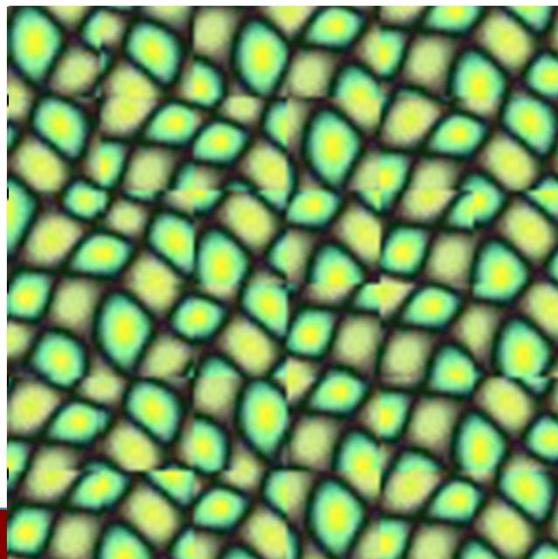
Neighboring blocks
constrained by overlap



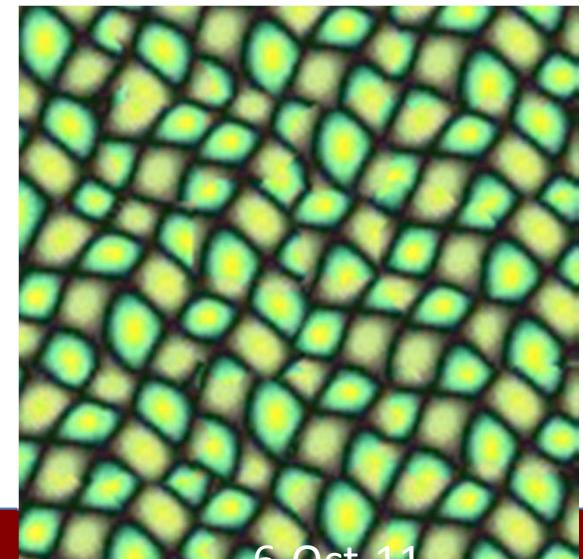
Minimal error
boundary cut



Referer Li



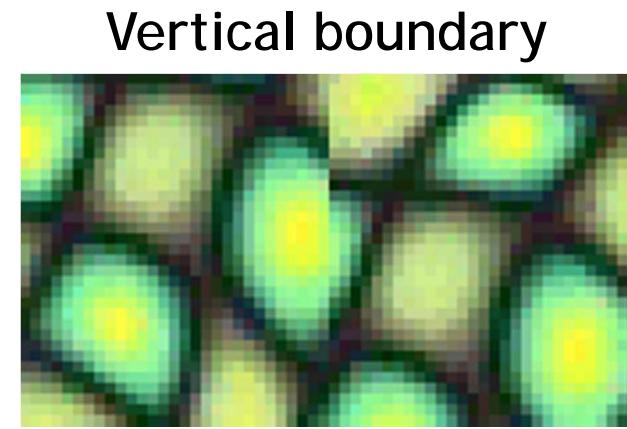
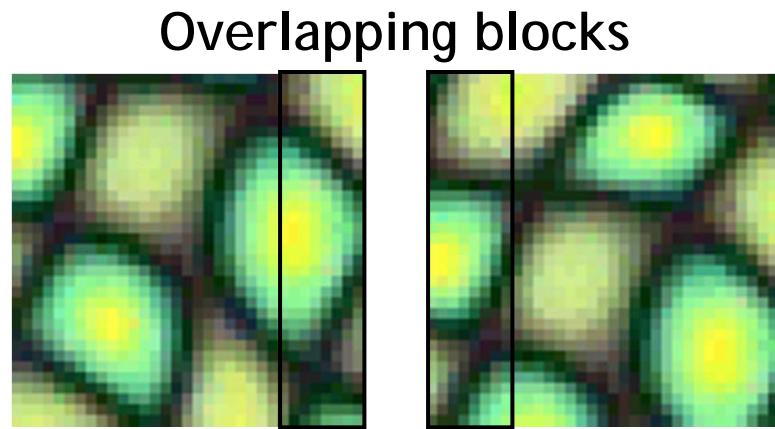
Lecture 6 -



6-Oct-11

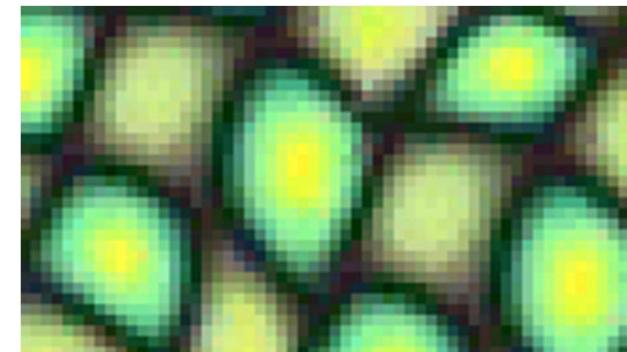
Slide from Alyosha Efros

Minimal Error Boundary



$$\left(\begin{array}{c} \text{[block]} \\ - \\ \text{[block]} \end{array} \right)^2 = \text{[red boundary]}$$

Overlap error



min. error boundary

Lecture 6 -

Fei-Fei Li

6-Oct-11

Slide from Alyosha Efros

GraphCut Texture Synthesis Results



Original fragments

Results (with perspective correction)

Lecture 6 -

6-Oct-11

Source: Vivek Kwatra

Application: Texture Synthesis in the Media



- Currently, still done manually...

How Do we Know?



<http://www.dailymail.co.uk/news/article-22442/878>

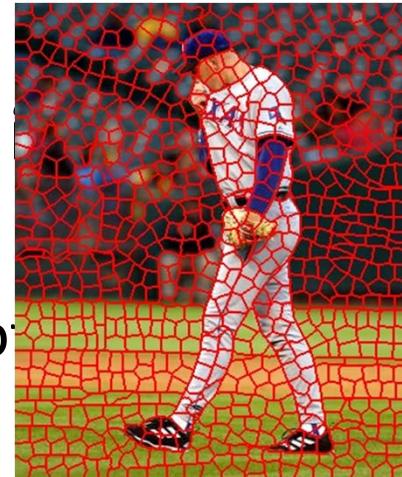
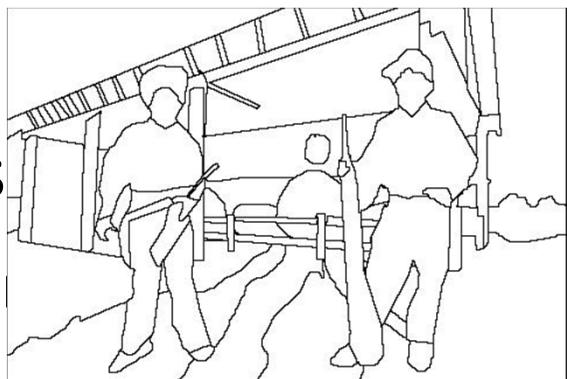
Another Example



<http://thelede.blogs.nytimes.com/2008/07/10/in-an-iranian-image-a-missile-too-many/>

Segmentation: Caveats

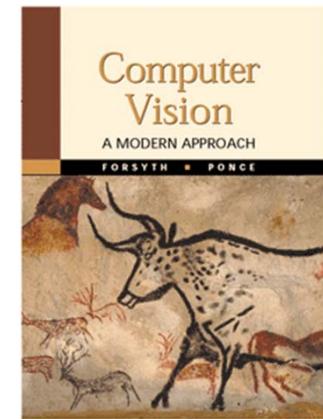
- We've looked at *bottom-up* ways to segment an image into regions, yet finding meaningful segments is intertwined with the recognition problem.
- Often want to avoid making hard decisions too soon
- Difficult to succeed
– Often easiest one-step pipeline.



en
st o
ation
pipeline.

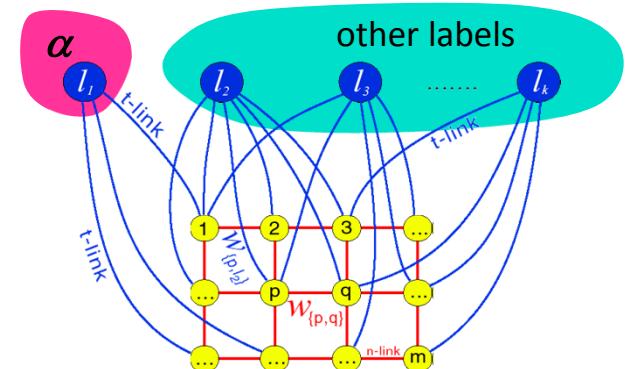
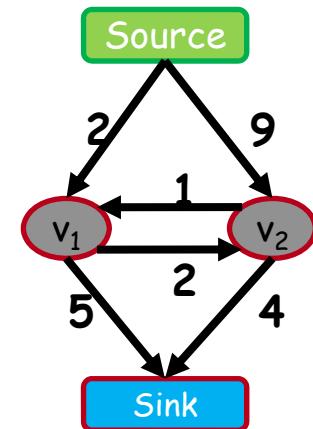
References and Further Reading

- Background information on Normalized Cuts can be found in Chapter 14 of
 - D. Forsyth, J. Ponce,
Computer Vision – A Modern Approach.
Prentice Hall, 2003
- Try the NCuts Matlab code at
 - <http://www.cis.upenn.edu/~jshi/software/>
- Try the GraphCut implementation at
<http://www.adastral.ucl.ac.uk/~vladkolm/software.html>

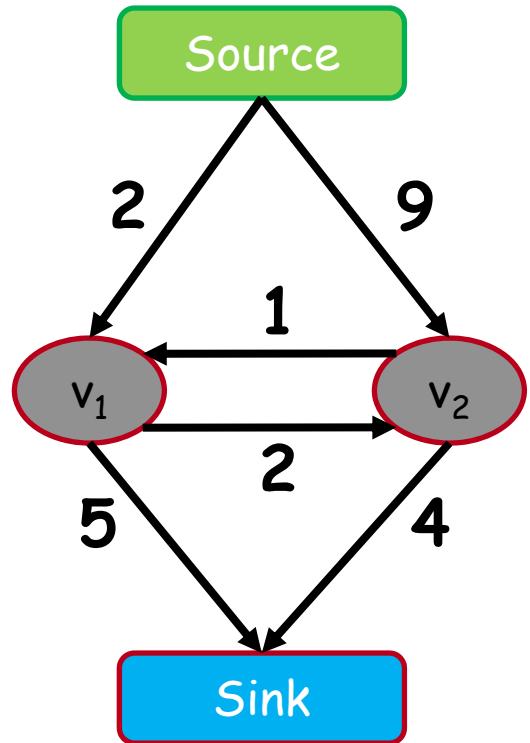


Supplementary materials

- Segmentation as Energy Minimization
 - Markov Random Fields
 - Graph cuts for image segmentation
 - s-t mincut algorithm
 - Extension to non-binary case
 - Applications



How Does it Work? The s-t-Mincut Problem



Graph (V, E, C)

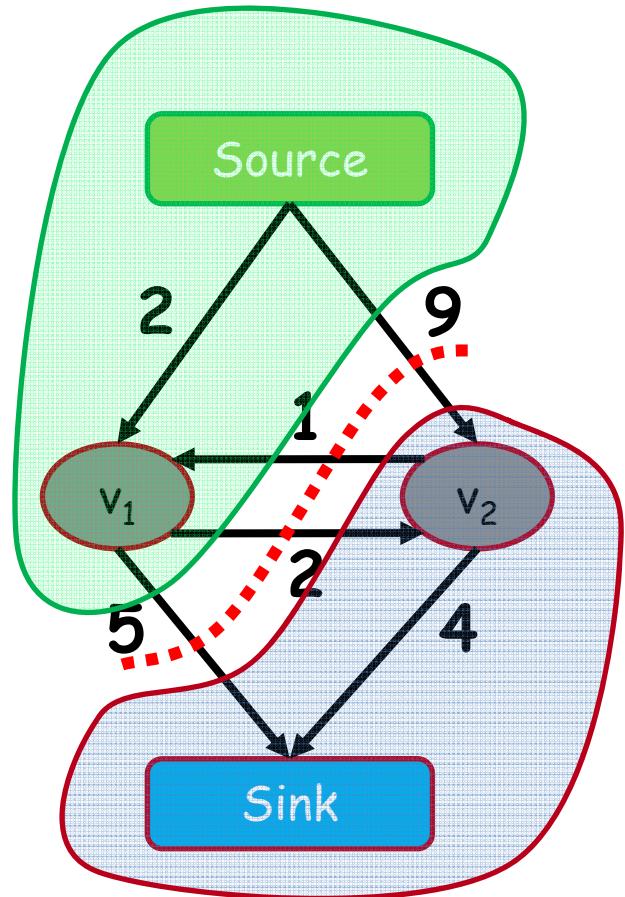
Vertices $V = \{v_1, v_2 \dots v_n\}$

Edges $E = \{(v_1, v_2) \dots\}$

Costs $C = \{c_{(1, 2)} \dots\}$

Slide credit: Pushmeet Kohli

The s-t-Mincut Problem



$$5 + 2 + 9 = 16$$

What is an st-cut?

An st-cut (S, T) divides the nodes between source and sink.

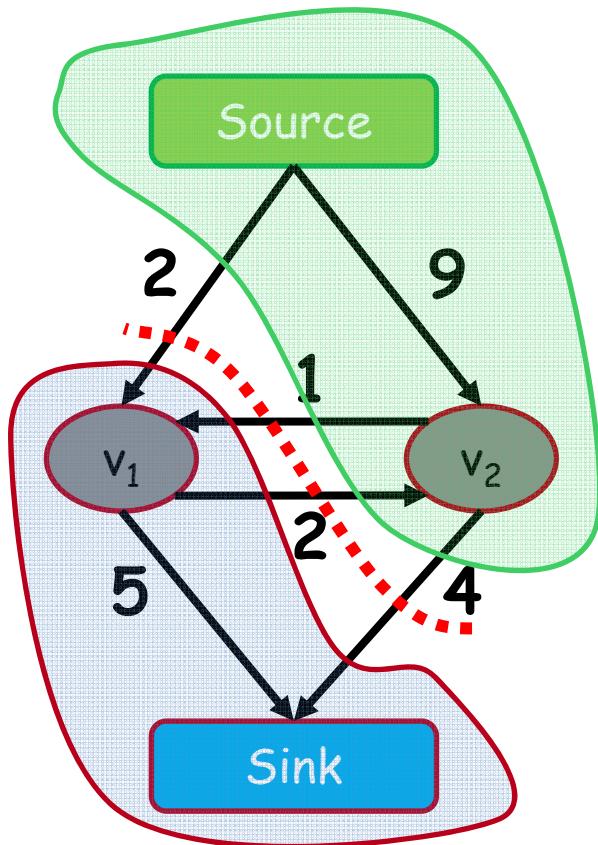
What is the cost of a st-cut?

Sum of cost of all edges going from S to T

Slide credit: Pushmeet Kohli

The s-t-Mincut Problem

What is an st-cut?



$$2 + 1 + 4 = 7$$

An st-cut (S, T) divides the nodes between source and sink.

What is the cost of a st-cut?

Sum of cost of all edges going from S to T

What is the st-mincut?

st-cut with the minimum cost

Slide credit: Pushmeet Kohli

History of Maxflow Algorithms

Augmenting Path and Push-Relabel

year	discoverer(s)	bound
1951	Dantzig	$O(n^2mU)$
1955	Ford & Fulkerson	$O(m^2U)$
1970	Dinitz	$O(n^2m)$
1972	Edmonds & Karp	$O(m^2 \log U)$
1973	Dinitz	$O(nm \log U)$
1974	Karzanov	$O(n^3)$
1977	Cherkassky	$O(n^2m^{1/2})$
1980	Galil & Naamad	$O(nm \log^2 n)$
1983	Sleator & Tarjan	$O(nm \log n)$
1986	Goldberg & Tarjan	$O(nm \log(n^2/m))$
1987	Ahuja & Orlin	$O(nm + n^2 \log U)$
1987	Ahuja et al.	$O(nm \log(n\sqrt{\log U}/m))$
1989	Cheriyan & Hagerup	$E(nm + n^2 \log^2 n)$
1990	Cheriyan et al.	$O(n^3 / \log n)$
1990	Alon	$O(nm + n^{8/3} \log n)$
1992	King et al.	$O(nm + n^{2+\epsilon})$
1993	Phillips & Westbrook	$O(nm(\log_{m/n} n + \log^{2+\epsilon} n))$
1994	King et al.	$O(nm \log_{m/(n \log n)} n)$
1997	Goldberg & Rao	$O(m^{3/2} \log(n^2/m) \log U)$ $O(n^{2/3} m \log(n^2/m) \log U)$

n : #nodes

m : #edges

U : maximum edge weight

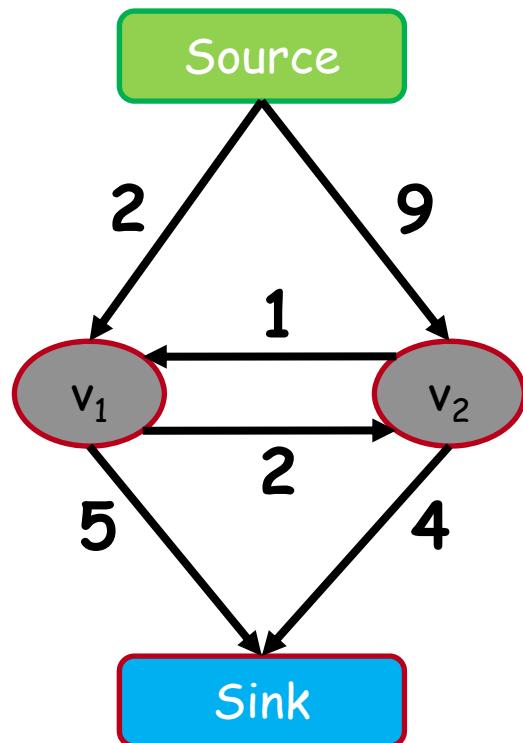
Algorithms assume non-negative edge weights

Slide credit: Andrew Goldberg

How to Compute the s-t-Mincut?

Solve the dual maximum flow problem

Compute the maximum flow between
Source and Sink



Constraints

Edges: Flow < Capacity

Nodes: Flow in = Flow out

Min-cut/Max-flow Theorem

In every network, the maximum flow equals the
cost of the st-mincut

Slide credit: Pushmeet Kohli

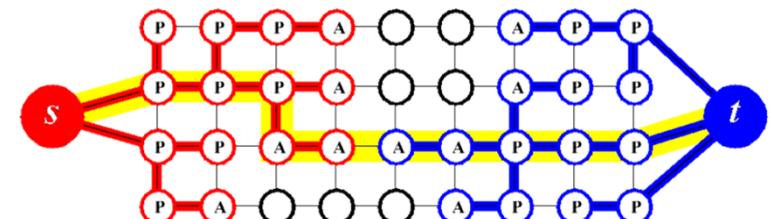
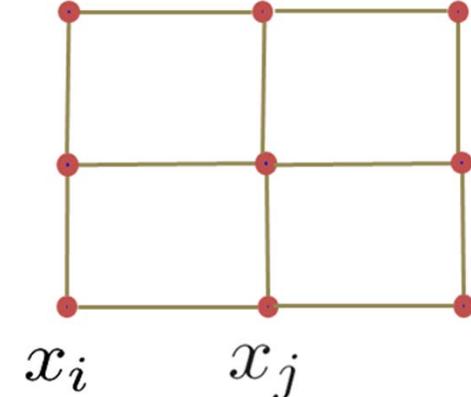
Maxflow in Computer Vision

- Specialized algorithms for vision problems
 - Grid graphs
 - Low connectivity ($m \sim O(n)$)
- Dual search tree augmenting path algorithm

[Boykov and Kolmogorov PAMI 2004]

- Finds approximate shortest augmenting paths efficiently
- High worst-case time complexity
- Empirically outperforms other algorithms on vision problems
- Efficient code available on the web

<http://www.adastral.ucl.ac.uk/~vladkolm/software.html>



Slide credit: Pushmeet Kohli

When Can s-t Graph Cuts Be Applied?

$$E(L) = \sum_p \underset{\text{t-links}}{E_p(L_p)} + \sum_{pq \in N} \underset{\text{n-links}}{E(L_p, L_q)}$$
$$L_p \in \{s, t\}$$

- s-t graph cuts can only globally minimize **binary energies** that are **submodular**. [Boros & Hummer, 2002, Kolmogorov & Zabih, 2004]

$E(L)$ can be minimized by *s-t* graph cuts

$$\iff E(s,s) + E(t,t) \leq E(s,t) + E(t,s)$$

Submodularity ("convexity")

- Non-submodular cases can still be addressed with some optimality guarantees.
 - Current research topic

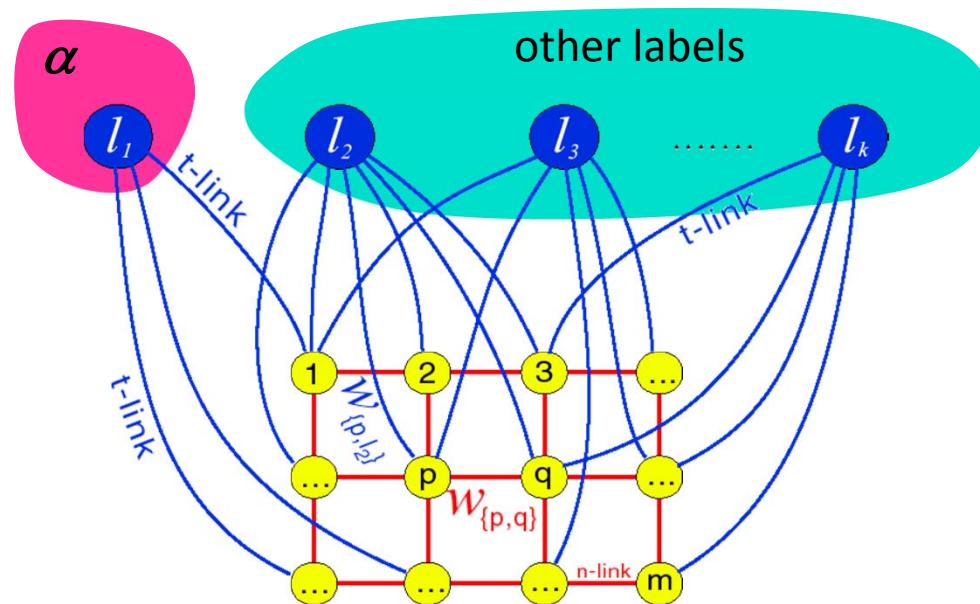
Slide credit: Bastian Leibe

Dealing with Non-Binary Cases

- For image segmentation, the limitation to binary energies is a nuisance.
 ⇒ Binary segmentation only
- We would like to solve also multi-label problems.
 - NP-hard problem with 3 or more labels
- There exist some approximation algorithms which extend graph cuts to the multi-label case
 - α -Expansion
 - $\alpha\beta$ -Swap
- They are no longer guaranteed to return the globally optimal result.
 - But α -Expansion has a guaranteed approximation quality (2-approx) and converges in a few iterations.

α -Expansion Move

- Basic idea:
 - Break multi-way cut computation into a sequence of binary s-t cuts.



Slide credit: Yuri Boykov

α -Expansion Algorithm

1. Start with any initial solution
2. For each label “ α ” in any (e.g. random) order
 1. Compute optimal α -expansion move (s-t graph cuts)
 2. Decline the move if there is no energy decrease
- Stop when no expansion move would decrease energy

Slide credit: Yuri Boykov

α -Expansion Moves

- In each α -expansion a given label “ α ” grabs space from other labels



initial solution

green-expansion

red-expansion

light green-expansion

orange-expansion

yellow-expansion

brown-expansion

dark green-expansion

For each move we choose the expansion that gives the largest decrease in the energy:
binary optimization problem

Slide credit: Yuri Boykov