

Lecture 13: Tracking motion features – optical flow


Professor Fei-Fei Li
Stanford Vision Lab

What we will learn today?

- Introduction
- Optical flow 
- Feature tracking
- Applications
- (Problem Set 3 (Q1))

From images to videos

- A video is a sequence of frames captured over time
- Now our image data is a function of space (x, y) and time (t)




Motion estimation techniques

- Optical flow
 - Recover image motion at each pixel from spatio-temporal image brightness variations (optical flow)
- Feature-tracking
 - Extract visual features (corners, textured areas) and “track” them over multiple frames




Optical flow




Picture courtesy of Selim Temizer - Learning and Intelligent Systems (LIS) Group, MIT

Feature-tracking



Courtesy of Jean-Yves Bouguet – Vision Lab, California Institute of Technology

Feature-tracking



Courtesy of Jean-Yves Bouguet – Vision Lab, California Institute of Technology

Optical flow


- Definition: optical flow is the *apparent* motion of brightness patterns in the image
- Note: apparent motion can be caused by lighting changes without any actual motion
 - Think of a uniform rotating sphere under fixed lighting vs. a stationary sphere under moving illumination

GOAL: Recover image motion at each pixel from optical flow

Source: Silvio Savarese

Estimating optical flow

*x_{init}
(t-1)*




$I(x,y,t)$

- Given two subsequent frames, estimate the apparent motion field $\mathbf{u}(x,y), \mathbf{v}(x,y)$ between them
- motion*
- Key assumptions
 - Brightness constancy:** projection of the same point looks the same in every frame
 - Small motion:** points do not move very far
 - Spatial coherence:** points move like their neighbors

Source: Silvio Savarese

The brightness constancy constraint



- Brightness Constancy Equation:

$$I(x, y, t-1) = I(x + u(x, y), y + v(x, y), t)$$

Linearizing the right side using Taylor expansion:

$$I(x + u, y + v, t) \approx I(x, y, t-1) + I_x \cdot u(x, y) + I_y \cdot v(x, y) + I_t$$

Image derivative along x

$$0 = I(x + u, y + v, t) - I(x, y, t-1) = I_x \cdot u(x, y) + I_y \cdot v(x, y) + I_t$$

Hence, $I_x \cdot u + I_y \cdot v + I_t \approx 0 \rightarrow \nabla I \cdot [u \ v]^T + I_t = 0$

Source: Silvio Savarese

The brightness constancy constraint

Can we use this equation to recover image motion (u, v) at each pixel?


$$\nabla I \cdot [u \ v]^T + I_t = 0$$

- How many equations and unknowns per pixel?
 - One equation (this is a scalar equation!), two unknowns (u, v)

The component of the flow perpendicular to the gradient (i.e., parallel to the edge) cannot be measured


If (u, v) satisfies the equation,
so does $(u+u', v+v')$ if

$$\nabla I \cdot [u' \ v']^T = 0$$





Source: Silvio Savarese

The aperture problem



Source: Silvio Savarese

The aperture problem



Perceived motion

Source: Silvio Savarese

The barber pole illusion



Source: Silvio Savarese

http://en.wikipedia.org/wiki/Barberpole_illusion


The barber pole illusion



Source: Silvio Savarese

http://en.wikipedia.org/wiki/Barberpole_illusion

Aperture problem cont'd



* From Marc Pollefeys COMP 256 2003

Solving the ambiguity...

B. Lucas and T. Kanade. An iterative image registration technique with an application to stereo vision. In *Proceedings of the International Joint Conference on Artificial Intelligence*, pp. 674–679, 1981.

- How to get more equations for a pixel?
- **Spatial coherence constraint:**
- Assume the pixel's neighbors have the same (u, v)
 - If we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix}$$

Source: Silvio Savarese

Lucas-Kanade flow

- Overconstrained linear system:

$$\mathbf{A} \begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix} \quad \begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \end{matrix}$$

A *d* *b*

Source: Silvio Savarese

Conditions for solvability

- When is this system solvable?
 - What if the window contains just a single straight edge?

Source: Silvio Savarese

Lucas-Kanade flow

- Overconstrained linear system

$$\begin{array}{c} \text{25x2} \\ \left[\begin{array}{cc} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{array} \right] \left[\begin{array}{c} u \\ v \end{array} \right] = - \left[\begin{array}{c} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{array} \right] \end{array} \quad \begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \end{matrix}$$

Least squares solution for d given by $(A^T A)^{-1} A^T b$

$$\left[\begin{array}{cc} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{array} \right] \left[\begin{array}{c} u \\ v \end{array} \right] = - \left[\begin{array}{c} \sum I_x I_t \\ \sum I_y I_t \end{array} \right]$$

$$\left. \begin{array}{c} A^T A \\ 25 \times 2 \end{array} \right\} 2 \times 2 \quad A^T b$$

The summations are over all pixels in the $K \times K$ window

Source: Silvio Savarese

Conditions for solvability

- Optimal (u, v) satisfies Lucas-Kanade equation

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A$ $A^T b$

When is This Solvable?

- $A^T A$ should be invertible
- $A^T A$ should not be too small due to noise
 - eigenvalues λ_1 and λ_2 of $A^T A$ should not be too small
- $A^T A$ should be well-conditioned
 - λ_1 / λ_2 should not be too large (λ_1 = larger eigenvalue)

Does this remind anything to you?

Source: Silvio Savarese

$M = A^T A$ is the *second moment matrix* ! (Harris corner detector...)


$$A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

- Eigenvectors and eigenvalues of $A^T A$ relate to edge direction and magnitude
 - The eigenvector associated with the larger eigenvalue points in the direction of fastest intensity change
 - The other eigenvector is orthogonal to it

Source: Silvio Savarese

Interpreting the eigenvalues

Classification of image points using eigenvalues of the second moment matrix:



Edge




$$\sum \nabla I (\nabla I)^T$$

- gradients very large or very small
- large λ_1 , small λ_2

Source: Silvio Savarese

Low-texture region



$$\sum \nabla I (\nabla I)^T$$

- gradients have small magnitude
- small λ_1 , small λ_2

Source: Silvio Savarese

High-texture region



$$\sum \nabla I (\nabla I)^T$$

- gradients are different, large magnitudes
- large λ_1 , large λ_2

Source: Silvio Savarese

What are good features to track?

- Can measure “quality” of features from just a single image
- Hence: tracking Harris corners (or equivalent) guarantees small error sensitivity!

→ Implemented in Open CV

Source: Silvio Savarese

Recap

- Key assumptions (Errors in Lucas-Kanade)

- **Small motion:** points do not move very far
- **Brightness constancy:** projection of the same point looks the same in every frame
- **Spatial coherence:** points move like their neighbors

Source: Silvio Savarese




Revisiting the small motion assumption



- Is this motion small enough?
 - Probably not—it's much larger than one pixel (2^{nd} order terms dominate)
 - How might we solve this problem?

* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Reduce the resolution!




* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Multi-resolution Lucas Kanade Algorithm

- Compute ‘simple’ LK at highest level
- At level i
 - Take flow u_{i-1}, v_{i-1} from level $i-1$
 - bilinear interpolate it to create u_i^*, v_i^* matrices of twice resolution for level i
 - multiply u_i^*, v_i^* by 2
 - compute f_t from a block displaced by $u_i^*(x,y), v_i^*(x,y)$
 - Apply LK to get $u_i'(x, y), v_i'(x, y)$ (the correction in flow)
 - Add corrections u_i', v_i' , i.e. $u_i = u_i^* + u_i'$,
 $v_i = v_i^* + v_i'$.

Source: Silvio Savarese

Coarse-to-fine optical flow estimation




Gaussian pyramid of image 1

$u=1.25 \text{ pixels}$

$u=2.5 \text{ pixels}$

$u=5 \text{ pixels}$

$u=10 \text{ pixels}$



Gaussian pyramid of image 2


Source: Silvio Savarese

Iterative Refinement


- Iterative Lukas-Kanade Algorithm
 1. Estimate velocity at each pixel by solving Lucas-Kanade equations
 2. Warp $I(t-1)$ towards $I(t)$ using the estimated flow field
 - use *image warping techniques*
 3. Repeat until convergence

* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Coarse-to-fine optical flow estimation




Optical Flow Results



* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Optical Flow Results



- <http://www.ces.clemson.edu/~stb/klt/>
- OpenCV

* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Recap

- Key assumptions (Errors in Lucas-Kanade)
 - **Small motion:** points do not move very far
 - **Brightness constancy:** projection of the same point looks the same in every frame
 - **Spatial coherence:** points move like their neighbors

Source: Silvio Savarese

Motion segmentation

- How do we represent the motion in this scene?





Source: Silvio Savarese

Motion segmentation

J. Wang and E. Adelson. Layered Representation for Motion Analysis. *CVPR 1993*.


- Break image sequence into “layers” each of which has a coherent (affine) motion



Source: Silvio Savarese

What are layers?

- Each layer is defined by an alpha mask and an affine motion model



J. Wang and E. Adelson. [Layered Representation for Motion Analysis](#). CVPR 1993.

Source: Silvio Savarese


Affine motion

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

- Substituting into the brightness constancy equation:

$$I_x \cdot u + I_y \cdot v + I_t \approx 0$$




Source: Silvio Savarese

Affine motion

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

- Substituting into the brightness constancy equation:



$$I_x(a_1 + a_2 x + a_3 y) + I_y(a_4 + a_5 x + a_6 y) + I_t \approx 0$$



- Each pixel provides 1 linear constraint in 6 unknowns
- Least squares minimization:

$$Err(\vec{a}) = \sum [I_x(a_1 + a_2 x + a_3 y) + I_y(a_4 + a_5 x + a_6 y) + I_t]^2$$

Source: Silvio Savarese

How do we estimate the layers?

- 1. Obtain a set of initial affine motion hypotheses
 - Divide the image into blocks and estimate affine motion parameters in each block by least squares
 - Eliminate hypotheses with high residual error
 - Map into motion parameter space
 - Perform k-means clustering on affine motion parameters
 - Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene



Source: Silvio Savarese

How do we estimate the layers?

- 1. Obtain a set of initial affine motion hypotheses
 - Divide the image into blocks and estimate affine motion parameters in each block by least squares
 - Eliminate hypotheses with high residual error
 - Map into motion parameter space
 - Perform k-means clustering on affine motion parameters
 - Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene


2. Iterate until convergence:

- Assign each pixel to best hypothesis
 - Pixels with high residual error remain unassigned

Source: Silvio Savarese

How do we estimate the layers?

- 1. Obtain a set of initial affine motion hypotheses
 - Divide the image into blocks and estimate affine motion parameters in each block by least squares
 - Eliminate hypotheses with high residual error
 - Map into motion parameter space
 - Perform k-means clustering on affine motion parameters
 - Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene



How do we estimate the layers?


- 1. Obtain a set of initial affine motion hypotheses
 - Divide the image into blocks and estimate affine motion parameters in each block by least squares
 - Eliminate hypotheses with high residual error
 - Map into motion parameter space
 - Perform k-means clustering on affine motion parameters
 - Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene

2. Iterate until convergence:

- Assign each pixel to best hypothesis
 - Pixels with high residual error remain unassigned
- Perform region filtering to enforce spatial constraints
- Re-estimate affine motions in each region

Source: Silvio Savarese

Example result



Source: Silvio Savarese

J. Wang and E. Adelson. [Layered Representation for Motion Analysis](#). CVPR 1993.

Tracking



Sources: Kristen Grauman, Deva Ramanan

What we will learn today?

- Introduction
- Optical flow
- Feature tracking
- Applications

Motion estimation techniques

- Optical flow
 - Recover image motion at each pixel from spatio-temporal image brightness variations (optical flow)
- Feature-tracking
 - Extract visual features (corners, textured areas) and “track” them over multiple frames
 - Shi-Tomasi feature tracker
 - Tracking with dynamics



Source: Silvio Savarese

Feature tracking

- So far, we have only considered optical flow estimation in a pair of images
- If we have more than two images, we can compute the optical flow from each frame to the next
- Given a point in the first image, we can in principle reconstruct its path by simply “following the arrows”

Source: Silvio Savarese

Tracking challenges

- Ambiguity of optical flow
 - Find good features to track
- Large motions
 - Discrete search instead of Lucas-Kanade
- Changes in shape, orientation, color
 - Allow some matching flexibility
- Occlusions, dis-occlusions
 - Need mechanism for deleting, adding new features
- Drift – errors may accumulate over time
 - Need to know when to terminate a track

Source: Silvio Savarese

Shi-Tomasi feature tracker

J. Shi and C. Tomasi. [Good Features to Track](#). CVPR 1994.

- Find good features using eigenvalues of second-moment matrix
 - Key idea: “good” features to track are the ones that can be tracked reliably
- From frame to frame, track with Lucas-Kanade and a pure *translation* model
 - More robust for small displacements, can be estimated from smaller neighborhoods
- Check consistency of tracks by *affine* registration to the first observed instance of the feature
 - Affine model is more accurate for larger displacements
 - Comparing to the first frame helps to minimize drift

Source: Silvio Savarese

Tracking example



Figure 1: Three frame details from Woody Allen's *Manhattan*. The details are from the 1st, 11th, and 21st frames of a subsequence from the movie.



Figure 2: The traffic sign windows from frames 1,6,11,16,21 as tracked (top), and warped by the computed deformation matrices (bottom).

Source: Silvio Savarese


Tracking with dynamics

- Key idea: Given a model of expected motion, predict where objects will occur in next frame, even before seeing the image
 - Restrict search for the object
 - Improved estimates since measurement noise is reduced by trajectory smoothness

Source: Silvio Savarese

Tracking with dynamics

initial position prediction measurement update



The Kalman filter:

- Method for tracking linear dynamical models in Gaussian noise
- The predicted/corrected state distributions are Gaussian
 - Need to maintain the mean and covariance
 - Calculations are easy (all the integrals can be done in closed form)

Source: Silvio Savarese

2D Target tracking using Kalman filter in MATLAB

by [AliReza KashaniPour](#)

<http://www.mathworks.com/matlabcentral/fileexchange/14243>



Source: Silvio Savarese


What we will learn today?

- Introduction
- Optical flow
- Feature tracking
- Applications

Uses of motion

- Tracking features
- Segmenting objects based on motion cues
- Learning dynamical models
- Improving video quality
 - Motion stabilization
 - Super resolution
- Tracking objects
- Recognizing events and activities


Estimating 3D structure



Source: Silvio Savarese

Segmenting objects based on motion cues



- Background subtraction
 - A static camera is observing a scene
 - Goal: separate the static *background* from the moving *foreground*



Source: Silvio Savarese

Segmenting objects based on motion cues


- Motion segmentation
 - Segment the video into multiple *coherently* moving objects



Source: Silvio Savarese

S. J. Pundlik and S. T. Birchfield, Motion Segmentation at Any Speed,
Proceedings of the British Machine Vision Conference (BMVC) 2006

Tracking objects



Z.Yin and R.Collins, "On-the-fly Object Modeling while Tracking," *IEEE Computer Vision and Pattern Recognition (CVPR '07)*, Minneapolis, MN, June 2007.

Source: Silvio Savarese

Synthesizing dynamic textures



Super-resolution

Example: A set of low quality images

Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S	Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S	Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S
Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S	Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S	Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S
Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S	Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S	Most of the test data o couple of exceptions. 1 low-temperature solder investigated (or some c manufacturing technol nonwetting of 40In40S microstructural coarse mal cycling of 58Bi42S

Source: Silvio Savarese

Super-resolution

Each of these images looks like this:

Most of the test data of a couple of exceptions. The low-temperature solder investigated (or some manufacturing technology) is wetting of 40Ni06Si microstructural coarse and cycling of 58Ti42S.

Source: Silvio Savarese


Super-resolution

The recovery result:

Most of the test data obtained in this work are in agreement with the model, with a couple of exceptions. The first exception concerns the low-temperature solder joints, which were not investigated (or some data were not available) in the manufacturing technology used. The second exception concerns the nonwetting of 40In40Sn solder joints, due to the microstructural coarsening induced by thermal cycling of 58Bi42Sb solder joints.

Source: Silvio Savarese


Recognizing events and activities



D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.

Source: Silvio Savarese


Recognizing events and activities



Juan Carlos Niebles, Hongcheng Wang and Li Fei-Fei, **Unsupervised Learning of Human Action Categories Using Spatial-Temporal Words**, ([BMVC](#)), Edinburgh, 2006.

Recognizing events and activities

Crossing – Talking – Queuing – Dancing – jogging




W. Choi & K. Shahid & S. Savarese WMC 2010

Source: Silvio Savarese



W. Choi, K. Shahid, S. Savarese, "What are they doing? : Collective Activity Classification Using Spatio-Temporal Relationship Among People", 9th International Workshop on Visual Surveillance (VSWS09) in conjunction with ICCV 09

Optical flow without motion!



What we have learned today?

- Introduction
- Optical flow
- Feature tracking
- Applications
- (Problem Set 3 (Q1))